

CSE 417T

# Introduction to Machine Learning

Lecture 10

Instructor: Chien-Ju (CJ) Ho

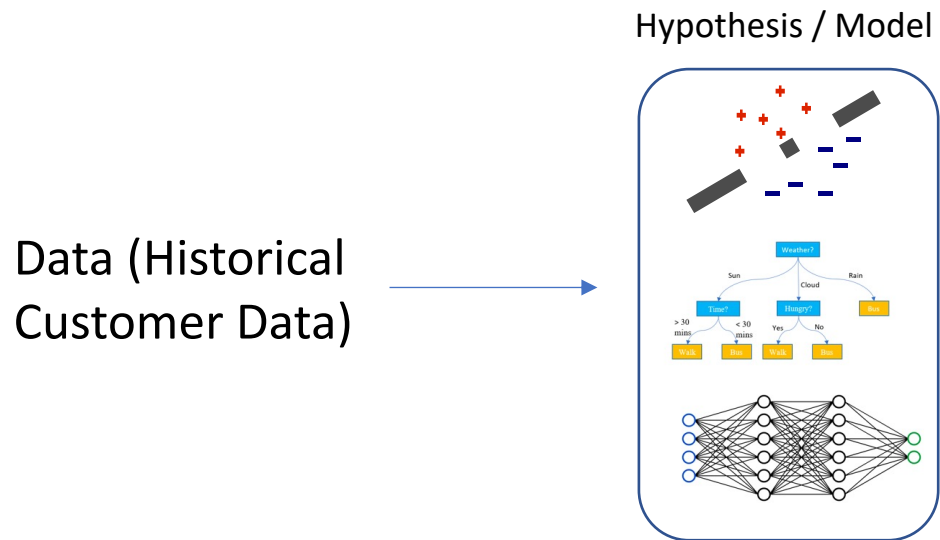
# Logistics

- Homework 2: due on **Oct 7** (Friday)
- Exam 1: **October 27 (Thursday)**
  - Topics: LFD Chapters 1 to 5
  - Timed exam (75 min) during lecture time
  - Location TBD
  - Closed-book exam with 2 letter-size cheat sheets allowed (4 pages in total)
    - No format limitations (it can be typed, written, or a combination)
- Homework 3 will be posted some time next week

Recap

# What We Have Taught So Far

- The explanation of “machine learning” from the first lecture





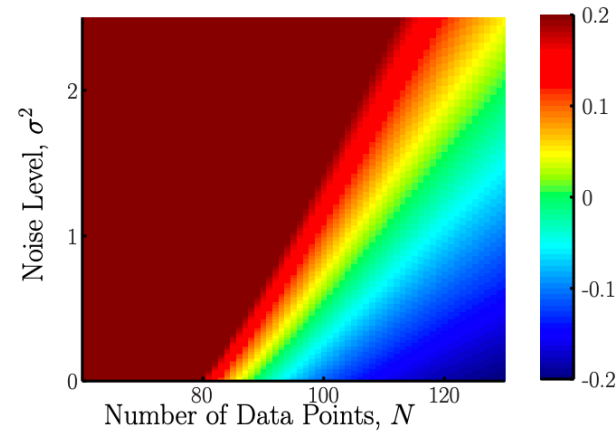

Find a hypothesis that “fits” the data  
(The process requires a lot of computation)

Our progress so far

- Generalization of learning
  - What to say about  $E_{out}(g)$  from  $E_{in}(g)$
- How to find  $g$ 
  - Using linear models as examples
  - Focus on  $g = \operatorname{argmin}_{h \in H} E_{in}(g)$

Seems to make sense, but...

# Overfitting



Number of data points $\uparrow$	Overfitting $\downarrow$
Noise $\uparrow$	Overfitting $\uparrow$
Target complexity $\uparrow$	Overfitting $\uparrow$

# Today's Lecture

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.  
Let me know if you spot errors.

# Overfitting and Its Cures

- Overfitting


- Fitting the data more than is warranted
- Fitting the noise instead of the pattern of the data
- Decreasing  $E_{in}$  but getting larger  $E_{out}$
- When  $H$  is too strong, but  $N$  is not large enough

- Regularization

- Intuition: Constrain  $H$  to make overfitting less likely to happen
- (Topic of this lecture)

- Validation

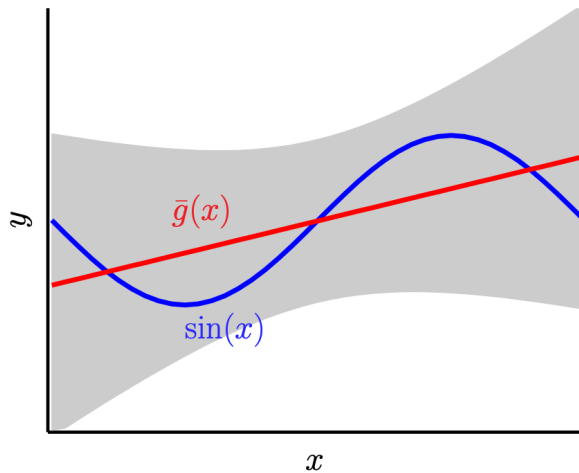
- Intuition: Reserve data to estimate  $E_{out}$
- (Focus of next lecture)



# Regularization (Constrain $H$ )

- Informal example:

- Regression;  $f = \sin(\pi x)$ ;  $H = \{h(x) = ax + b\}$ ;  $N = 2$

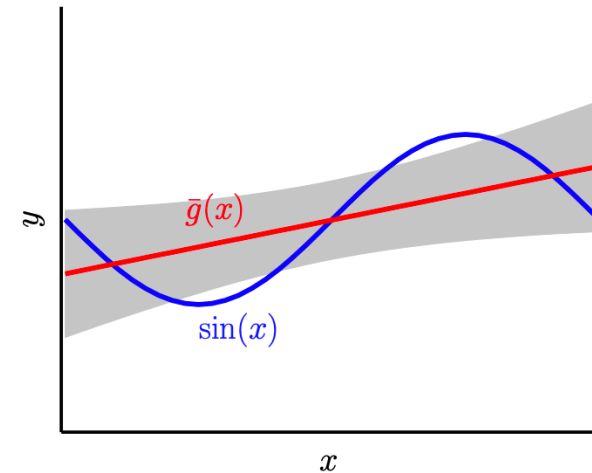


no regularization

bias = 0.21

var = 1.69

Regularization:  
**Constrain** the hypothesis set to  
avoid large  $a$  and  $b$



regularization

bias = 0.23

var = 0.33

How to do this in a principled way?



# Hard Constraints

- We have seen hard constraints already

$$H_2 = \{h(x) = w_0 + w_1x + w_2x^2\}$$

$$H_{10} = \{h(x) = w_0 + w_1x + w_2x^2 + \cdots + w_{10}x^{10}\}$$

- $H_2$  can be written as constrained  $H_{10}$

$$H_2 = \{h \in H_{10} \text{ and } w_3 = w_4 = \cdots = w_{10} = 0\}$$



Constraints

# Soft-Order Constraints

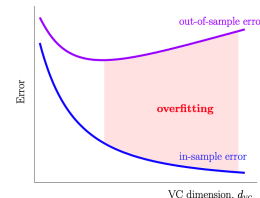
Hard constraints

$$H_2 = \{h \in H_{10} \text{ and } w_3 = w_4 = \dots = w_{10} = 0\}$$

- Instead of setting the weights to 0

$$H(C) = \left\{ h \in H_Q \text{ and } \sum_{q=0}^Q w_q^2 \leq C \right\}$$
$$= \{ h \in H_Q \text{ and } \vec{w}^T \vec{w} \leq C \}$$

- Observations
  - When  $C \rightarrow \infty$ ,  $H(C) = H_Q$
  - When  $C_1 \leq C_2$ ,  $H(C_1) \subseteq H(C_2)$  and therefore  $d_{vc}(H(C_1)) \leq d_{vc}(H(C_2))$
  - A smoother way to tune the complexity of hypothesis set



# Soft-Order Constraints

$$H(\mathcal{C}) = \{h \in H_Q \text{ and } \vec{w}^T \vec{w} \leq \mathcal{C}\}$$

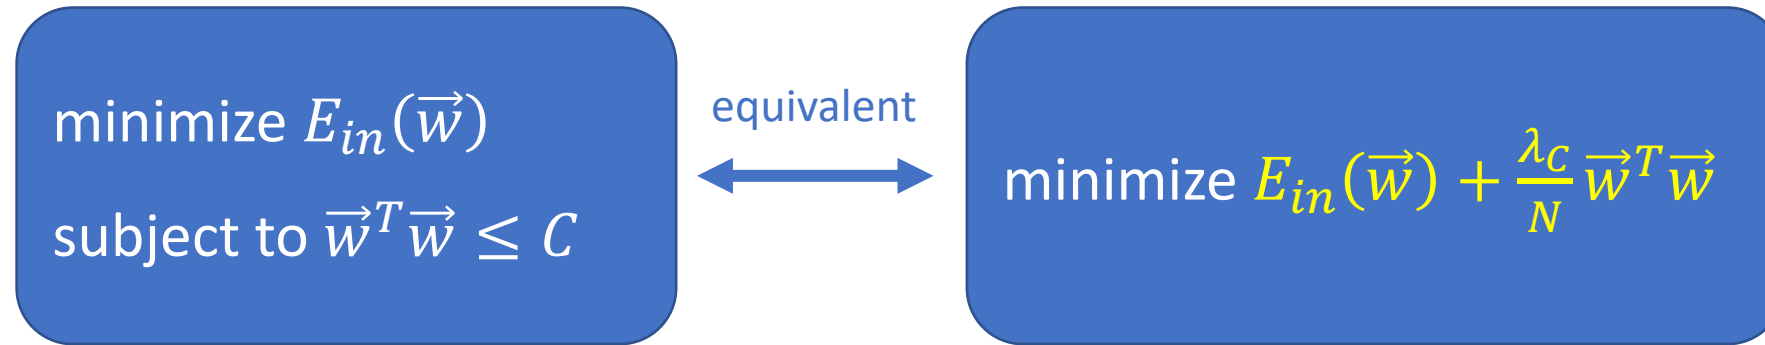
- Two main questions
  - How do we choose  $\mathcal{C}$ 
    - Model selection: The same question as selecting  $H$
    - The focus of the next lecture
  - How do we perform learning, i.e., find a  $g \in H(\mathcal{C})$  such that  $g \approx f$ 
    - Solve the following **constrained optimization** problem

minimize  $E_{in}(\vec{w})$

subject to  $\vec{w}^T \vec{w} \leq \mathcal{C}$

# Constrained to Unconstrained Optimization

- Constrained optimization  $\Leftrightarrow$  Unconstrained optimization



- Why the above is true?
  - Will talk about how to utilize Lagrangian relaxation to get this in the 2<sup>nd</sup> half of the semester
  - For now, let's think about it graphically

minimize  $E_{in}(\vec{w})$  subject to  $\vec{w}^T \vec{w} \leq C$

- Notations

- $\vec{w}_{lin}$ : the solution for  $\min E_{in}(\vec{w})$
- $\vec{w}_{reg}$ : the solution for  $\min E_{in}(\vec{w})$  subject to  $\vec{w}^T \vec{w} \leq C$

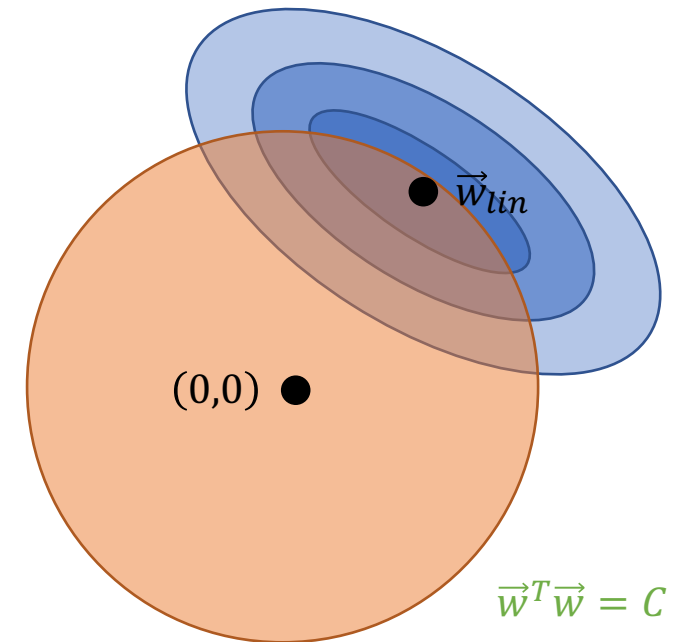
minimize  $E_{in}(\vec{w})$  subject to  $\vec{w}^T \vec{w} \leq C$

- Notations

- $\vec{w}_{lin}$ : the solution for  $\min E_{in}(\vec{w})$
- $\vec{w}_{reg}$ : the solution for  $\min E_{in}(\vec{w})$  subject to  $\vec{w}^T \vec{w} \leq C$

- When  $C$  is large enough, i.e.,  $\vec{w}_{lin}^T \vec{w}_{lin} \leq C$ 
  - $\vec{w}_{reg} = \vec{w}_{lin}$

- When  $C$  is not large enough, i.e.,  $\vec{w}_{lin}^T \vec{w}_{lin} > C$ 
  - $\vec{w}_{reg}^T \vec{w}_{reg} = C$



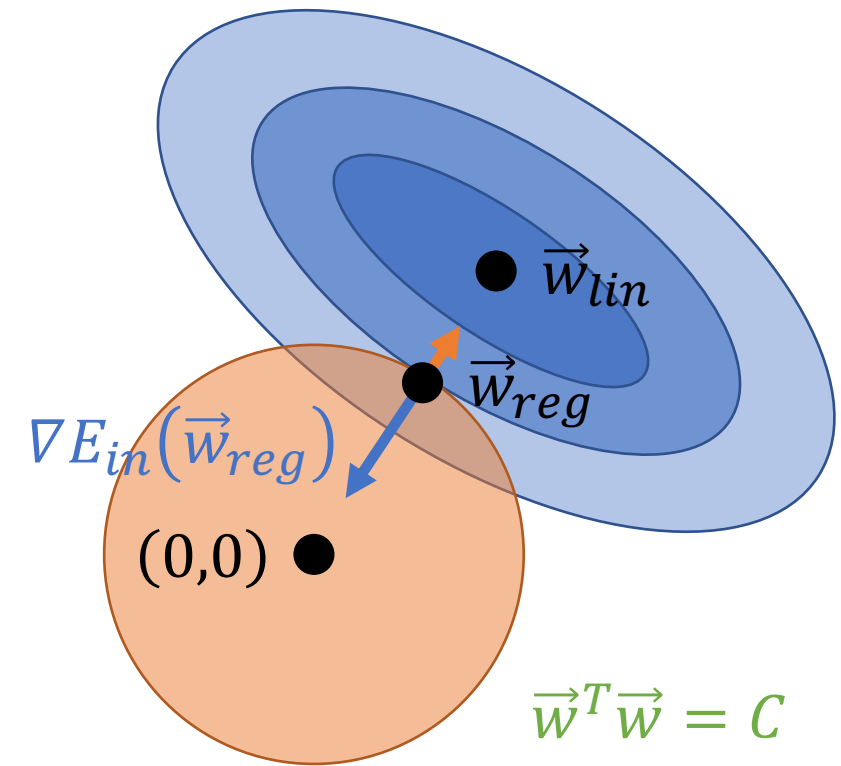
minimize  $E_{in}(\vec{w})$  subject to  $\vec{w}^T \vec{w} \leq C$

- When  $C$  is not large enough

- $\vec{w}_{lin}$ : the solution for  $\min E_{in}(\vec{w})$
- $\vec{w}_{reg}$ : the solution for  $\min E_{in}(\vec{w})$  subject to  $\vec{w}^T \vec{w} \leq C$

minimize  $E_{in}(\vec{w})$  subject to  $\vec{w}^T \vec{w} \leq C$

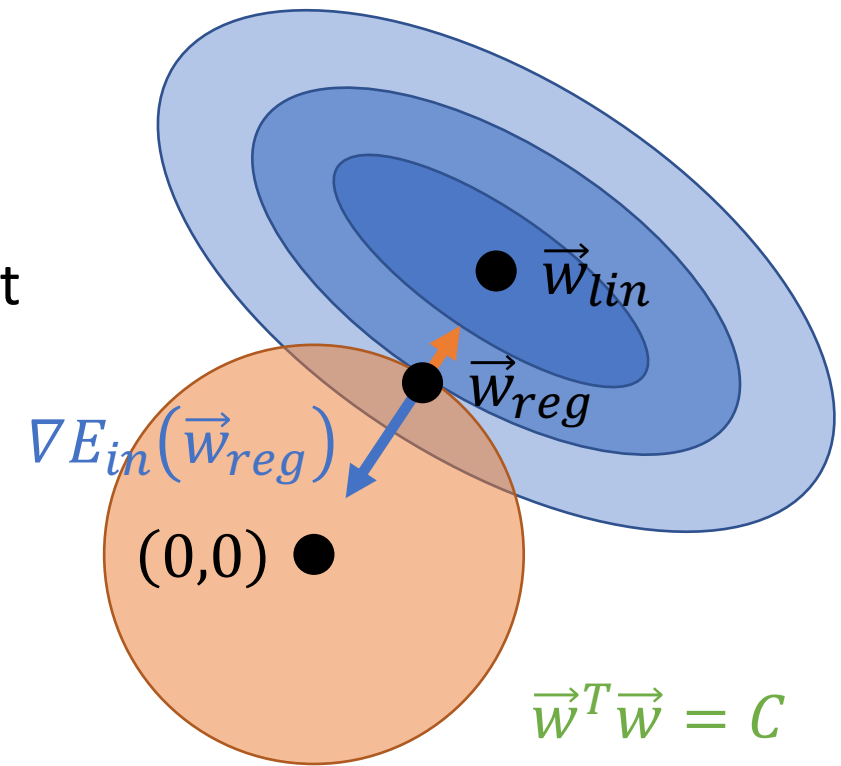
- When  $C$  is not large enough
  - Using graphical arguments
    - $\vec{w}_{reg} \propto -\nabla_{\vec{w}} E_{in}(\vec{w}_{reg})$





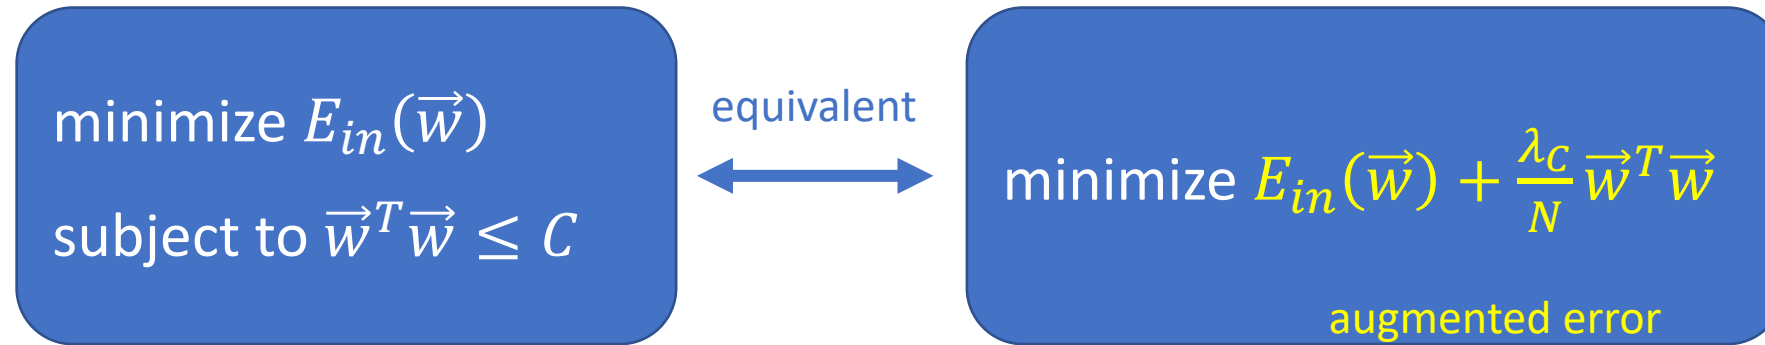
minimize  $E_{in}(\vec{w})$  subject to  $\vec{w}^T \vec{w} \leq C$

- When  $C$  is not large enough
  - Using graphical arguments
    - $\vec{w}_{reg} \propto -\nabla_{\vec{w}} E_{in}(\vec{w}_{reg})$
  - That is, we can find some constant  $\lambda_C \geq 0$  such that
    - $\nabla_{\vec{w}} E_{in}(\vec{w}_{reg}) = -\frac{2\lambda_C}{N} \vec{w}_{reg}$
- Therefore,
  - $\nabla_{\vec{w}} \left( E_{in}(\vec{w}_{reg}) + \frac{\lambda_C}{N} \vec{w}_{reg}^T \vec{w}_{reg} \right) = 0$
- This implies,  $\vec{w}_{reg}$  is the solution for
  - minimize  $E_{in}(\vec{w}) + \frac{\lambda_C}{N} \vec{w}^T \vec{w}$



# Constrained to Unconstrained Optimization

- Constrained optimization  $\Leftrightarrow$  Unconstrained optimization



- Interpretations of regularization
  - Constraining  $H$  (by adding constraints)
  - Adding penalty to complex hypothesis in augmented errors

# Augmented Error

$\vec{w}^T \vec{w}$ : weight decay

- Define augmented error

- $E_{aug}(\vec{w}) = E_{in}(\vec{w}) + \frac{\lambda_C}{N} \vec{w}^T \vec{w}$
- Algorithm: Find  $\vec{w}^* = \operatorname{argmin} E_{aug}(\vec{w})$

minimize  $E_{in}(\vec{w})$   
subject to  $\vec{w}^T \vec{w} \leq C$

minimize  $E_{in}(\vec{w}) + \frac{\lambda_C}{N} \vec{w}^T \vec{w}$

- A bit more discussion

- When  $C \rightarrow \infty$ ,  $\lambda_C = 0$
- Smaller  $C$  (stronger constraints)
  - $\Rightarrow$  larger  $\lambda_C$
  - $\Rightarrow$  smaller  $H$
  - $\Rightarrow$  stronger regularization
- Use  $\lambda_C$  to tune the level of regularization

Side note:

You will see people/us interchangeably use  $\lambda_C$  and  $\frac{\lambda_C}{N}$  to be the constant, depending on whether the dependency on  $N$  is emphasized.

# General Form of Regularization

$$E_{aug}(h, \lambda, \Omega) = E_{in}(\vec{w}) + \frac{\lambda}{N} \Omega(h)$$

- Key parameters
  - $\Omega$ : Regularizer
  - $\lambda$ : Amount of regularization
- Does the form look familiar: VC Theory
  - $E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{vc} \frac{\ln N}{N}}\right)$
- If we pick the right  $\Omega$ ,  $E_{aug}$  can be a better proxy for  $E_{out}$

# How to Pick the Right $\Omega$

- No definite answer, but generally
  - We like to pick  $\Omega$  that leads to “smoother” hypothesis
    - Overfitting is due to noise
    - Informally, noise is usually “high frequency”
  - We prefer  $\Omega$  that makes the optimization easier (e.g., convex/differentiable)
    - Similar to picking the error measure
  - We might have some other objective in mind
    - Ex: L-1 regularizer leads to weight vectors with more 0s
      - $E_{aug}(\vec{w}) = E_{in}(\vec{w}) + \lambda \|\vec{w}\|_1 = E_{in}(\vec{w}) + \lambda \sum_i |w_i|$
- What if we pick the wrong  $\Omega$  (Think about weight growth)
  - We might still fix it by picking the right  $\lambda$  – using validation in the next lecture

More Discussion on Regularization

# Why $\vec{w}^T \vec{w}$ is Called Weight Decay

- Run gradient descent on  $E_{aug}(\vec{w}) = E_{in}(\vec{w}) + \lambda_c \vec{w}^T \vec{w}$
- The update rule would be

$$\vec{w}(t+1) \leftarrow \vec{w}(t) - \eta \nabla_{\vec{w}} E_{aug}(\vec{w}(t))$$

$$\Rightarrow \vec{w}(t+1) \leftarrow (1 - 2\eta\lambda_c)\vec{w}(t) - \eta \nabla_{\vec{w}} E_{in}(\vec{w}(t))$$

Proving this is in your HW3.

We are **decaying** the weights first, then do the update

# Linear Regression with Weight Decay

- $E_{aug}(\vec{w}) = E_{in}(w) + \frac{\lambda_C}{N} \vec{w}^T \vec{w} = \frac{1}{N} \|X\vec{w} - \vec{y}\|^2 + \frac{\lambda_C}{N} \vec{w}^T \vec{w}$
- Solve  $\nabla_{\vec{w}} E_{aug}(\vec{w})|_{\vec{w}=\vec{w}_{reg}} = 0$ , we get
  - $\frac{2}{N} (X^T X \vec{w}_{reg} - X^T \vec{y} + \lambda_C \vec{w}_{reg}) = 0$
  - $(X^T X + \lambda_C I) \vec{w}_{reg} = X^T \vec{y}$
  - $\vec{w}_{reg} = (X^T X + \lambda_C I)^{-1} X^T \vec{y}$

Notation:  $I$  is an identity matrix: only the elements in the diagonals are 1, and all others are 0.

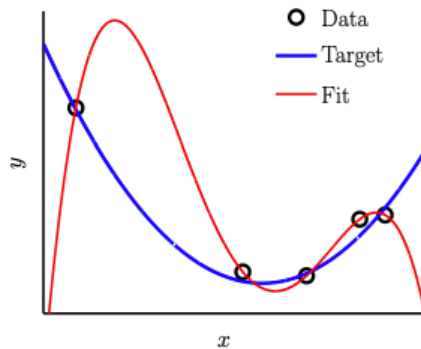
This is called “Ridge Regression” in statistics.



# Effect of Regularization (Different $\lambda$ )


- Minimizing  $E_{in}(\vec{w}) + \frac{\lambda}{N} \vec{w}^T \vec{w}$  with different  $\lambda$

$$\lambda = 0$$




Overfitting


$$\lambda = 0.0001$$



$$\lambda = 0.01$$

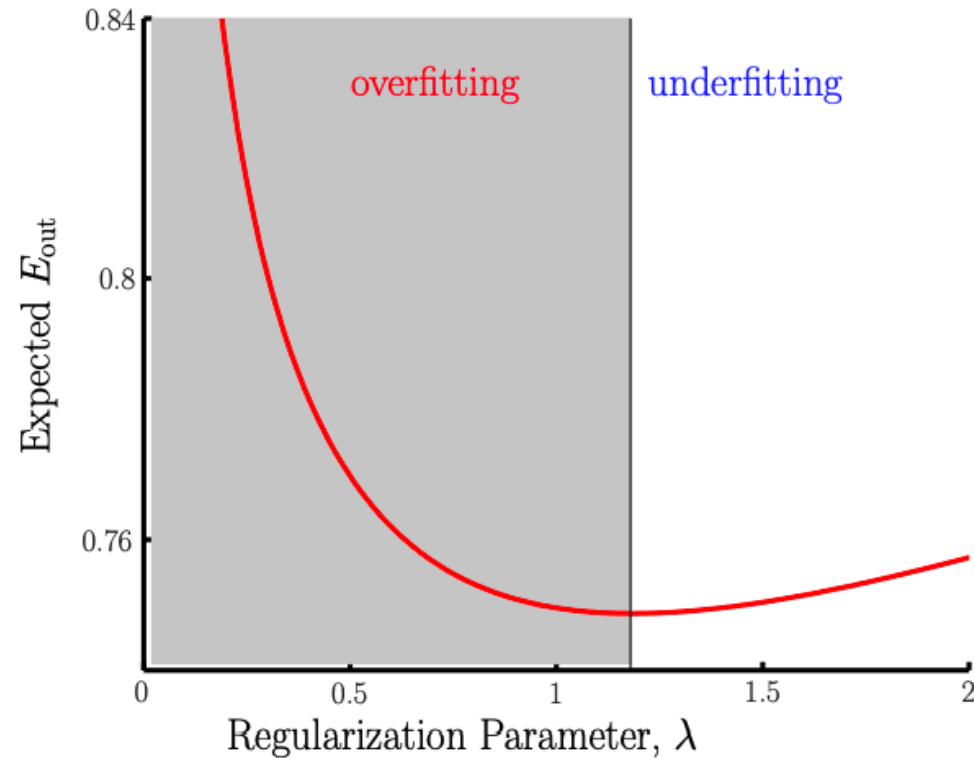


$$\lambda = 1$$


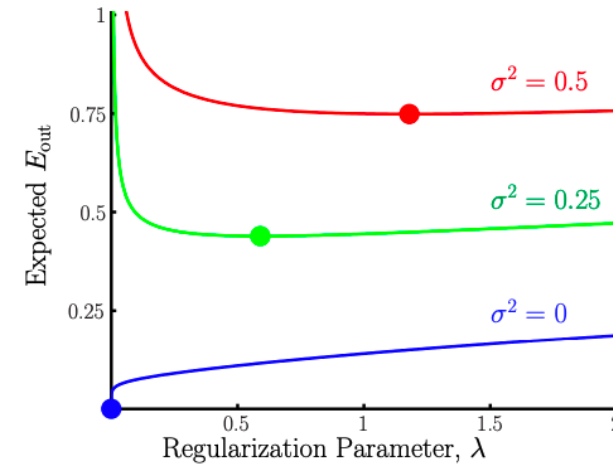


Underfitting

# Overfitting and Underfitting




Need to pick the right  $\lambda$ :  
Using validation: Focus of next lecture




# Variations on Weight Decay (Different $\Omega$ )

Uniform Weight Decay



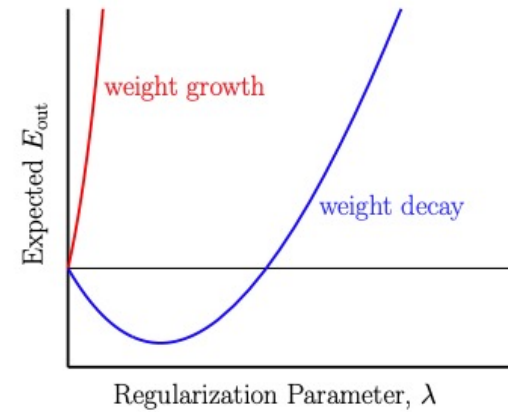
$$\sum_{q=0}^Q w_q^2$$

Low Order Fit



$$\sum_{q=0}^Q q w_q^2$$

Weight Growth!



$$\sum_{q=0}^Q \frac{1}{w_q^2}$$

# How to Pick the Right $\Omega$

- As discussed earlier
  - Intuition: pick  $\Omega$  that leads to “smoother” hypothesis
    - Overfitting is due to noise
    - Informally, noise is generally “high frequency”
  - Computation: prefer  $\Omega$  that makes the optimization easier (e.g., convex/differentiable)
    - Similar to picking the error measure
  - We might have some other objective in mind
    - Ex: L-1 regularizer leads to weight vectors with more 0s
- What if we pick the **wrong  $\Omega$**  (weight growth)
  - We might still fix it by picking the right  $\lambda$  – using **validation**

# Summarizing Regularization

- Regularization is **everywhere** in machine learning
- Two main ways of thinking about regularization
  - **Constrain  $H$**  to make overfitting less likely to happen
    - Will discuss more regularization methods in the 2nd half of the semester
    - Pruning for decision trees, early stopping / dropout for neural networks, etc
  - Define **augmented error**  $E_{aug}$  to better approximate  $E_{out}$ 
    - $E_{aug}(h, \lambda, \Omega) = E_{in}(\vec{w}) + \frac{\lambda}{N} \Omega(h)$
- We show the **equivalence** of the two for weight decay
  - The conceptual equivalence is general with Lagrangian relaxation (will cover later in the semester)