

CSE 417T

Introduction to Machine Learning

Lecture 3

Instructor: Chien-Ju (CJ) Ho

Logistics

- Course website and Piazza
 - Website: <http://chienjuho.com/courses/cse417t/>
 - Piazza: <http://piazza.com/wustl/spring2020/cse417t>
 - Make sure you follow both regularly
- Office hours
 - Will be announced later this week
 - Will start next week

Logistics

- Homework 1
 - Will be announced tomorrow or before lecture on Thursday
 - Expected due: Feb 19 (Friday)
 - Mixture of math questions and programming questions (implement PLA)
 - Programming language: Python
 - We won't teach you how to program python
 - Basic sessions? (e.g., for environment setup, etc)
- Exam and Grades
 - Two exams (one in the middle of semester, one in the last day)
 - What to expect for the final grades

Recap

UNKNOWN TARGET FUNCTION

$$f: \mathcal{X} \mapsto \mathcal{Y}$$

(ideal credit approval formula)

$$y_n = f(\mathbf{x}_n)$$

TRAINING EXAMPLES

$$(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$$

(historical records of credit customers)

Given by the learning problem

**LEARNING
ALGORITHM**

\mathcal{A}

**FINAL
HYPOTHESIS**

$$g \approx f$$

(learned credit approval formula)

Goal of learning

HYPOTHESIS SET

\mathcal{H}

(set of candidate formulas)

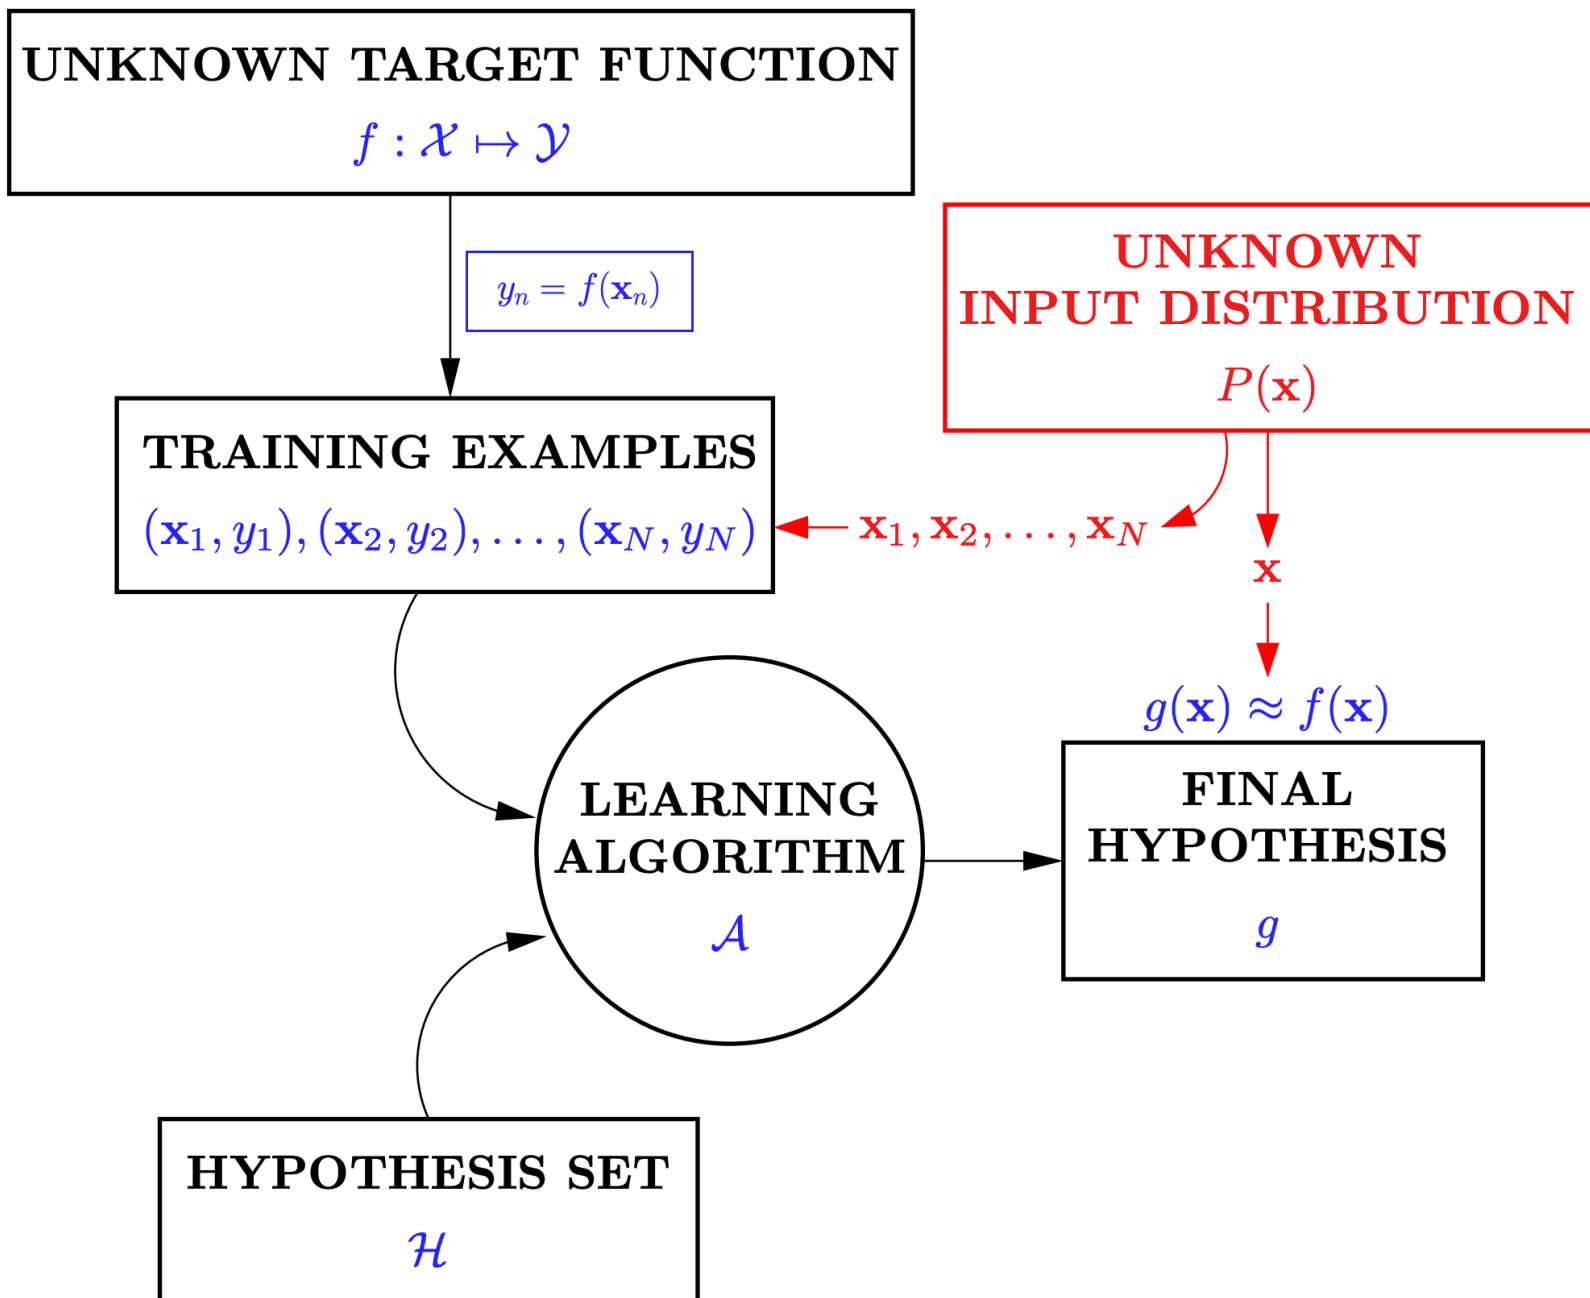
learning model
(example:
H: Perceptron
A: PLA)

Goal of Learning: Generalization


- Given **training data**, find $g \approx f$ on the **unseen testing data**.
- This goal is generally impossible without assumptions.

Key assumption of ML

Training data points and **testing** data points are **i.i.d.**
drawn from the same (unknown) distribution



A Thought Experiment about Probability



What can we say about μ from ν ?

Law of large numbers

- When $N \rightarrow \infty$, $\nu \rightarrow \mu$


Hoeffding's Inequality

- $\Pr[|\mu - \nu| > \epsilon] \leq 2e^{-2\epsilon^2 N}$ for any $\epsilon > 0$

Connection to Learning

- Let each marble represent a point \vec{x} , drawn from unknown $P(\vec{x})$
 - Dataset $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$
 - Recall that $y_n = f(\vec{x}_n)$ (will discuss noisy target function f later in the semester)

- “Fix” a hypothesis h
 - For each marble \vec{x} , color it as below
 - If $h(\vec{x}) = f(\vec{x})$, color it as green marble [h is correct on \vec{x}]
 - If $h(\vec{x}) \neq f(\vec{x})$, color it as red marble [h is wrong on \vec{x}]



- With the above coloring

$$\mu = \Pr_{\vec{x} \sim P(\vec{x})} [h(\vec{x}) \neq f(\vec{x})]$$

$$\stackrel{\text{def}}{=} E_{out}(h) \quad \text{[Out-of-sample error of } h]$$

$$\nu = \frac{1}{N} \sum_{n=1}^N \mathbb{I}[h(\vec{x}_n) \neq f(\vec{x}_n)]$$

$$\stackrel{\text{def}}{=} E_{in}(h) \quad \text{[in-sample error of } h]$$

Connection to Learning

- $E_{out}(h)$: What we really want to know but unknown to us
- $E_{in}(h)$: What we can calculate from dataset
- Fixed a h , What can we say about $E_{out}(h)$ from $E_{in}(h)$?

Hoeffding's Inequality

$$\Pr[|E_{out}(h) - E_{in}(h)| > \epsilon] \leq 2e^{-2\epsilon^2 N} \quad \text{for any } \epsilon > 0$$

- This is verification, not learning!

Verification vs. Learning

- Verification

- I have a hypothesis h .
- I know $E_{in}(h)$, i.e., how well h performs in my dataset.
- I can infer what $E_{out}(h)$ (how well h will perform for unseen data) might be.

- Learning

- Given a dataset D and hypothesis set H .
- Apply some learning algorithm, that outputs a $g \in H$.
- Know $E_{in}(g)$.
- Want to infer $E_{out}(g)$


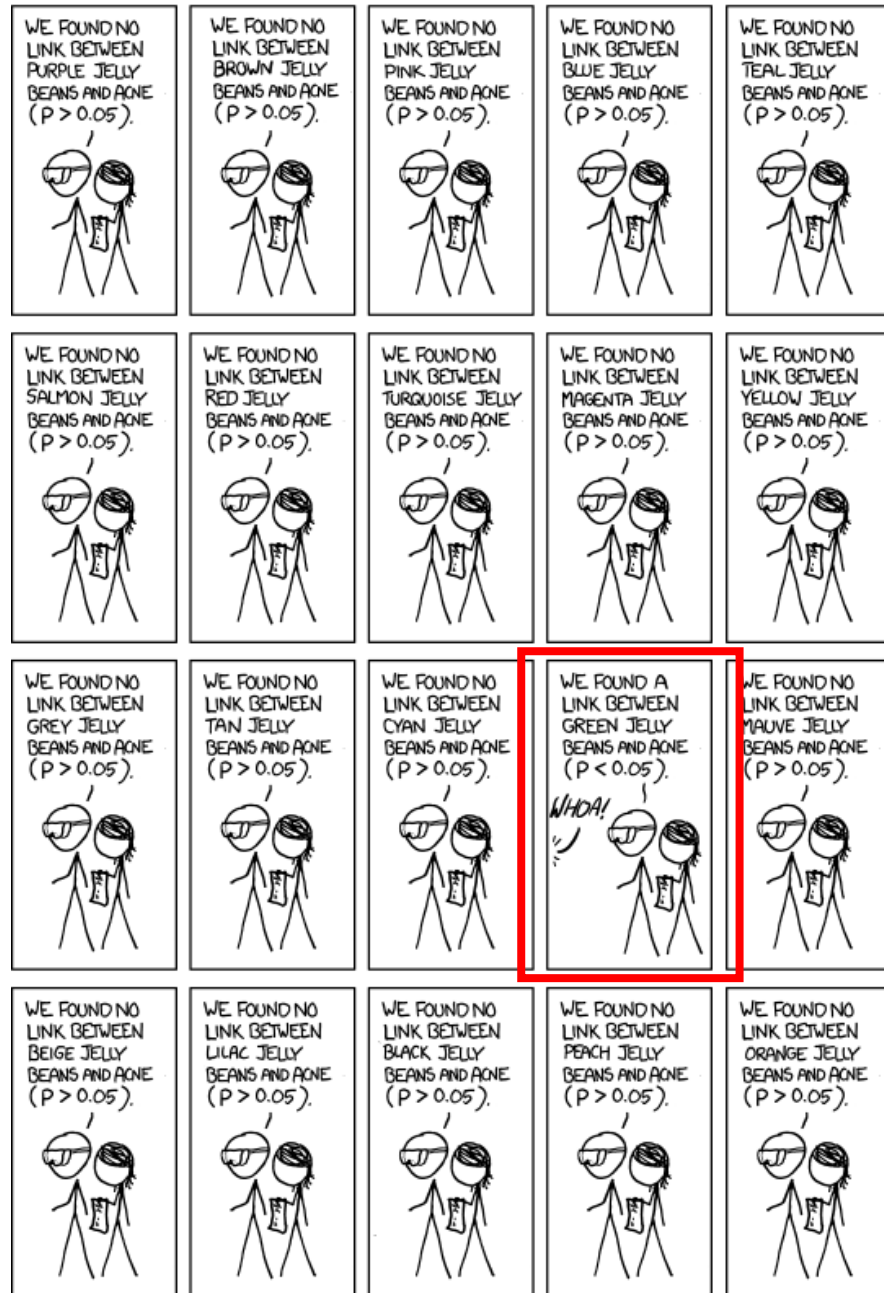
Connection to “Real” Learning

- Given a **finite** hypothesis set $H = \{h_1, \dots, h_M\}$
 - Will discuss the infinite case in the next few lectures.
- Apply some learning algorithm on D , output a $g \in H$
 - For example, choosing the hypothesis that minimizes in-sample error
 - $g = \operatorname{argmin}_{h \in H} E_{in}(h)$
- Can we apply Hoeffding’s inequality and claim
$$\Pr[|E_{out}(g) - E_{in}(g)| > \epsilon] \leq 2e^{-2\epsilon^2 N} \quad \text{for any } \epsilon > 0$$
- **No!**

Today's Lecture

The notes are not intended to be comprehensive.
Let me know if you spot errors.





Another Analogy

- If you toss a fair coin 10 times, the prob that you get heads 10 times is

$$2^{-10} = \frac{1}{1024}$$

- If you toss 1000 fair coins 10 times each, the probability that at least one coin comes up heads 10 times is

$$1 - \left(\frac{1023}{1024}\right)^{1000} \approx 62.36\%$$

- If each hypothesis is doing random guessing (i.e., tossing a fair coin), if we have 1000 hypothesis with 10 data points, more than 60% chance there will be at least one hypothesis with **zero in-sample error**
 - But that hypothesis is still random guessing and has 50% out-of-sample error

One More Analogy

- Three fair coins, numbered by 1, 2, 3. Flip each 10 times
- Question: (choosing from >5 , $=5$, or <5)

Ans: = 5 • For coin 1, what's the expected number of heads among 10 flips?

Ans: = 5 • Randomly choose a coin, what's the expected number of heads for this coin?

Ans: < 5 • After observing the realized flips, choosing the coin with the smallest number of heads, what is the expected number of heads for the coin?

Ans: = 5 • Without observing the flips, choose the coin anyway you like, what is the expected number of heads of the 10 flips for this coin?

- You will simulate this process (with 1,000 coins) in HW1.

One More Analogy

- Connects to learning
 - Coin -> Hypothesis
 - Coin flips -> Performance of hypothesis in training data D
- Choosing the hypothesis “before” or “after” looking at the data (knowing the realization of the data drawing) makes a very big difference!

What Can We Do?

Connection to “Real” Learning

- Given a **finite** hypothesis set $H = \{h_1, \dots, h_M\}$
- Apply some learning algorithm on D , output a $g \in H$
- Question: What can we say about $E_{out}(g)$ from $E_{in}(g)$?

Derivations

- Define “bad event of h ” $B(h)$ as $|E_{out}(h) - E_{in}(h)| > \epsilon$
 - Informally, you can interpret “bad event of h ” as the event that we draw a “unrepresentative dataset D ” that makes the in-sample errors of h to be far away from out-of-sample error of h

For each fixed $h \in H$, we have $\Pr[B(h)] \leq 2e^{-2\epsilon^2 N}$

- Recall g is selected from H (it could be any $h \in H$)
- What can we say about $\Pr[B(g)]$?

Derivations

- Define “bad event of h ” $B(h)$ as $|E_{out}(h) - E_{in}(h)| > \epsilon$
 - Informally, you can interpret “bad event of h ” as the event that we draw a “unrepresentative dataset D ” that makes the in-sample errors of h to be far away from out-of-sample error of h

For each fixed $h \in H$, we have $\Pr[B(h)] \leq 2e^{-2\epsilon^2 N}$

- Recall g is selected from H (it could be any $h \in H$)
- What can we say about $\Pr[B(g)]$?

$$\begin{aligned}\Pr[B(g)] &\leq \Pr[B(h_1) \text{ or } B(h_2) \text{ or } \dots \text{ or } B(h_M)] \\ &\leq \Pr[B(h_1)] + \Pr[B(h_2)] + \dots + \Pr[B(h_M)] \\ &\leq M 2e^{-2\epsilon^2 N}\end{aligned}$$

Connection to “Real” Learning

- Given a **finite** hypothesis set $H = \{h_1, \dots, h_M\}$
- Apply some learning algorithm on D , output a $g \in H$
- Question: What can we say about $E_{out}(g)$ from $E_{in}(g)$?

$$\Pr[|E_{out}(g) - E_{in}(g)| > \epsilon] \leq 2\mathbf{M}e^{-2\epsilon^2 N} \quad \text{for any } \epsilon > 0$$

- \mathbf{M} can be considered as a proxy of the “complexity” of the hypothesis set
 - Will talk about what happens when $\mathbf{M} \rightarrow \infty$ in the next few lectures

Interpreting $\Pr[|E_{out}(g) - E_{in}(g)| > \epsilon] \leq 2Me^{-2\epsilon^2 N}$

- Playing around with the math
 - Define $\delta = \Pr[|E_{out}(g) - E_{in}(g)| > \epsilon]$
 - We have $\delta \leq 2Me^{-2\epsilon^2 N} \Rightarrow \epsilon \leq \sqrt{\frac{1}{2N} \ln \frac{2M}{\delta}}$
- This means, with probability at least $1 - \delta$
 - $E_{out}(g) \leq E_{in}(g) + \epsilon \leq E_{in}(g) + \sqrt{\frac{1}{2N} \ln \frac{2M}{\delta}}$

More Discussion

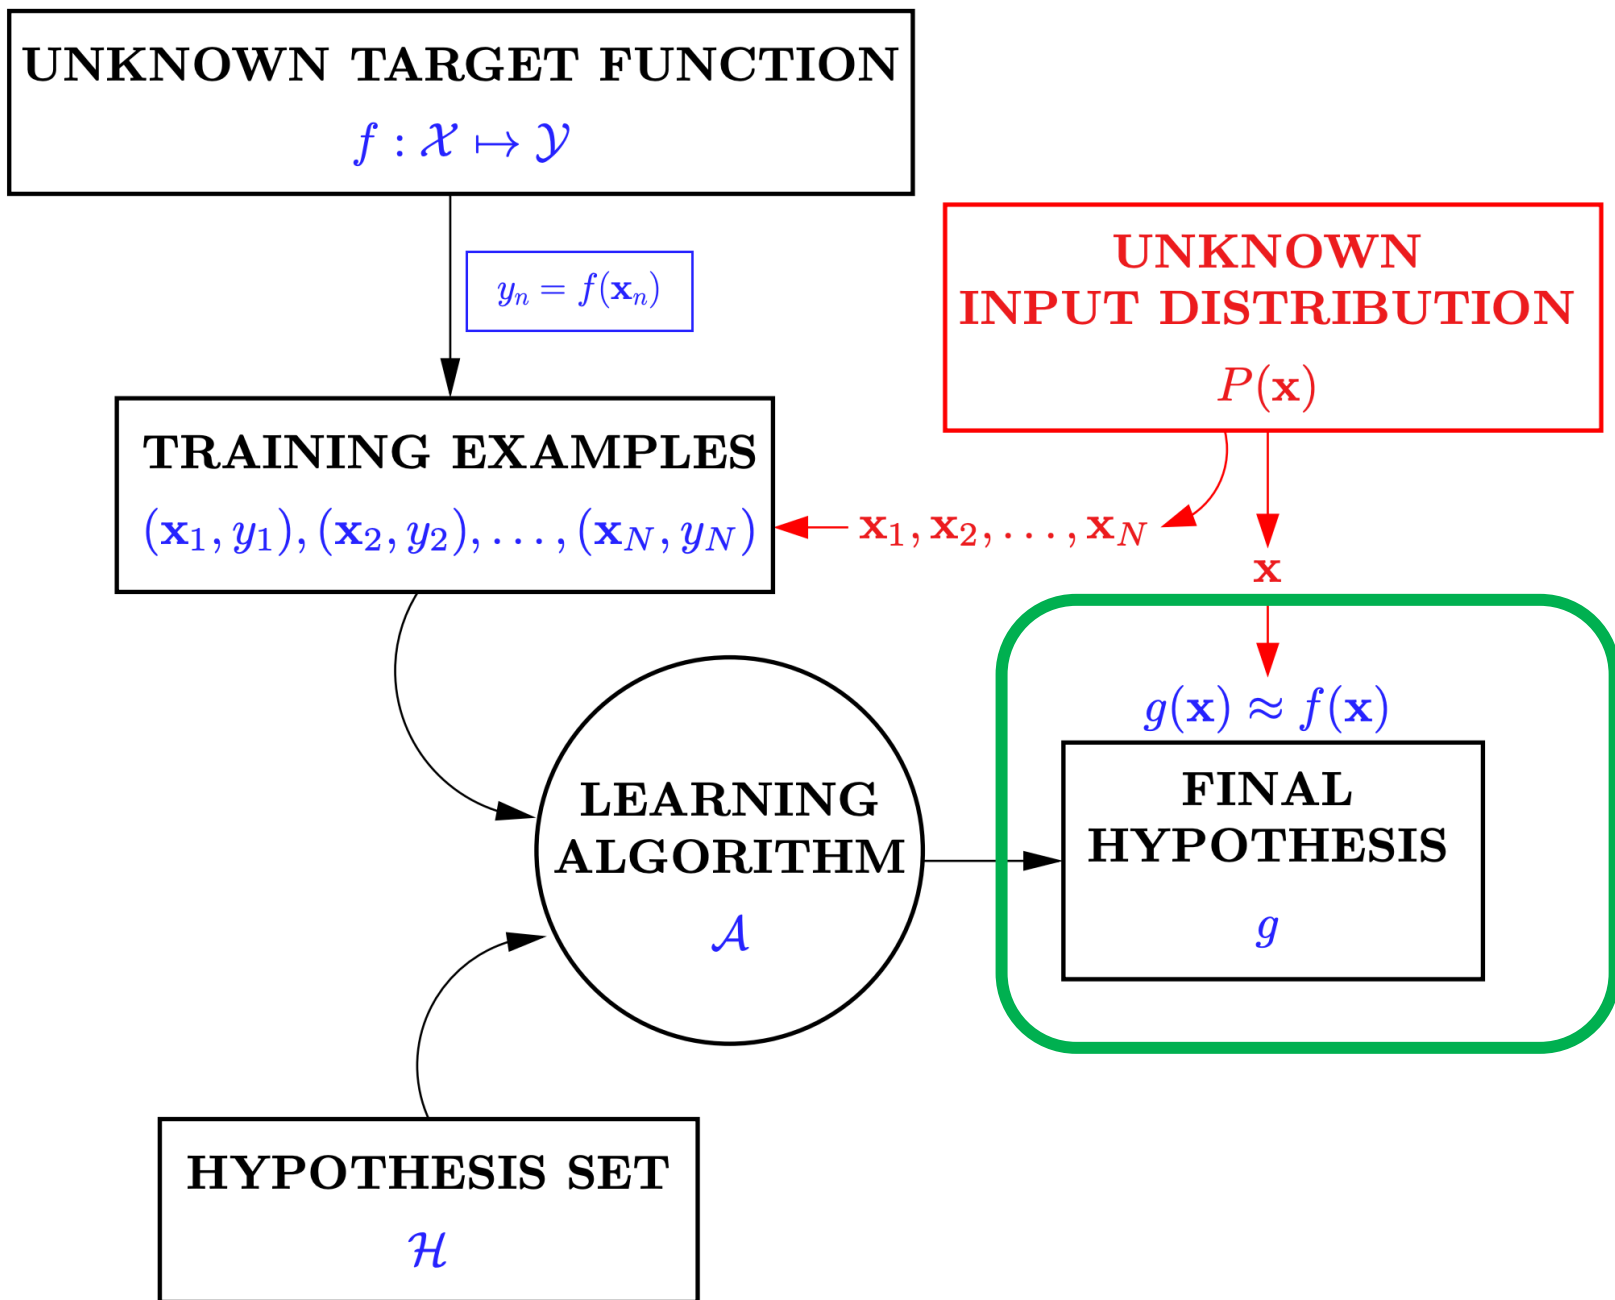
- With probability at least $1 - \delta$

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{1}{2N} \ln \frac{2M}{\delta}}$$

Consider M as a proxy measure on the “complexity” of H

- Our ultimate goal is to have a small $E_{out}(g)$
 - There is a tradeoff of choosing M (what “learning model” to use)
 - Increase M -> Smaller $E_{in}(g)$ (more hypothesis to “fit” the training data)
 - Increase M -> Larger ϵ
 - It also depends on N , the number of data points you have
 - A small number of data points => use simple models (e.g., linear models)
 - Complex models (e.g., deep learning) work when you have a lot of data

Revisit the Learning Problem



Goal: $g \approx f$

- A general approach:
 - Define an error function $E(h, f)$ that quantify how far away g is to f
 - Choose the one with the smallest error (empirical risk minimization)
 - For example: $g = \operatorname{argmin}_{h \in \mathcal{H}} E(h, f)$
- E is usually defined in terms of a pointwise error function $e(h(\vec{x}), f(\vec{x}))$
 - Binary error (classification): $e(h(\vec{x}), f(\vec{x})) = \mathbb{I}[h(\vec{x}_n) \neq f(\vec{x}_n)]$ (What we have discussed so far)
 - Squared error (regression): $e(h(\vec{x}), f(\vec{x})) = (f(\vec{x}) - h(\vec{x}))^2$
- In-sample and out-of-sample errors
 - $E_{in}(h) = \frac{1}{N} \sum_{n=1}^N e(h(\vec{x}_n), f(\vec{x}_n))$
 - $E_{out}(h) = \mathbb{E}_{\vec{x}}[e(h(\vec{x}), f(\vec{x}))]$

The discussion on the Hoeffding's inequality applies for general (bounded) error functions.

How to choose the error function?

- Consideration 1: Properties of application problems
- Example: Fingerprint recognition
 - Input: fingerprints
 - Outputs: whether the person is authorized

		$f(\vec{x})$	
		+1	-1
$h(\vec{x})$	+1	No error	False positive
	-1	False negative	No error

- Errors assigned to false negative/positive differ depending on applications
 - Supermarket coupons vs FBI
 - False positive is a big issue for FBI but probably fine for supermarket coupons

How to choose the error function?

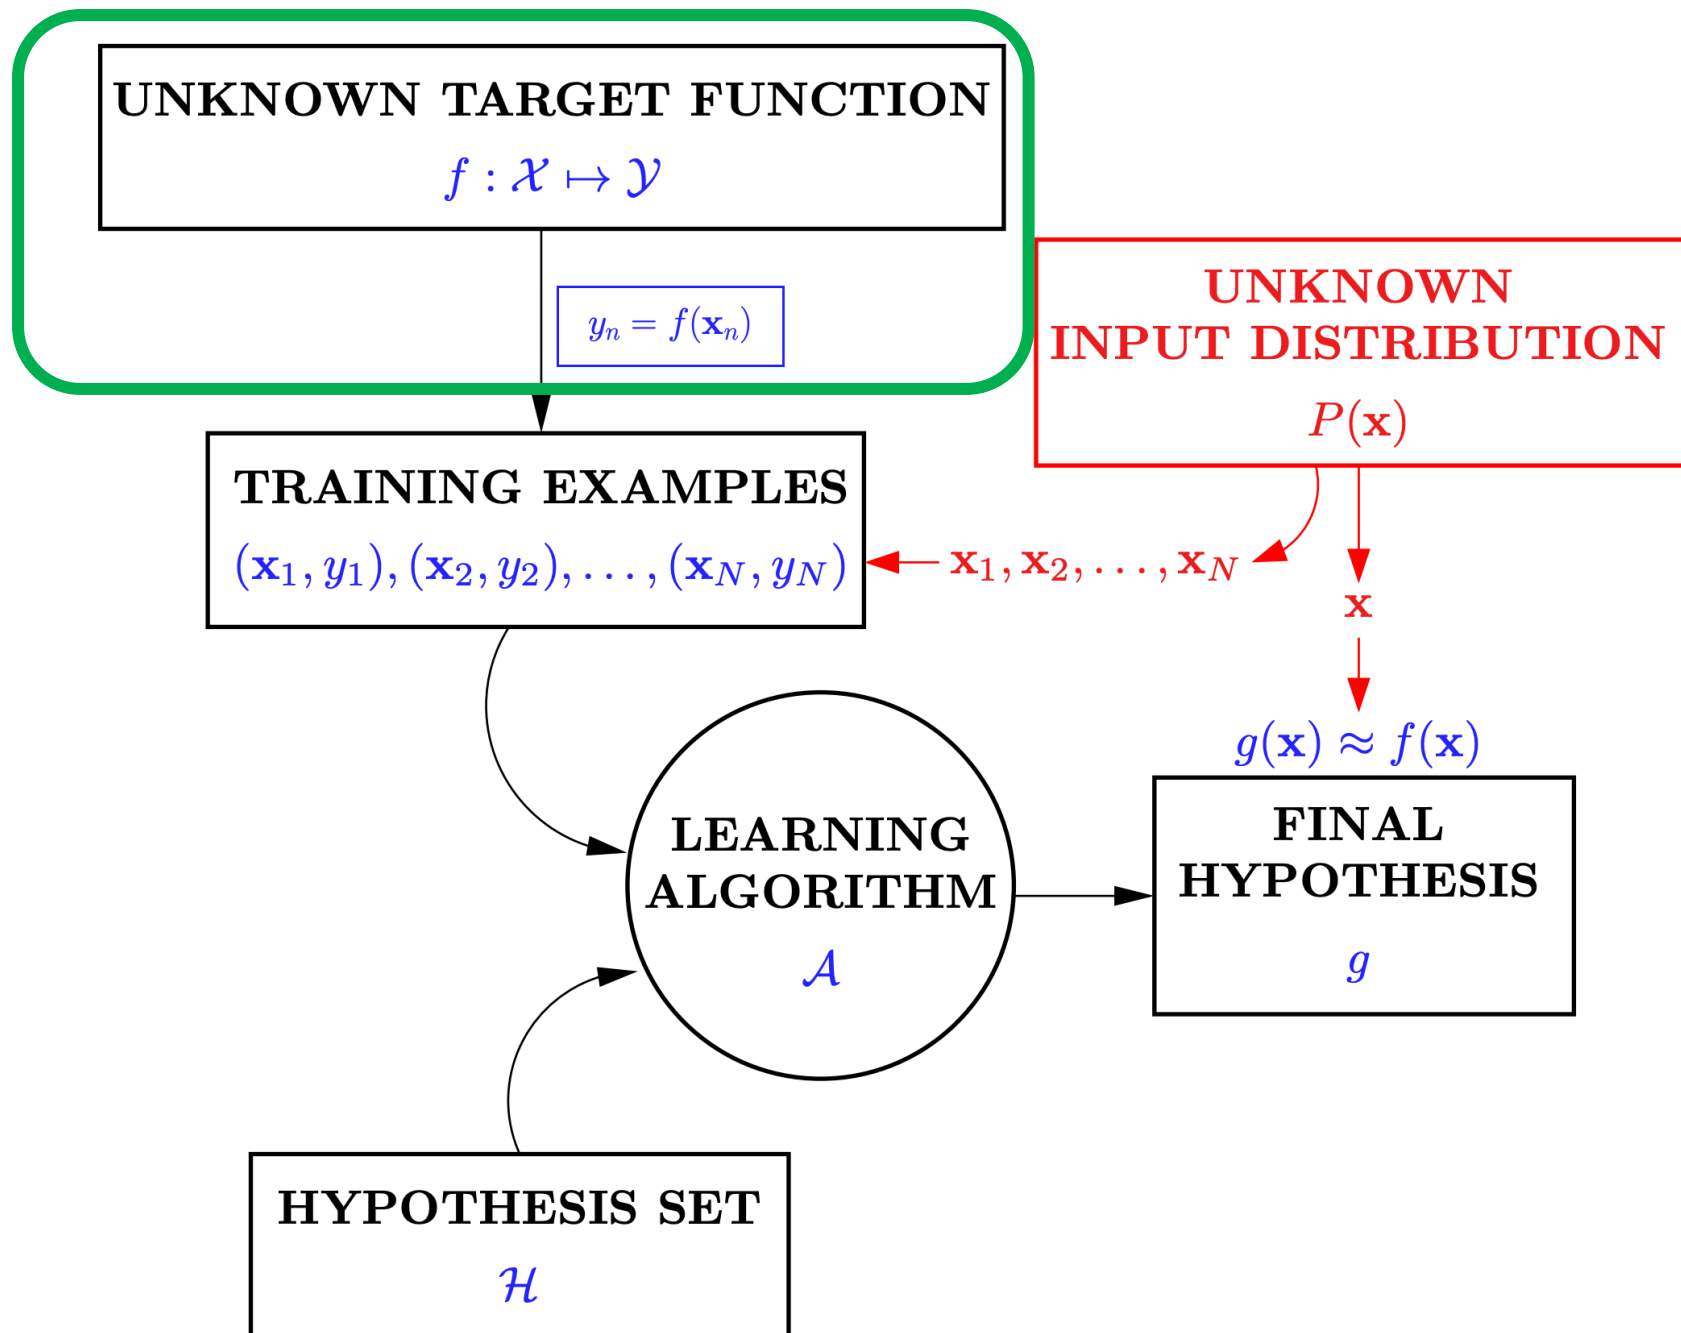
- Consideration 1: Properties of application problems
- Consideration 2: Computation
 - ML Algorithm is essentially doing **optimization** (finding g with smallest error)

$$g = \operatorname{argmin}_{h \in \mathcal{H}} E(h, f)$$

- Choosing the error that is “easier” to optimize

How to choose the error function?

- Consideration 1: Properties of application problems
- Consideration 2: Computation
- Specifying the error function is part of setting up the learning problem
 - It impacts what you eventually learn



Noisy Target

- What if there doesn't exist f such that $y = f(\vec{x})$?
 - f is stochastic instead of deterministic
- Common approach
 - Instead of a target function, define a target **distribution**
 - Instead of $y = f(\vec{x})$, y is drawn from a conditional distribution $P(y|\vec{x})$
 - $y = f(\vec{x}) + \epsilon$ where ϵ is zero-mean noise

The discussion on the Hoeffding's inequality applies for noisy targets.

UNKNOWN TARGET DISTRIBUTION
(target function f plus noise)
 $P(y | \mathbf{x})$

$$y_n \sim P(y | \mathbf{x}_n)$$

**UNKNOWN
INPUT DISTRIBUTION**
 $P(\mathbf{x})$

TRAINING EXAMPLES
 $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$

$\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$

\mathbf{x}

**ERROR
MEASURE**

$$g(\mathbf{x}) \approx f(\mathbf{x})$$

**LEARNING
ALGORITHM**
 \mathcal{A}

**FINAL
HYPOTHESIS**
 g

HYPOTHESIS SET
 \mathcal{H}

