

CSE 417T

# Introduction to Machine Learning

Lecture 9

Instructor: Chien-Ju (CJ) Ho

# Logistics

- Homework 2 is due on **Mar 8, Monday**
  - Implement gradient descent for logistic regression
  - Several math questions
- Return of Homework
  - We plan to return each homework around 2 weeks after the deadline
  - Regrade requests
    - You will have up to 7 days to submit regrade requests after homework return.
    - We might check the entire homework for each request, so the grades might go down as well if we find new mistakes
- Exam 1: **Mar 23 (Tuesday)**

Recap

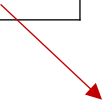
# Linear Models

This is why it's called linear models



- $H$  contains hypothesis  $h(\vec{x})$  as **some function of  $\vec{w}^T \vec{x}$**


	Domain	Model	Credit Card Example
Linear Classification	$y \in \{-1, +1\}$	$H = \{h(\vec{x}) = \text{sign}(\vec{w}^T \vec{x})\}$	Approve or not
Linear Regression	$y \in \mathbb{R}$	$H = \{h(\vec{x}) = \vec{w}^T \vec{x}\}$	Credit line
Logistic Regression	$y \in [0,1]$	$H = \{h(\vec{x}) = \theta(\vec{w}^T \vec{x})\}$	Prob. of default


$$\theta(s) = \frac{e^s}{1 + e^s}$$

- Algorithm:
  - Focus on  $g = \operatorname{argmin}_{h \in H} E_{in}(h)$

# Logistic Regression

- Predict a probability
  - Interpreting  $h(\vec{x}) \in [0,1]$  as the prob for  $y = +1$  given  $\vec{x}$
- Hypothesis set  $H = \{h(\vec{x}) = \theta(\vec{w}^T \vec{x})\}$ 
  - $\theta(s) = \frac{e^s}{1+e^s} = \frac{1}{1+e^{-s}}$
- Algorithm
  - Find  $g = \operatorname{argmin}_{h \in H} E_{in}(h)$
- Two key questions
  - How to define  $E_{in}(h)$ ?
  - How to perform the optimization (minimizing  $E_{in}$ )?



Define  $E_{in}(\vec{w})$ : Cross-Entropy Error

$$E_{in}(\vec{w}) = \frac{1}{N} \sum_{n=1}^N \ln(1 + e^{-y_n \vec{w}^T \vec{x}_n})$$

- Minimizing cross entropy error is the same as maximizing likelihood
- Likelihood:  $\Pr(D|\vec{w})$ 
  - $\vec{w}^* = \operatorname{argmax}_{\vec{w}} \Pr(D|\vec{w})$  (maximizing likelihood)  
   $= \operatorname{argmin}_{\vec{w}} E_{in}(\vec{w})$  (minimizing cross-entropy error)

# Min Cross-Entropy Error $\Leftrightarrow$ Max Likelihood

- $\vec{w}^* = \operatorname{argmax}_{\vec{w}} \Pr(D|\vec{w})$   
 $= \operatorname{argmax}_{\vec{w}} \prod_{n=1}^N \Pr(y_n|\vec{x}_n, \vec{w})$   
 $= \operatorname{argmax}_{\vec{w}} \prod_{n=1}^N \theta(y_n \vec{w}^T \vec{x}_n)$   
 $= \operatorname{argmax}_{\vec{w}} \ln(\prod_{n=1}^N \theta(y_n \vec{w}^T \vec{x}_n))$   
 $= \operatorname{argmax}_{\vec{w}} \sum_{n=1}^N \ln(\theta(y_n \vec{w}^T \vec{x}_n))$   
 $= \operatorname{argmin}_{\vec{w}} - \sum_{n=1}^N \ln(\theta(y_n \vec{w}^T \vec{x}_n))$   
 $= \operatorname{argmin}_{\vec{w}} \sum_{n=1}^N \ln \frac{1}{\theta(y_n \vec{w}^T \vec{x}_n)}$   
 $= \operatorname{argmin}_{\vec{w}} \sum_{n=1}^N \ln(1 + e^{-y_n \vec{w}^T \vec{x}_n})$   
 $= \operatorname{argmin}_{\vec{w}} \frac{1}{N} \sum_{n=1}^N \ln(1 + e^{-y_n \vec{w}^T \vec{x}_n})$

1. data independence assumption
2.  $1 - \theta(s) = \theta(-s)$

$$\begin{aligned} \operatorname{argmax} A(x)B(x) \\ = \operatorname{argmax} \ln A(x) + \ln B(x) \end{aligned}$$

$$E_{in}(\vec{w}) = \frac{1}{N} \sum_{n=1}^N \ln(1 + e^{-y_n \vec{w}^T \vec{x}_n})$$

# Optimizing $E_{in}(\vec{w})$ : Gradient Descent

An iterative method:  $\vec{w}(t+1) \leftarrow \vec{w}(t) + \eta_t \vec{v}_t$


- $\vec{v}_t$ : a unit vector, determining the direction of the update
- $\eta_t$ : a scalar, determining how much to update

- How to choose  $\vec{v}_t$


- Move towards the “steepest” direction
- Approaching the minimum faster
- Taylor approximation:
  - $E_{in}(\vec{w}(t+1)) - E_{in}(\vec{w}(t)) \approx \eta_t \nabla_{\vec{w}} E_{in}(\vec{w}(t))^T \vec{v}_t$
- Choose  $\vec{v}_t$  to be the opposite direction of  $\nabla_{\vec{w}} E_{in}$ 
  - $\vec{v}_t = \frac{-\nabla_{\vec{w}} E_{in}(\vec{w}(t))}{\|\nabla_{\vec{w}} E_{in}(\vec{w}(t))\|}$

- How to choose  $\eta_t$


$\eta$  too small



$\eta$  too large



variable  $\eta_t$  – just right



- $\eta_t = \eta \|\nabla_{\vec{w}} E_{in}(\vec{w}(t))\|$



# Gradient Descent (GD) for Logistic Regression

- Initialize  $\vec{w}(0)$
- For  $t = 0, \dots$ 
  - Compute gradient  $\nabla_{\vec{w}} E_{in}(\vec{w}(t)) = -\frac{1}{N} \sum_{n=1}^N \frac{y_n \vec{x}_n}{1 + e^{y_n \vec{w}(t)^T \vec{x}_n}}$
  - $\vec{w}(t + 1) \leftarrow \vec{w}(t) - \eta \nabla_{\vec{w}} E_{in}(\vec{w}(t))$  [Take gradient, then descent]
  - Terminate if the **stop conditions** are met
- Return the final weights

$\eta$ : learning rate  
A parameter the learner can choose

We focus on fixed learning rate GD  
There are other variants

# Gradient Descent (GD) for Logistic Regression

- Initialization
  - In HW2, you are asked to initialize  $\vec{w}(0)$  to  $\vec{0}$
  - In practice, random initialization is a good idea and a common approach
- Stop conditions (a mix of the following criteria)
  - When the **number of iteration** exceeds the pre-set threshold
  - When the **improvement on  $E_{in}$**  (e.g., check  $\nabla_{\vec{w}} E_{in}$ ) is too small
  - When  **$E_{in}$  is small** enough
  - (We use the first two in HW2)

# Using Logistic Regression for Classification

- Let  $\vec{w}^*$  or  $g$  be the learned logistic regression model, how can we make classification predictions using  $\vec{w}^*$ ?
- Set a cutoff probability  $C\%$  (e.g., 50%).
  - Classify +1 if  $g(\vec{x}) = \theta(\vec{w}^{*T} \vec{x}) > C\%$
  - Classify -1 if  $g(\vec{x}) = \theta(\vec{w}^{*T} \vec{x}) < C\%$
- When  $C$  is 50 (a common choice)
  - $\theta(\vec{w}^{*T} \vec{x}) > 50\% \Rightarrow \vec{w}^{*T} \vec{x} > 0$
  - Equivalent to using  $\vec{w}^*$  as the linear classification hypothesis, i.e.,  $g(\vec{x}) = \text{sign}(\vec{w}^{*T} \vec{x})$

# Today's Lecture

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.  
Let me know if you spot errors.

# More on Cross Entropy [This Page is Safe to Skip]

- Cross entropy of  $q$  related to  $p$ :  $H(p, q) = \sum_{i=1}^n p(x_i) \log \frac{1}{q(x_i)}$ 
  - Distance measure between two distributions
  - Fix  $p$ ,  $H(p, q)$  is minimized when  $q = p$  [Solve for  $\nabla_q H(p, q) = 0$ ]
- Cross-entropy error
  - $$E_{in}(\vec{w}) = \frac{1}{N} \sum_{n=1}^N \ln(1 + e^{-y_n \vec{w}^T \vec{x}_n})$$
$$= \frac{1}{N} \sum_{n=1}^N \left[ \mathbb{I}[y_n = 1] \ln \frac{1}{\theta(\vec{w}^T \vec{x}_n)} + \mathbb{I}[y_n \neq 1] \ln \frac{1}{1 - \theta(\vec{w}^T \vec{x}_n)} \right]$$
  - Interpretations
    - $p$ : empirical distribution of  $y_n$  in training data
    - $q$ : predicted probability distribution of  $y_n$  of hypothesis  $h$
    - Minimizing  $E_{in} \Rightarrow$  Make  $q \approx p \Rightarrow$  Make prediction align with data

# Computation of Gradient Descent

- Gradient descent algorithm
  - Initialize  $\vec{w}(0)$
  - For  $t = 0, \dots$ 
    - Compute  $\nabla_{\vec{w}} E_{in}(\vec{w}(t)) = -\frac{1}{N} \sum_{n=1}^N \frac{y_n \vec{x}_n}{1 + e^{y_n \vec{w}(t)^T \vec{x}_n}}$
    - $\vec{w}(t + 1) \leftarrow \vec{w}(t) - \eta \nabla_{\vec{w}} E_{in}(\vec{w}(t))$
    - Terminate if the stop conditions are met
  - Return the final weights
- Which step is the most computationally heavy?
  - Calculate the gradient  $\nabla_{\vec{w}} E_{in}(\vec{w}) = -\frac{1}{N} \sum_{n=1}^N \frac{y_n \vec{x}_n}{1 + e^{y_n \vec{w}^T \vec{x}_n}}$
  - The time complexity is  $O(N)$ 
    - $N$  is large for big datasets

# Deal with Heavy Computation of $\nabla_{\vec{w}} E_{in}(\vec{w})$

- Speed up the implementation of  $\nabla_{\vec{w}} E_{in}(\vec{w}) = -\frac{1}{N} \sum_{n=1}^N \frac{y_n \vec{x}_n}{1 + e^{y_n \vec{w}^T \vec{x}_n}}$ 
  - E.g., “vectorization”
- Solve  $\nabla_{\vec{w}} E_{in}(\vec{w})$  “in expectation”
  - Define  $e_n(\vec{w}) = \ln(1 + e^{-y_n \vec{w}^T \vec{x}_n})$ , the point-wise error caused by  $(\vec{x}_n, y_n)$
  - Observe that
    - $E_{in}(\vec{w}) = \frac{1}{N} \sum_{n=1}^N e_n(\vec{w})$
    - $\nabla_{\vec{w}} E_{in}(\vec{w}) = \frac{1}{N} \sum_{n=1}^N \nabla_{\vec{w}} e_n(\vec{w})$  (gradient of dataset is the average gradient of points)
- Draw a point  $\vec{x}_n$  from  $\{\vec{x}_1, \dots, \vec{x}_N\}$  uniformly at random
  - $E_{\vec{x}_n}[\nabla_{\vec{w}} e_n(\vec{w})] = \nabla_{\vec{w}} E_{in}(\vec{w})$

# Stochastic Gradient Descent (SGD)

- Algorithm
  - Initialize  $\vec{w}(0)$
  - For  $t = 0, \dots$ 
    - Randomly choose a data point  $n$  from  $\{1, \dots, N\}$
    - $\vec{w}(t + 1) \leftarrow \vec{w}(t) - \eta \nabla_{\vec{w}} e_n(\vec{w}(t))$
    - Terminate if the stop conditions are met
  - Return the final weights
- $\mathbb{E}[\nabla_{\vec{w}} e_n(\vec{w})] = \nabla_{\vec{w}} E_{in}(\vec{w})$ 
  - SGD is doing the same thing as GD **in expectation**
    - More efficient (scale to large dataset), suitable for online data, helps escaping local min, etc.
    - Noisier, harder to define stop criteria




# Mini-Batch Gradient Descent

- GD: Computationally heavy, stable updates
- SGD: Computationally light, noisy updates
- Middle ground: Mini-Batch Gradient Descent
  - In each iteration, randomly choose  $k$  points  $\{n(1), \dots, n(k)\}$
  - Update rule
    - $\vec{w}(t + 1) \leftarrow \vec{w}(t) - \eta \frac{1}{k} \sum_{i=1}^k \nabla_{\vec{w}} e_{n(i)}(\vec{w}(t))$
- Side-note about HW2
  - Please report your results on GD (non-stochastic version).
    - You should feel free to play around with SGD or mini-batch on your own.


# Non-Linear Transformation

# Limitations of Linear Models




# Using Non-Linear Transformations

- Find a feature transform  $\Phi$  that map data from  $\vec{x}$  space to  $\vec{z}$  space




# Using Non-Linear Transformations

- Learn a linear classifier in  $\vec{z}$  space:  $g^{(z)}(\vec{z}) = \text{sign}(\vec{w}^{(z)T} \vec{z})$



$$\vec{x} = (1, x_1, x_2)$$

$$\vec{z} = \Phi(\vec{x})$$




$$\vec{z} = (1, x_1^2, x_2^2)$$

$$g^{(z)}(\vec{z}) = \text{sign}(-0.6 + z_1 + z_2)$$

# Using Non-Linear Transformations


- Transform the learned hypothesis back to  $\vec{x}$  space

- $g(\vec{x}) = g^{(z)}(\Phi(\vec{x})) = \text{sign}\left(\vec{w}^{(z)T} \Phi(\vec{x})\right)$



$$\vec{x} = (1, x_1, x_2)$$


$$g(\vec{x}) = \text{sign}(-0.6 + x_1^2 + x_2^2)$$



$$\vec{z} = (1, x_1^2, x_2^2)$$


$$g^{(z)}(\vec{z}) = \text{sign}(-0.6 + z_1 + z_2)$$

# Nonlinear Transformation




# Generalization of Nonlinear Transformation

- Fact (We'll prove this later)
  - The VC Dimension of d-dim perceptron is  $d + 1$
- VC dimension of perceptron on input space  $\vec{x} = (x_0, \dots, x_d)$ 
  - $d+1$
- VC dimension of perceptron on input space  $\vec{z} = (z_0, \dots, z_{d'})$ 
  - $\leq d' + 1$  (usually treated as  $\approx d' + 1$ )
- Careful: Non-linear transform might lead to "nonsense" behavior









# How to Choose Feature Transform $\Phi$



# How to Choose Feature Transform $\Phi$



$$x_1^2 + x_2^2 = 0.6$$



Something Seems Wrong!

Must choose  $\Phi$   
**BEFORE** looking at the data

Otherwise, you are doing “data snooping”

The hypothesis set  $H$  is as large as anything your brain can think of

# Choose $\Phi$ Before Seeing Data

- Rely on domain knowledge (feature engineering)
  - Handwriting digit recognition example
- Use common sets of feature transformation
  - Polynomial transformation
  - 2nd order Polynomial transformation
    - $\vec{x} = (1, x_1, x_2)$
    - $\Phi_2(\vec{x}) = (1, x_1, x_2, x_1x_2, x_1^2, x_2^2)$
    - Pros: more powerful (contains circle, ellipse, hyperbola, etc)
    - Cons: 2-d  $\Rightarrow$  5-d
      - More computation/storage
      - Worse generalization error

The VC dimension of d-dim perceptron is  $d+1$

# Q-th Order Polynomial Transform


- $\vec{x} = (1, x_1, \dots, x_d)$
- From 1-st order to Q-th order polynomial transform:
  - $\Phi_1(\vec{x}) = \vec{x}$
  - $\Phi_2(\vec{x}) = (\Phi_1(\vec{x}), x_1^2, x_1x_2, x_1x_3, \dots, x_d^2)$
  - ...
  - $\Phi_Q(\vec{x}) = (\Phi_{Q-1}(\vec{x}), x_1^Q, x_1^{Q-1}x_2, \dots, x_d^Q)$
- Number of elements in  $\Phi_Q(\vec{x})$

# Q-th Order Polynomial Transform

- $\vec{x} = (1, x_1, \dots, x_d)$
- From 1-st order to Q-th order polynomial transform:
  - $\Phi_1(\vec{x}) = \vec{x}$
  - $\Phi_2(\vec{x}) = (\Phi_1(\vec{x}), x_1^2, x_1x_2, x_1x_3, \dots, x_d^2)$
  - ...
  - $\Phi_Q(\vec{x}) = (\Phi_{Q-1}(\vec{x}), x_1^Q, x_1^{Q-1}x_2, \dots, x_d^Q)$
- Number of elements in  $\Phi_Q(\vec{x})$ 
  - $\binom{Q+d}{Q}$

# Structural Hypothesis Sets


- Let  $H_Q$  be the linear model for the  $\Phi_Q(\vec{x})$  space



- Let  $g_Q = \operatorname{argmin}_{h \in H_Q} E_{in}(h)$ 
  - $H_0 \quad H_1 \quad H_2 \dots$
  - $d_{vc}(H_0) \quad d_{vc}(H_1) \quad d_{vc}(H_2) \dots$
  - $E_{in}(g_0) \quad E_{in}(g_1) \quad E_{in}(g_2) \dots$

# Structural Hypothesis Sets

- Let  $H_Q$  be the linear model for the  $\Phi_Q(\vec{x})$  space



- Let  $g_Q = \operatorname{argmin}_{h \in H_Q} E_{in}(h)$ 
  - $H_0 \subset H_1 \subset H_2 \dots$
  - $d_{vc}(H_0) \leq d_{vc}(H_1) \leq d_{vc}(H_2) \dots$
  - $E_{in}(g_0) \geq E_{in}(g_1) \geq E_{in}(g_2) \dots$

