

CSE 417T

Introduction to Machine Learning

Lecture 23

Instructor: Chien-Ju (CJ) Ho

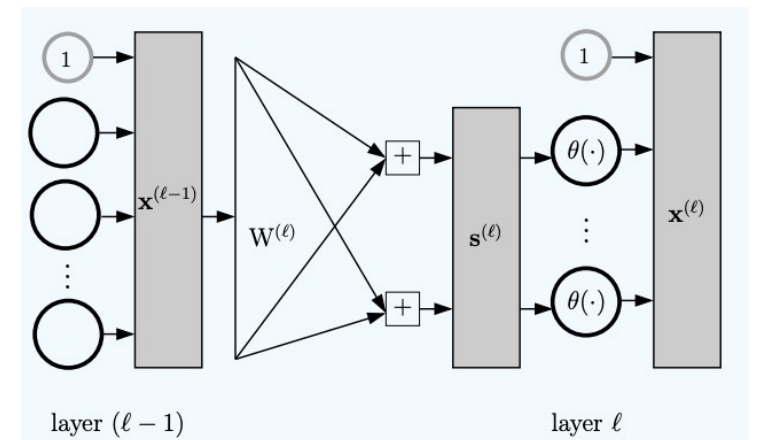

- Homework 5: due **April 30** (Friday)
 - Again, please start it early
- Exam 2: (**May 4**, Tuesday)
 - Duration: 75+5 Minutes
 - Content: Focus on the content of 2nd half of the semester
 - Though knowledge is cumulative
 - Time: by default, everyone is expected to take it **during lecture time**
 - If you can't, let me know **by this Friday** through [this google form](#)
 - Review lecture: Apr 29
 - Practice questions will be posted next week
 - Other logistics are the same as Exam 1
 - Format: Gradescope online exam + Zoom (with camera on)
 - Information access during exam:
 - Allowed: Textbook, slides, hardcopy materials (e.g., your own notes)
 - Not allowed: search for information online during exam, talk to any other persons
 - **Follow Piazza announcements** for updates/information

Recap

Neural Networks (NN)

- Notations:
 - $\ell = 0$ to L : layer
 - $d^{(\ell)}$: dimension of layer ℓ
 - $\vec{x}^{(\ell)}$: the nodes in layer ℓ
 - $w_{i,j}^{(\ell)}$: weights; characterize hypothesis in NN
 - $s_j^{(\ell)} = \sum_{i=0}^{d^{(\ell-1)}} w_{i,j}^{(\ell)} x_i^{(\ell-1)}$: linear signals
 - θ : activation function
 - $x_j^{(\ell)} = \theta(s_j^{(\ell)})$

Feed-forward network




Forward Propagation and Backpropagation

- Evaluate $h(\vec{x})$ given h (characterized by $\{w_{i,j}^{(\ell)}\}$)
 - Forward propagation

$$\mathbf{x} = \mathbf{x}^{(0)} \xrightarrow{w^{(1)}} \mathbf{s}^{(1)} \xrightarrow{\theta} \mathbf{x}^{(1)} \xrightarrow{w^{(2)}} \mathbf{s}^{(2)} \xrightarrow{\theta} \mathbf{x}^{(2)} \dots \xrightarrow{w^{(L)}} \mathbf{s}^{(L)} \xrightarrow{\theta} \mathbf{x}^{(L)} = h(\mathbf{x}).$$

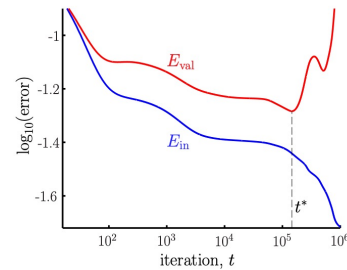
- Given D , learn the weights $W = \{w_{i,j}^{(\ell)}\}$
 - Backpropagation
 - Initialize $w_{i,j}^{(\ell)}$ randomly
 - For $t = 1$ to T
 - Randomly pick a point from D (for stochastic gradient descent)
 - **Forward propagation:** Calculate all $x_i^{(\ell)}$ and $s_i^{(\ell)}$
 - **Backward propagation:** Calculate all $\delta_j^{(\ell)}$
 - Update the weights $w_{i,j}^{(\ell)} \leftarrow w_{i,j}^{(\ell)} - \eta \delta_j^{(\ell)} x_i^{(\ell-1)}$
 - Return the weights



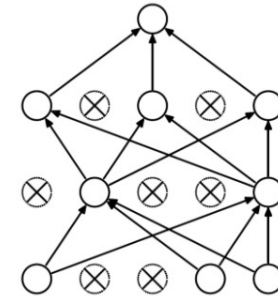
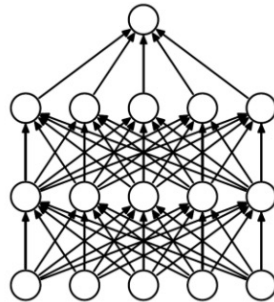
Regularizations in Neural Networks

- Weight-based regularization

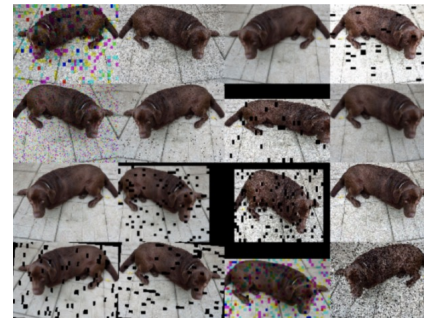
- Early stopping



- Dropout




- Adding noises



Deep Learning (NN with many layers)

- “Shallow” neural network is powerful (universal approximation theorem holds with a single hidden layer). Why “deep” neural networks?



Each layer captures **features** of the previous layers.

We can use “raw data” (e.g., pixels of an image) as input. The hidden layer are extracting the **features**.

Design different **network architectures** to incorporate domain knowledge.

Some Techniques in Improving Deep Learning

- Regularization to mitigate overfitting
 - Weight-based, early stopping, dropout, etc
- Incorporating domain knowledges
 - Network architectures (e.g., Convolutional Neural Nets)
- Improving computation with huge amount of data
 - Hardware architecture to improve parallel computation
- Improving gradient-based optimization
 - Choosing better **initialization** points

Today's Lecture

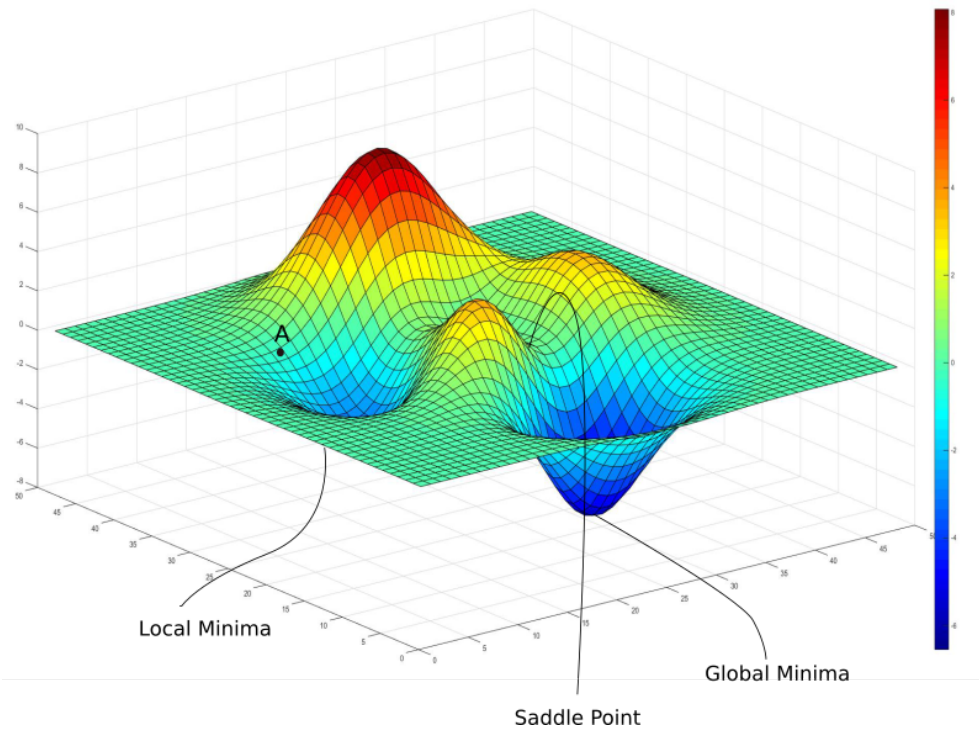
The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.
Let me know if you spot errors.

Initialization

Why initialization matters in deep learning

- Error is nonconvex in NN
- Vanishing/exploding gradient problem

Error is Nonconvex in Neural Networks



- We mostly adopt gradient-descent-style algorithms for optimization.
- No guarantee to converge to global optimal.
- Need to run it many times.
- Initialization matters!

Vanishing Gradient Problem

- Backpropagation

- $\frac{\partial e_n(W)}{\partial w_{i,j}^{(\ell)}} = \delta_j^{(\ell)} x_i^{(\ell-1)}$

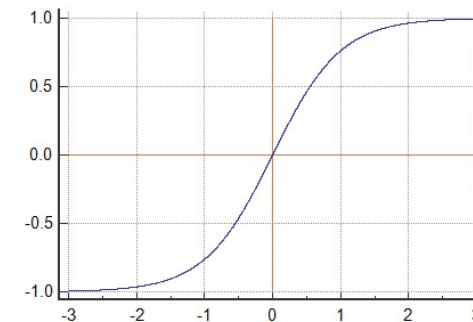
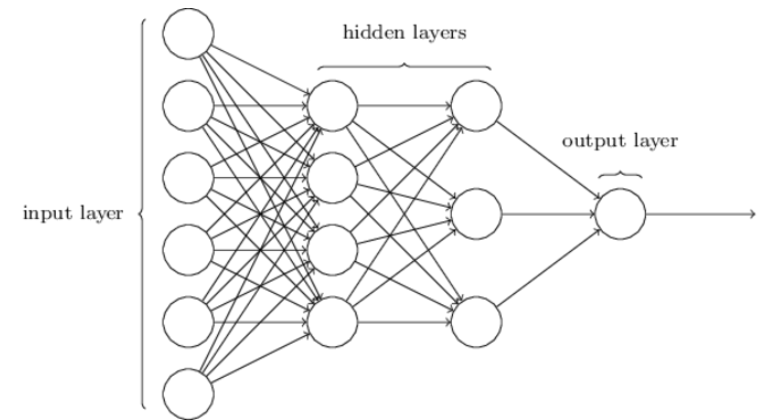
- $\delta_j^{(\ell)} = \theta' \left(s_j^{(\ell)} \right) \sum_{k=1}^{d^{(\ell+1)}} \delta_k^{(\ell+1)} w_{j,k}^{(\ell+1)}$

- If we use activation function $\theta(s) = \tanh(s)$

- $\theta'(s) = 1 - \theta(s)^2 < 1$

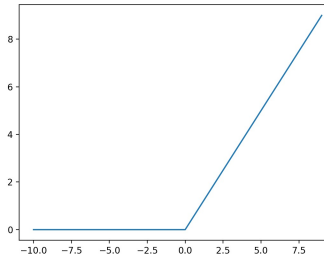
- In deep learning with a lot of layers,

- the gradient might vanish
 - hard to update the early layers



Vanishing Gradient Problem

- $\delta_j^{(\ell)} = \theta' \left(s_j^{(\ell)} \right) \sum_{k=1}^{d^{(\ell+1)}} \delta_k^{(\ell+1)} w_{j,k}^{(\ell+1)}$
- There is also a corresponding “exploding gradient problem”
- What can we do
 - Choose more suitable activation functions
 - One common choice is Rectified Linear Unit (ReLU) and its variant
 - $\theta(s) = \max(0, s)$
 - Choose better **initialization**
 - Many approaches

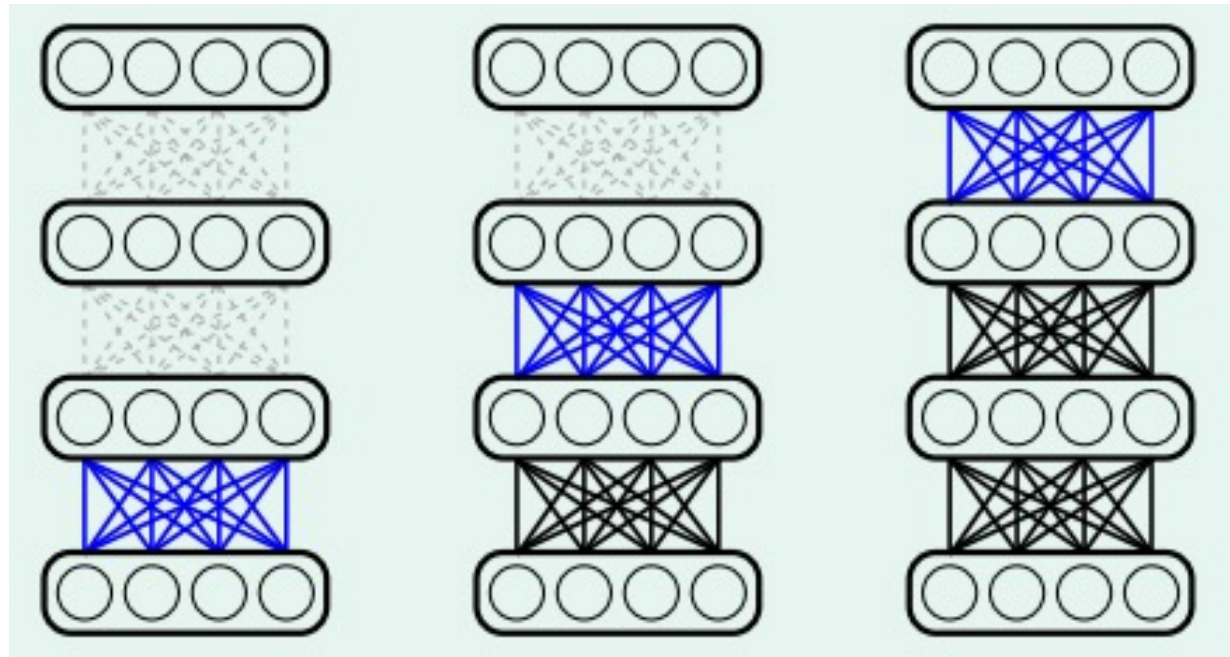


Weight Initialization


- Initializing weights to all 0 is a bad idea
 - Q5 of HW1
 - Hint: Look at the backpropagation formulation
- Randomly Initializing weights to regions so that vanishing/exploding gradients are less likely to happen
 - Activation-function dependent
 - e.g., Xavier initialization for tanh
- Learning the initialization that might be closer to the optimal
 - E.g., using autoencoder

Initialization

- Hard to initialize the entire network well.
- Intuition: Initialize the weights **layer by layer** such that each layer **preserves** the properties of the previous layer.



Autoencoder



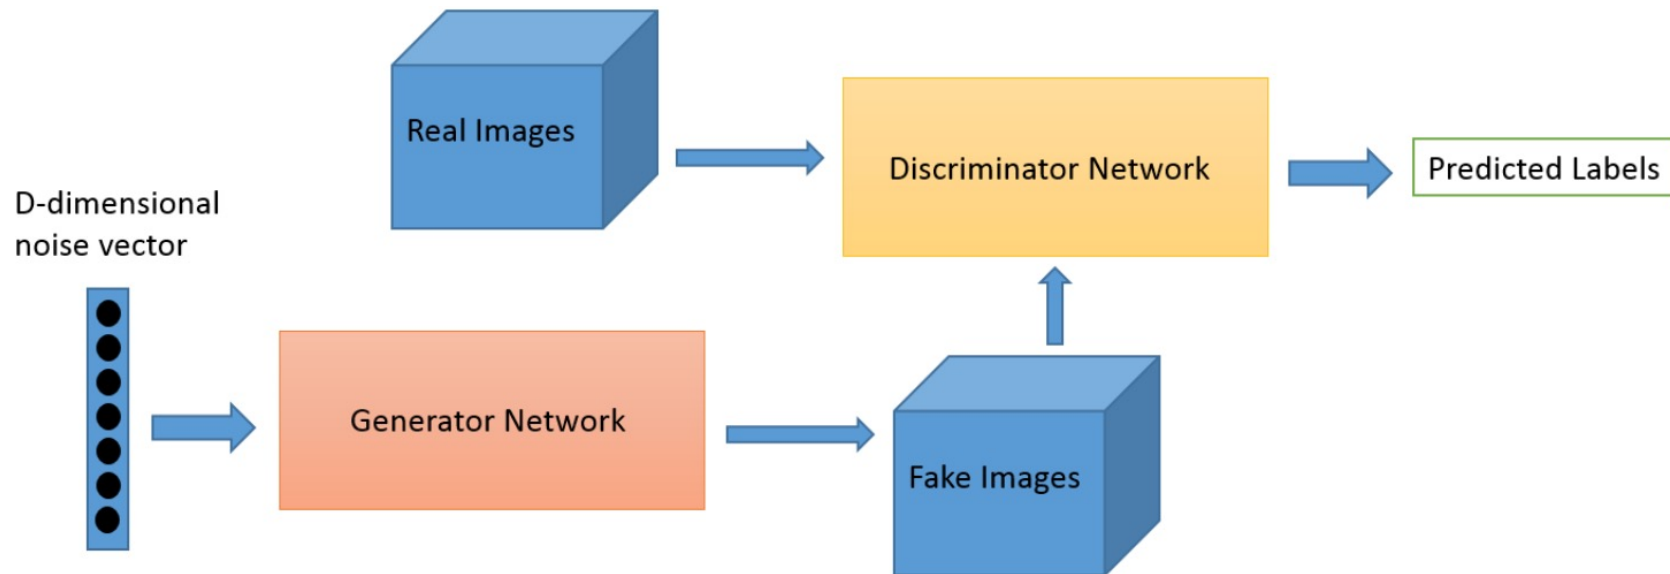
Unsupervised learning!

Cool Stuffs for Deep Learning

[Safe to Skip for the exam]

Generative Adversarial Nets (GAN)

- A Competition: Generator vs Discriminator
 - Discriminator wants to correctly classify the images (true images or not)
 - Generator wants to generate images that discriminator can't classify




[Safe to Skip for the Exam]



<https://thisPersonDoesNotExist.com/>

Style Transfer



- Informal intuitions:
 - Recall that we can treat hidden layers as **feature transforms**
 - Deep learning is learning **representation** of data
 - How to achieve style transfer:
 - Learn a **content representation** for an image using hidden layers
 - Learn a **style representation** for an image using hidden layers
 - Compute an image that jointly minimizes the distance from the content image's content representation and the style image's style representation
 - <https://arxiv.org/pdf/1508.06576.pdf>

Radial Basis Functions (RBF)

Instance-Based Learning

- Make predictions based on data instances
- k -nearest neighbor (K-NN)
 - $g(\vec{x}) = \text{sign}(\sum_{i=1}^k y_{[i]}(\vec{x}))$
 - Predict according to the nearest neighbors
- Kernel SVM
 - $g(\vec{x}) = \text{sign}(\sum_{\alpha_n^* > 0} \alpha_n^* y_n K(\vec{x}_n, \vec{x}) + b^*))$
 - Predict according to support vectors
- Radial Basis Functions (RBF)
 - Focus of today

Radial Basis Functions

- Think about k -nearest neighbor (K-NN) again
 - $g(\vec{x}) = \text{sign}(\sum_{i=1}^k y_{[i]}(\vec{x}))$
 - Make predictions based on k nearest data points
 - Each of the k data points has the same weight
- Natural questions:
 - Can we use more (or even all) data?
 - **Weight** them based on **how close** data points are to \vec{x}

Radial Basis Functions

- Given dataset $D = \{\vec{x}_1, \dots, \vec{x}_N\}$
- Task: Make a prediction on \vec{x}

- Radial Basis Function:

- $g(\vec{x}) = \frac{1}{Z(\vec{x})} \sum_{n=1}^N \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right) y_n$

- $\phi(s)$: a monotonically decreasing function

- This is for regression. We can take a sign and make it a classification.
- $Z(\vec{x}) = \sum_{m=1}^N \phi\left(\frac{\|\vec{x} - \vec{x}_m\|}{r}\right)$ is for normalization

- It's called **radial** basis function since it takes the **distance** to the points as the basis function

Radial Basis Functions

- Radial Basis Function:


- $g(\vec{x}) = \sum_{n=1}^N \frac{1}{Z(\vec{x})} \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right) y_n$

- Example of ϕ

- Gaussian RBF (we have seen this in SVM): $\phi(s) = e^{-s}$


- Intuitions

- The impact of \vec{x}_n to \vec{x} is higher if it's closer to \vec{x}
 - The role of r is similar to k in k -NN
 - $r = 0$: 1-NN
 - $r \rightarrow \infty$: N -NN (i.e., $k = N$)




Changing the Viewpoints

- ϕ centered around \vec{x}



$$g(\vec{x}) = \sum_{n=1}^N \frac{y_n}{Z(\vec{x})} \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$$

- ϕ centered around \vec{x}_n



$$g(\vec{x}) = \sum_{n=1}^N w_n(\vec{x}) \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$$

From Nonparametric to Parametric RBF

- Nonparametric RBF

- $g(\vec{x}) = \sum_{n=1}^N \frac{y_n}{Z(\vec{x})} \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$

- $g(\vec{x}) = \sum_{n=1}^N w_n(\vec{x}) \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$

- The hypothesis is defined by dataset

- Parametric RBF hypothesis set

- $h(\vec{x}) = \sum_{n=1}^N w_n \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$

- Learn w_n from data

Parametric RBF => Linear Models

- Parametric RBF is linear model with nonlinear transformation

- $h(\vec{x}) = \sum_{n=1}^N w_n \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$

- The projection $\Phi(\vec{x}) = \begin{bmatrix} \phi\left(\frac{\|\vec{x} - \vec{x}_1\|}{r}\right) \\ \phi\left(\frac{\|\vec{x} - \vec{x}_2\|}{r}\right) \\ \vdots \\ \phi\left(\frac{\|\vec{x} - \vec{x}_N\|}{r}\right) \end{bmatrix}$


- We can apply what we learned in linear models to learn w_n
- However, this seems to be **overfitting** (N parameters for N points)

From N points to K points

- Use only K points $(\vec{\mu}_1, \dots, \vec{\mu}_K)$

- $h(\vec{x}) = \sum_{k=1}^K w_k \phi\left(\frac{\|\vec{x} - \vec{\mu}_k\|}{r}\right)$

- Which K points?
 - We can find K representative points
 - Use clustering algorithms, e.g., Lloyd algorithm as introduced earlier
 1. Randomly pick K points as centers
 2. Create the Voronoi regions as clusters
 3. Update the centers (calculating the mean)
 4. Update the region
 5. Repeat 3 and 4



More Discussion on RBF


- $h(\vec{x}) = \sum_{k=1}^K w_k \phi\left(\frac{\|\vec{x} - \vec{\mu}_k\|}{r}\right)$

- Connection to linear models
 - Parametric RBF is essentially linear model with nonlinear transformation
- Connection to nearest neighbor
 - Radial Basis Function is defined by “similarity”
 - A prediction for a point is based on the “similarity” of the points to be predicted and other points


More Discussion on RBF

- $h(\vec{x}) = \sum_{k=1}^K w_k \phi\left(\frac{\|\vec{x} - \vec{\mu}_k\|}{r}\right)$
- Connection to SVM
 - Gaussian RBF Kernel: $g(\vec{x}) = \text{sign}\left(\sum_{\alpha_n^* > 0} \alpha_n^* y_n K(\vec{x}_n, \vec{x}) + b^*\right)$
 - Use **cluster centers** or **support vectors** to characterize a hypothesis

RBF




SVM




More Discussion on RBF

- $h(\vec{x}) = \sum_{k=1}^K w_k \phi\left(\frac{\|\vec{x} - \vec{\mu}_k\|}{r}\right)$
- Connection to Neural Network
 - RBF can be graphically presented:
 - (Similar to SVM, it's a single-hidden layer NN)

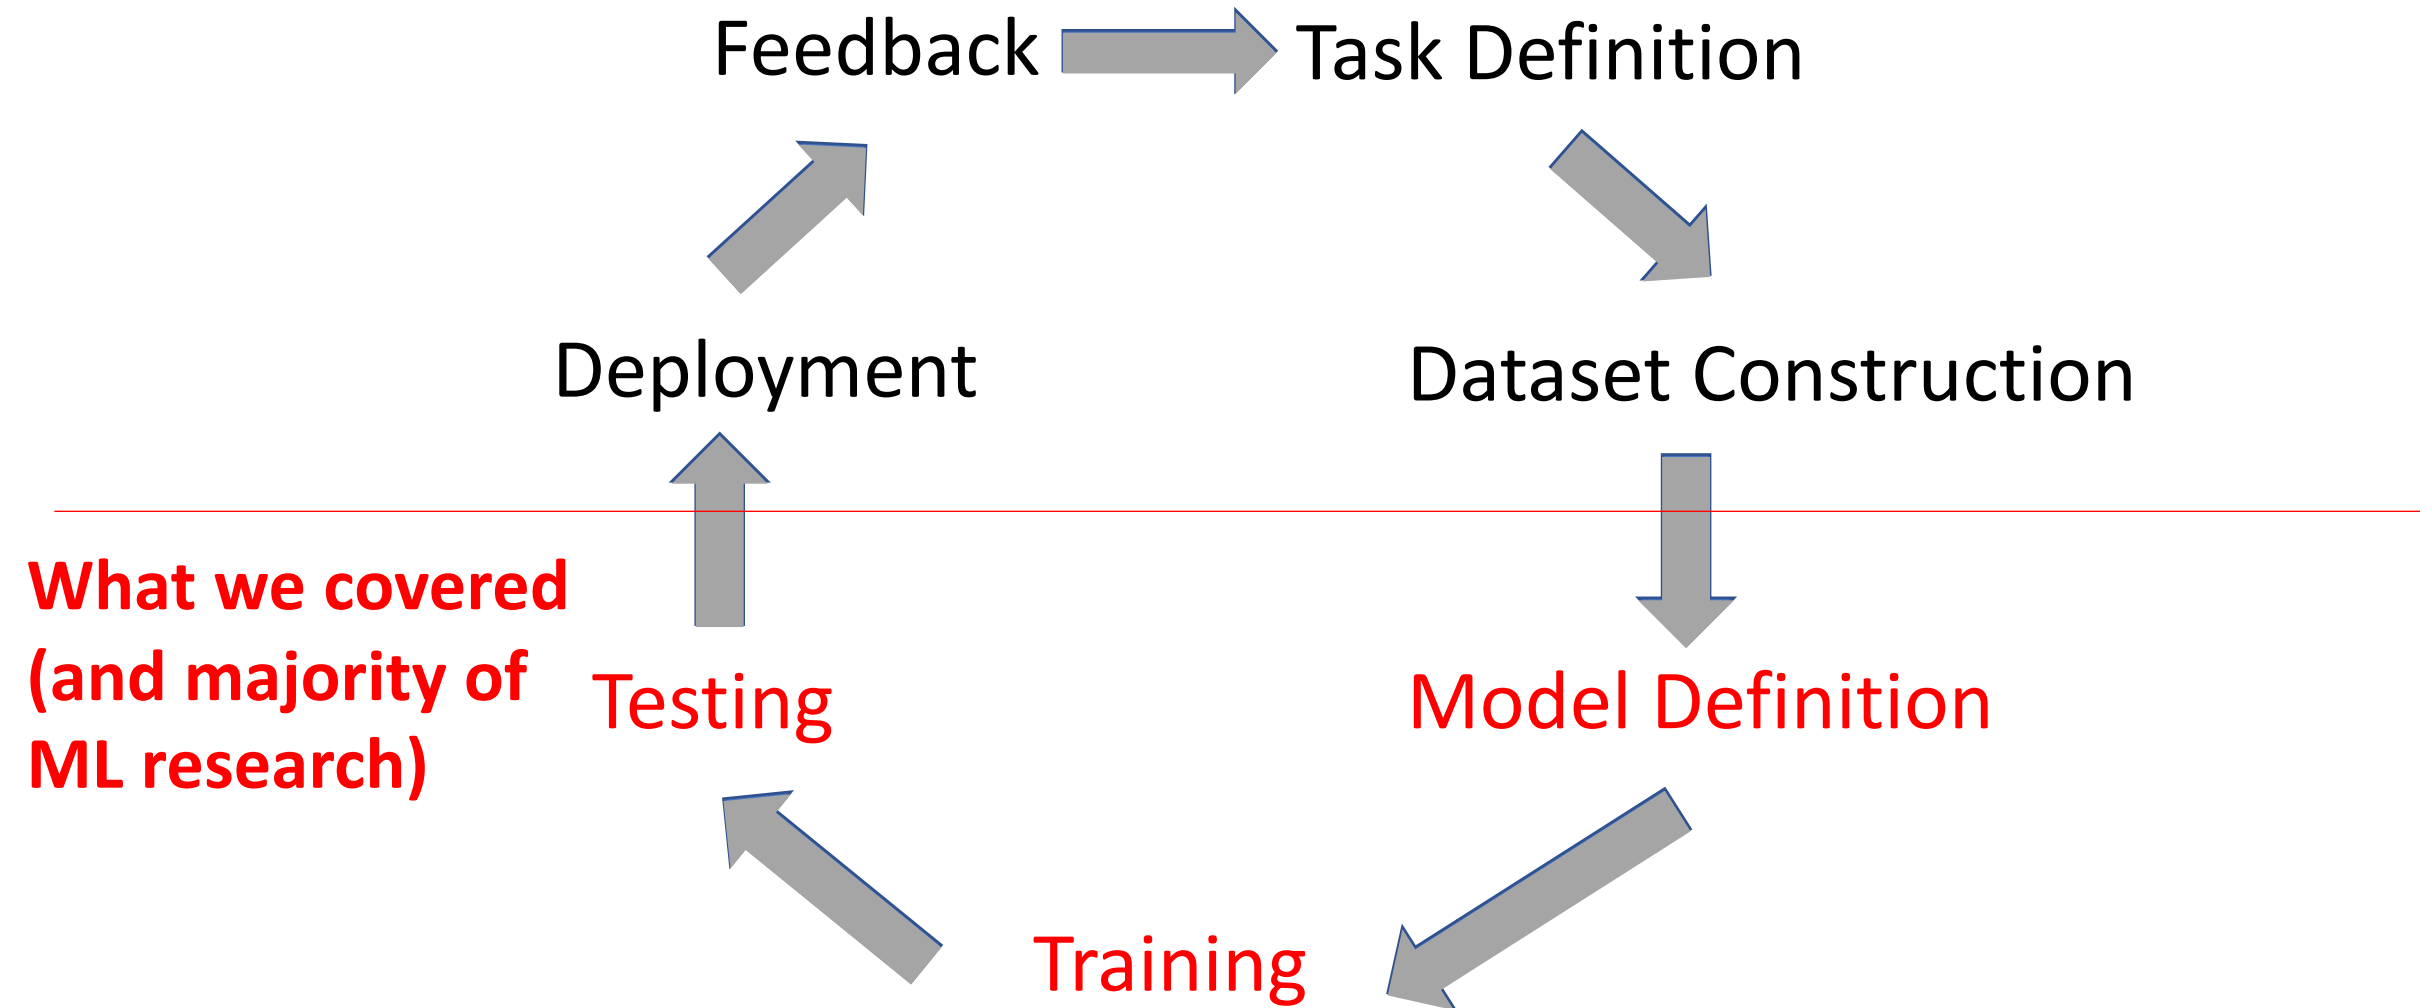


Discussion Beyond This Course

Machine Learning Lifecycle



Machine Learning Lifecycle



Machine Learning Lifecycle

To have “positive” impacts,
we need to be careful in
every stage

Feedback → Task Definition

Deployment




Dataset Construction

Testing

Model Definition

Training

What we covered
(and majority of
ML research)




Strategic Classification

Or more recently, people start to call it “performative prediction”

Classification

- Standard setup of (supervised) machine learning




- Finding patterns from the given training datasets
 - Use the pattern to make predictions on new testing data
-
- Fundamental assumption:
 - Training and testing data points are i.i.d. drawn from the same distribution


Example: Spam Filter

honest emails


**spam emails
(if no detection)**







Example: Spam Filter



Example: Spam Filter




Social Program Eligibility [Camacho and Conover, 2012]




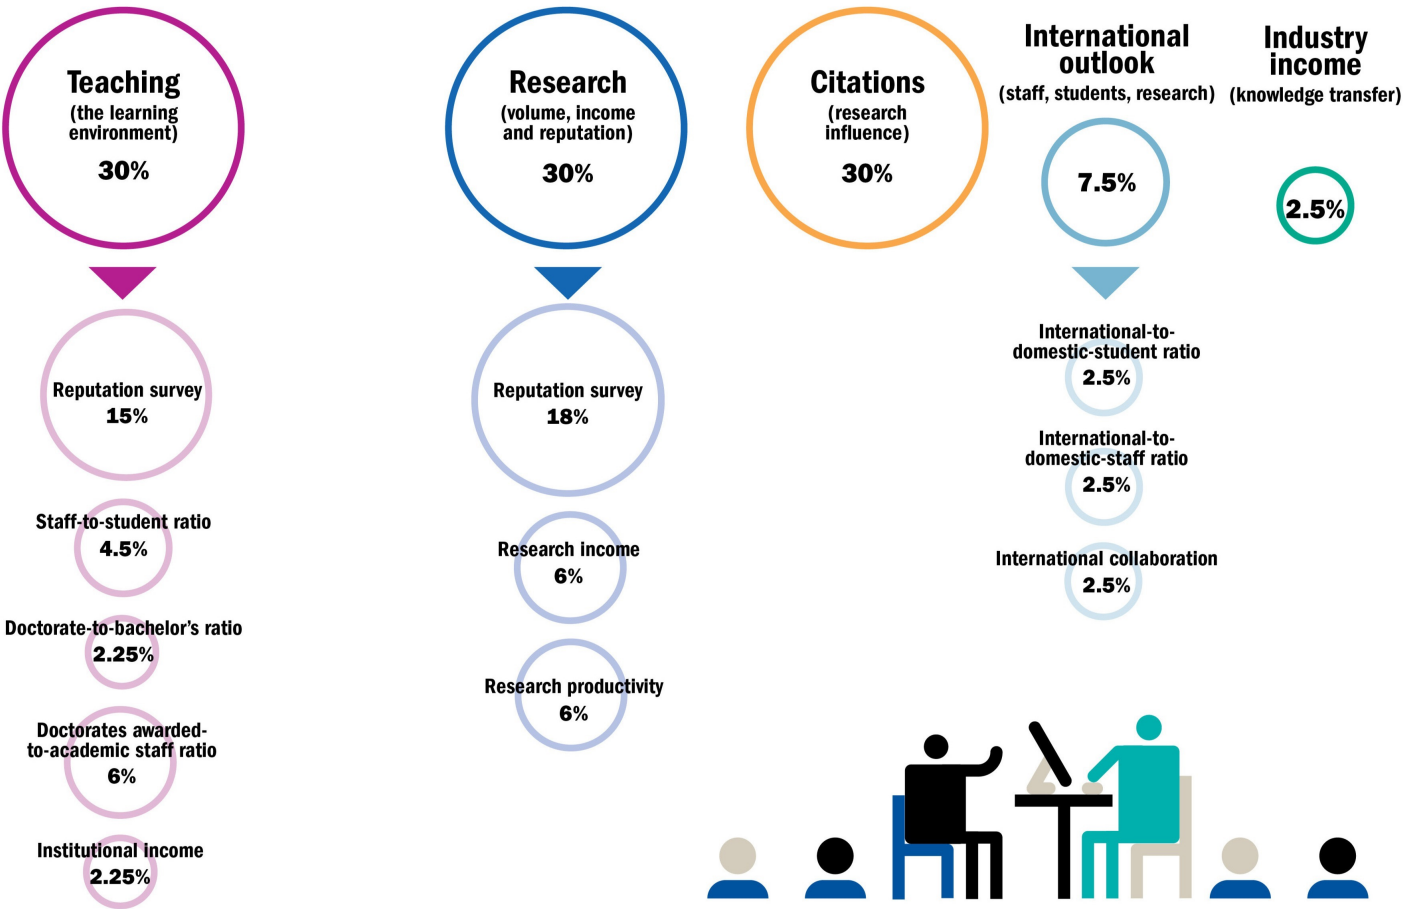
Goodhart's law:

“If a measure becomes the public's goal,
it is no longer a good measure.”

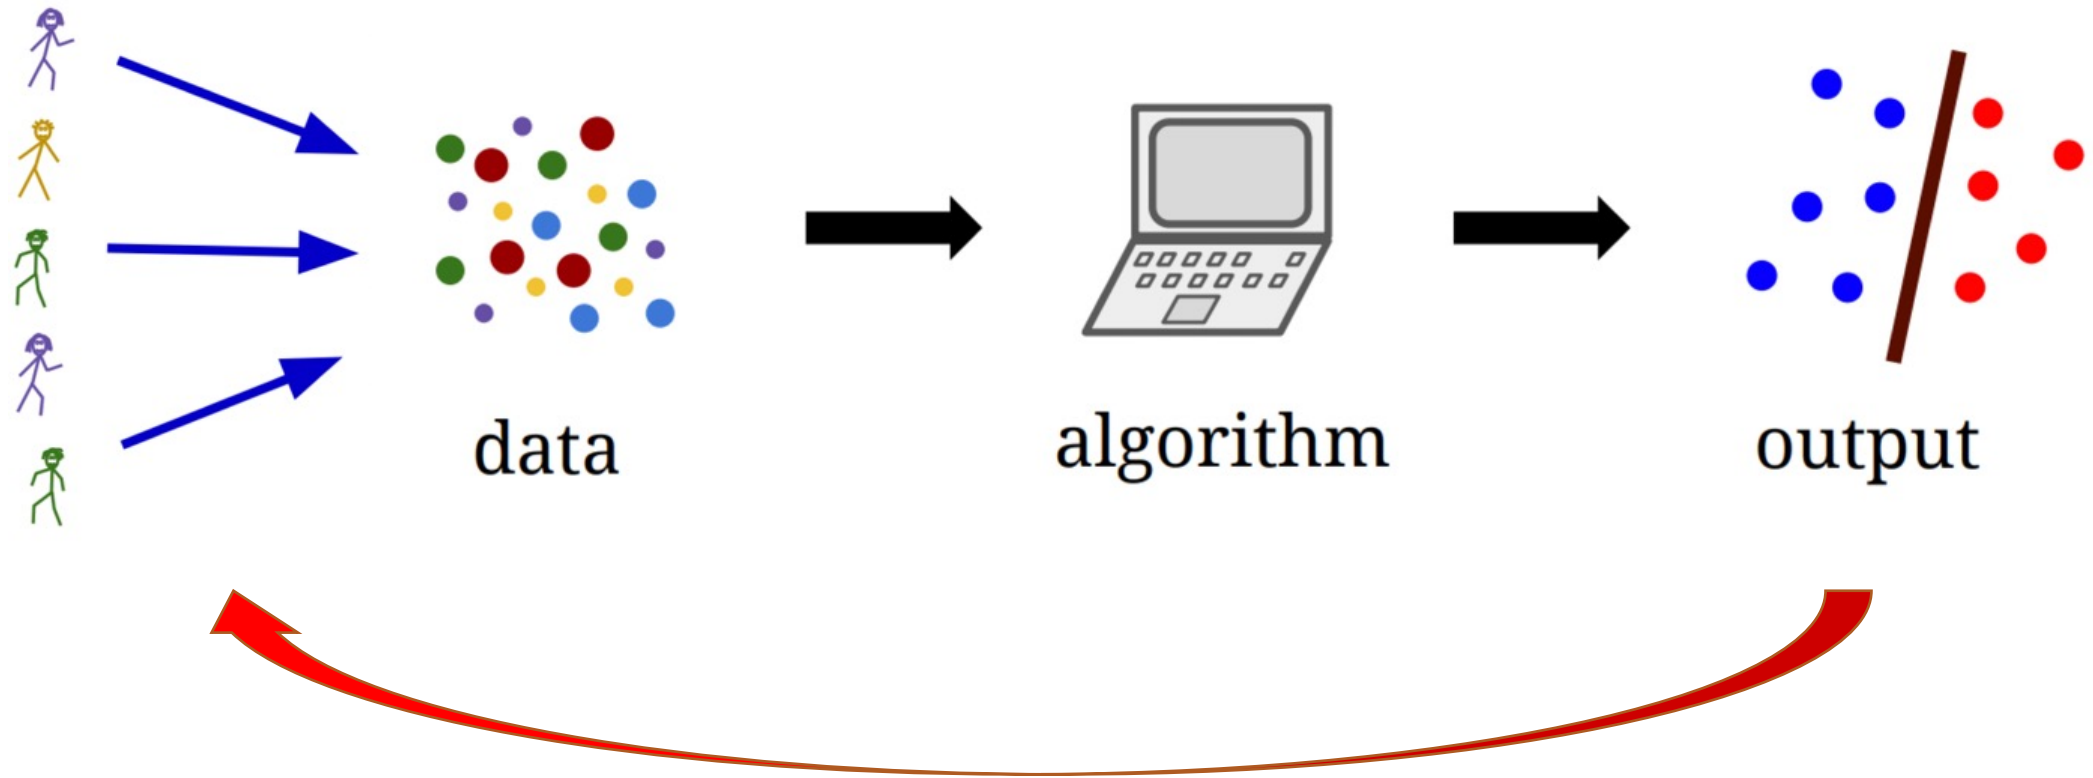
COLLEGE ADMISSIONS



Methodology



Strategic Classification



How to take this interaction between ML algorithms and data-holders into account?

Game theoretical modeling

Game Theoretical Modeling

- Key elements:
 - Players, actions, payoffs
- Players: Jury (e.g., university) and Contestants (student applicants)
- Actions:
 - First, Jury decides on the machine learning model (binary classification)
 - Then, Contestant decides how to alter their features based on the model
- Payoffs
 - Jury wants to maximize the probability of correct predictions
 - Contestants want to be selected (being predicted as 1)