Modeling Humans: Incentive Design

Instructor: Chien-Ju (CJ) Ho

Logistics: Presentations

- For presenters
  - No need to submit the review questions.
  - Talk to me **one week before your presentation**.
    - Default: talk to me after class
    - Finish reading the required paper and at least 1~2 optional papers.
    - Have an idea of your presentation plan.
    - Come up with at least two discussion questions.
    - Try to look for materials outside of the reading list to include in your presentation (e.g., real-world examples/platforms related to the topics you present)

Logistics: Presentations

- For non-presenters:
  - Attend the lecture and engage in discussion.
  - Everyone needs to fill in peer review forms.
  - Comments are not anonymous to me but anonymous to the presenters.
    - Please try to give constructive comments.

CSE 518A Review Form: Sep 25

ID:

How much do you like the presentations:
(Score 1-10, with 10 the highest)

Comments to the presenters:
Try to give constructive comments. They will be anonymized before giving to the presenters.

Logistics: Projects

- I have reserved some blocks of time if you want to discuss:
  - https://calendly.com/chienjuho/30min

- Milestone 1:
  - Due: Oct 9
  - Literature survey
  - Concrete plan on what you plan to do (including a timeline)

Logistics: Projects

- If you plan to run MTurk experiments:
  - Recommend to start with a small-scale design so that MTurk is not necessary
    - Have a small-scale / prototype design that requires small data (from you and/or friends)
    - Being able to utilize existing datasets or simulated data

  - Schedule a meeting with me well in advance
    - You might need to go through the university IRB process, which requires quite a bit of logistics.
    - Be prepared to take online classes and complete required documents/applications.
    - Do-able, but require additional planning.
    - Spare a few more weeks for completing the process.

Logistics: Why IRB?

- “Institutional Review Boards (IRBs) were created, as required by federal law, to review and oversee research involving humans. The mission of the IRB is to protect the rights and welfare of individuals recruited for, or participating in, human subject research.”

- Before the IRB age…
  - The Twin Study
  - Stanford Prisoners’ experiments

Logistics: Assignment 3

- Will be announced around the deadline of Assignment 2.
- Mostly mathematical questions on game theory / scoring rules / peer predictions (content of today's lecture).

Today’s Lecture: Modeling Incentives

Warm-Up Discussion

- What are the incentives that motivate you to do things in your life?

- Can you try to model the incentives? For example, how do you mathematically specify the relationships of the "strength" of the incentives and your actions?

- Mechanism design: Given your understanding of how humans respond to incentives, how can you design mechanisms (set of rules) that encourage people to do what you want them to do? (You can try to think about these in some specific applications).

Today's Lecture: Modeling Incentives

- Game theory basics
  - Utility, Games, Equilibrium
  - Example usage in crowdsourcing
    - Contract design (Principal Agent Model)

- Proper scoring rules (Eliciting Truthful Probability Estimates)
  - Example usage in crowdsourcing
    - Prediction markets

- Peer prediction
  - Example usage in crowdsourcing
    - Peer grading in MOOCs

Game Theory

- Mathematical study of interactions between **rational** and **self-interested** agents.

- Agents are often assumed to be rational and choose actions to maximize their **expected utility.**

Utility

- A way to quantify agents' preferences over the state of the world.
- Example

$$\Omega = \{\text{Sunny}, \text{Cloudy}, \text{Rainy}\}$$

$$\text{Sunny} \succ \text{Cloudy} \succ \text{Rainy}$$

(Sunny is preferred over Cloudy and Rainy, and Cloudy is preferred over Rainy)

- Using von Neumann–Morgenstern utility

u(Sunny) = 10, u(Cloudy) = 5, u(Rainy) = 3

Expected Utility Theory

- Agents take actions to maximize their expected utility

$$\sum_{\omega \in \Omega} p(\omega) u(\omega)$$

- There have been discussions/extensions on risk attitudes, irrational behavior, etc.

- Game theory deals with situations in which $p(\omega)$ and $u(\omega)$ are influenced by **agents' joint actions**
  - For example, states could be {hungry, not hungry}, and the states depend on your actions {eat, not eat}

Example 1: Prisoner's Dilemma

B

Normal-form game

A

	Stay Silent	Confess
Stay Silent	A: 6 months B: 6 months	A: 10 years B: free
Confess	A: free B: 10 years	A: 5 years B: 5 years

Solution Concept

What should the prisoners do?

"Confess" is a **dominant strategy** – it maximizes the prisoner's utility no matter what action the other player chooses.

Normal-Form Game

- Players take actions simultaneously
- The elements of a normal-form game
  - Players: (prisoner A, prisoner B)
  - Strategies: (stay silent, confess)
  - Payoff: (sentences for all strategy combinations)

	B Stay Silent	B Confess
A Stay Silent	A: 6 months B: 6 months	A: 10 years B: free
A Confess	A: free B: 10 years	A: 5 years B: 5 years

Example 2: Coordination Game

- Two friends A and B are deciding what to do on Friday night
  - A prefers to go to the movie
  - B prefers to go to the bar
  - Both prefer to do something together than doing something separately

B

A

	Movie	Bar
Movie	(2, 1)	(0, 0)
Bar	(0, 0)	(1, 2)

"Nash equilibrium" of this game

What should A and B do?

(Movie, Movie) and (Bar, Bar) seem to be two ***stable*** outcomes

Example 3: Rock, Paper, Scissors

- A zero-sum game

	Rock	Paper	Scissors
Rock	(0, 0)	(-1, 1)	(1, -1)
Paper	(1, -1)	(0, 0)	(-1, 1)
Scissors	(-1, 1)	(1, -1)	(0, 0)

Is there a dominant strategy?

Are there stable pairs of actions?

Need a more general solution concept.

Definitions

- Mixed strategy
  - Let S_i be the set of all probability distributions over A_i
  - Each $s_i \in S_i$ is a ***mixed strategy***, where
    - $s_i(a_i)$ denotes the probability for agent i choosing action a_i

- Payoffs can be calculated using expected utility

Definitions

- Best response
  - $\vec{s}$: a strategy profile, i.e., the set of strategies for all agents
  - $\vec{s}_{-i}$: the strategies of all agents except i

  - A strategy $s_i^* \in S_i$ is a ***best response*** to $\vec{s}_{-i}$ if

$$u_i(s_i^*, s_{-i}) \geq u_i(s_i, s_{-i}) \text{ for all } s_i \in S_i$$

  - There could be multiple best responses.
  - If a best response put positive probability on 2 or more actions, the agent must be indifferent between these actions. (Think about why)

Definitions

- Nash equilibrium
  - A strategy profile $\vec{s}$ is a ***Nash equilibrium*** if, for all agent i, $\mathrm{s_i}$ is a best response to $\vec{s}_{-i}$
- Intuitive interpretations:
  - If all agents except i follow the strategies in the Nash equilibrium, agent i would maximize her payoff by following the strategy in Nash.
  - No incentive to deviate if everyone else follows the strategy in Nash

Let’s look at the examples again

	B Stay Silent	B Confess
A Stay Silent	A: 6 months B: 6 months	A: 10 years B: free
A Confess	A: free B: 10 years	A: 5 years B: 5 years

- (Confess, Confess) is the dominant strategy equilibrium
- Strongest solution concept

Let’s look at the examples again

	Movie	Bar
Movie	(2, 1)	(0, 0)
Bar	(0, 0)	(1, 2)

- (Movie, Movie) and (Bar, Bar) are pure strategy Nash equilibria
- Are there other equilibria?
  - Using the indifference property to derive a mixed-strategy equilibria

Let’s look at the examples again

	Rock	Paper	Scissors
Rock	(0, 0)	(-1, 1)	(1, -1)
Paper	(1, -1)	(0, 0)	(-1, 1)
Scissors	(-1, 1)	(1, -1)	(0, 0)

- Both players play each action with 1/3 probability is the mixed strategy Nash equilibrium
- It is the unique equilibrium

Theorem (Nash, 51):

Every game with a finite number of players and actions has at least one equilibrium.

Notes:

- This is the "existence" proof. Computing the equilibrium could be hard (computationally expensive).
- If it's hard to calculate the equilibrium, can we really expect humans to follow the equilibrium?
  - There have recent studies developing new solution concepts which assume humans are "learning" to adapt to the game.

Extensive-Form Game

- Agents take decisions in a sequential manner

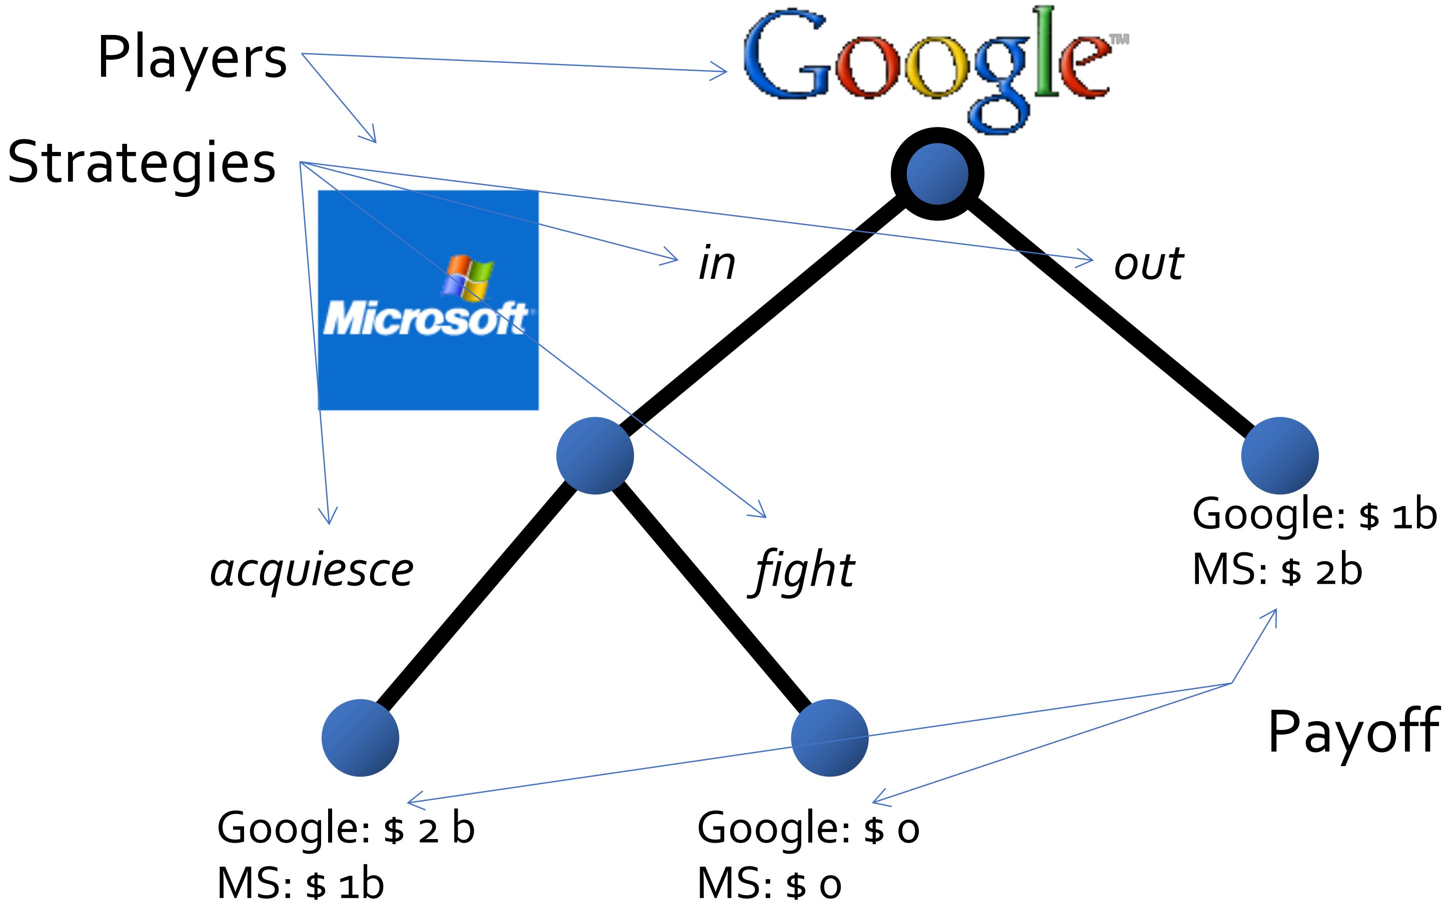

	Acquiesce	Fight
In	(2, 1)	(0, 0)
Out	(1, 2)	(1, 2)

Solution Concepts in Extensive Games

- There are some weird cases by directly extending Nash equilibrium.

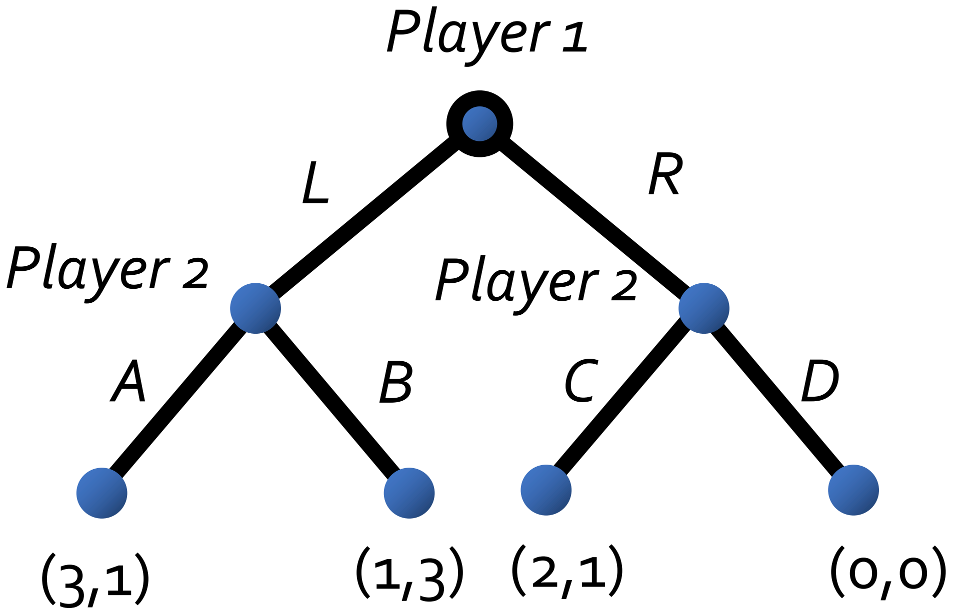

- Nash equilibrium:
  - Player 1 chooses L
  - Player 2 chooses (B, D)

Is it stable?

Why is this the weird case?

- Player 2 gives a “non-credible threat”
- I’ll choose D if you choose R

Solution Concepts in Extensive Games

- Subgame Perfect Equilibrium (SPE)
  - Play in each subgame is a Nash equilibrium.
  - Rule out the "non-credible threat"

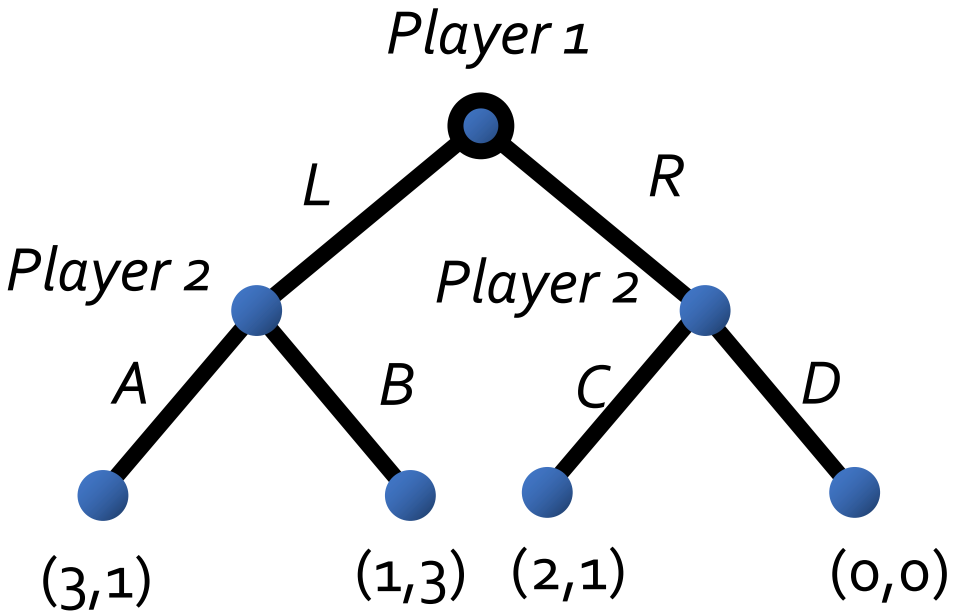

- SPE
  - Player 1 chooses R
  - Player 2 chooses (B, C)

- Can usually be calculated using backwards induction.

Applying Game Theory in this Course

- Formulate users' incentives.
- Describe the game structure
  - Sometime there are no interactions between workers. It becomes a simpler optimization problem.
- Analyze the equilibrium as the prediction of the outcome

- Mechanism design
  - Specify the desired outcome
  - Design incentive and game structures such that the outcome is the equilibrium

Practice Example – Contract Design

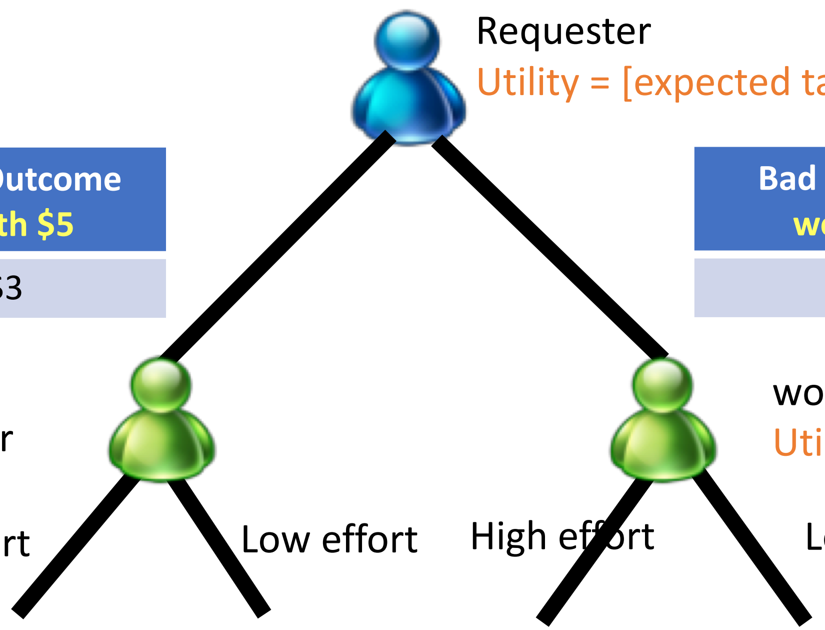

Bad Outcome worth $1	Good Outcome worth $5
$1	$3

Bad Outcome worth $1	Good Outcome worth $5
$1	$1

	Bad Outcome	Good Outcome	
High effort	Prob: 0.1	Prob: 0.9	Cost: $1
Low effort	Prob: 0.8	Prob: 0.2	Cost: $0

Which contract should the requester choose?

Can be generalized to any real-valued contracts (Principal Agent problem)

Discussion

- The main assumption we usually make in incentive design is that users are rational. While it approximates human behavior fine in some cases, this assumption has often been criticized.
- Questions
  - What are the examples of the “non-rational” human behavior you can think of?
  - Can you still try to model the non-rational behavior?

Example of Human Bias

- Present bias:
  - Humans value immediate payoffs heavier than future payoffs.
  - What will happen if we don’t require project milestone reports?
- Herding bias:
  - Humans tend to follow what others do/say.
  - Will you eat at a restaurant with a lot of people lining up or another with no one inside.
- Prospect theory:
  - Nobel-winning theory in explaining biases in decision-making with uncertainty events.
- And more…

- An growing research direction to account for human biased behavior in computational systems.
  - First workshop on Behavioral EC

Sep 25	Incentive Design: Financial Incentives Presenter: Christine, Gus, and Jeremy	Are workers really rational in financial incentives?
Sep 30	Incentive Design: Badges and Attention Presenter: CJ	Can we formulate non-financial incentives?
Oct 2	Application: Darpa Network Challenge Presenter: Alex, Junbin, and Pengqiu	Incentivizing users through social networks
Oct 7	Application: Prediction Markets Presenter: Ali, Davis, and Kristen	Using market mechanisms for eliciting users' beliefs

Proper Scoring Rules

Incentivizing Truthful Reports

- Example scenarios:
  - Ask a weather forecaster: will it rain tomorrow?
  - Ask a political researcher: will Trump win 2020 election?
  - Ask a Microsoft employer: will the new version of Office be shipped on time?

- Want to obtain forecasts about future events
- How do we make sure we obtain “truthful” reports?

Incentivizing Truthful Reports

- Setting
  - Consider a rational agent with linear utility for cash
  - Suppose there are n mutually exclusive and exhaustive states of the world $\Omega = \{w_1, w_2, \dots, w_n\}$ (e.g., Sun, Rain, Snow)
  - p_i is the subjective belief of the agent that state w_i will occur
- Question
  - How do we motivate this agent to tell us her beliefs about the likelihood of each state?

Scoring Rules

- A scoring rule rewards an agent $S(\vec{r}, w)$ when his reported distribution is $\vec{r}$ and the realized outcome is w

Scoring Rules

- Let's consider a linear scoring rule

$$S(\vec{r}, w_i) = r_i$$

- If a risk-neutral agent believes the probability for Rain, Sun, and Snow are

$$\vec{p} = (0.7, 0.2, 0.1)$$

What report should the agent provide?

Scoring Rules

- A scoring rule rewards an agent $S(\vec{r}, w)$ when his reported distribution is $\vec{r}$ and the realized outcome is w

- A scoring rule is called ***proper*** if the agent maximizes her utility by providing truthful report

$$\vec{p} = \mathrm{argmax}_{\vec{r}} \sum_{i=1}^{n} p_i S(\vec{r}, w_i)$$

- A scoring rule is ***strictly proper*** if honestly reporting is the **unique** maximizer.

Examples of Strictly Proper Scoring Rules

- Quadratic scoring rule (Brier score):

$$S(\vec{r}, w_i) = r_i - \frac{1}{2}\sum_j r_j^2$$

We can verify this by taking the gradient of the expected payoff

- Affine transformation of the proper scoring rule is still proper.

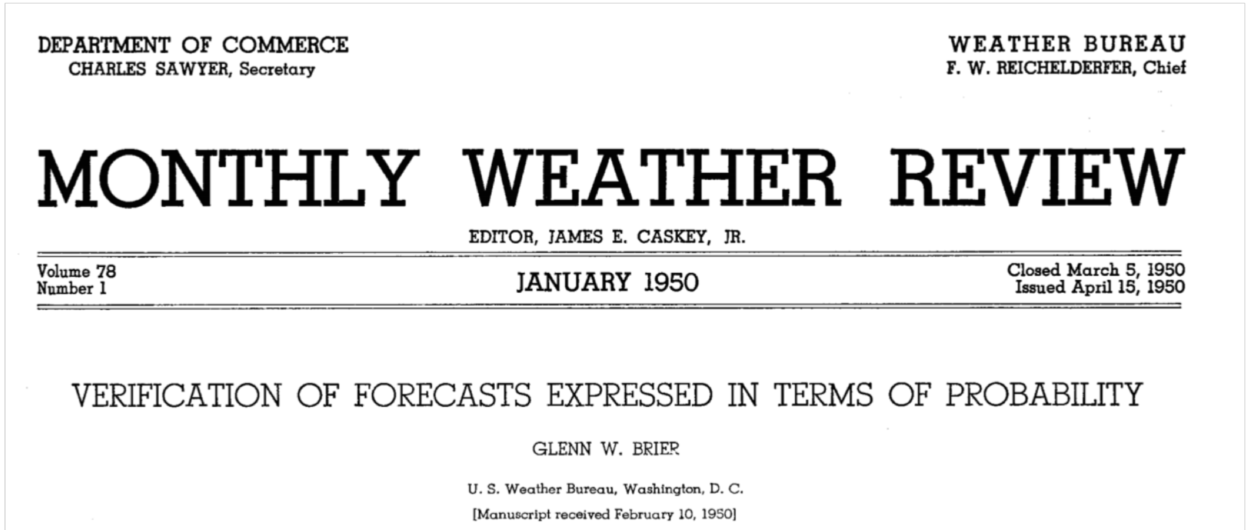
DEPARTMENT OF COMMERCE
CHARLES SAWYER, Secretary

WEATHER BUREAU
F. W. REICHELDERFER, Chief

MONTHLY WEATHER REVIEW

EDITOR, JAMES E. CASKEY, JR.

Volume 78
Number 1

JANUARY 1950

Closed March 5, 1950
Issued April 15, 1950

VERIFICATION OF FORECASTS EXPRESSED IN TERMS OF PROBABILITY

GLENN W. BRIER

U. S. Weather Bureau, Washington, D. C.

[Manuscript received February 10, 1950]

Examples of Strictly Proper Scoring Rules

- Logarithmic scoring rule:

$$S(\vec{r}, w_i) = \log r_i$$

We can verify this by taking a gradient of the expected payoff

- In logarithmic scoring rule, the score for outcome w_i only depends on the report r_i and not r_j for $j \neq i$

More examples?

- How do we construct a strictly proper scoring rule?

- How many strictly proper scoring rules are there?

Characterization of Proper Scoring Rules

- Connections between convex functions and proper scoring rules.

- A scoring rule $S(\vec{r}, w_i)$ is (strictly) proper if and only if

$$S(\vec{r}, w_i) = G(\vec{r}) - \sum_{j \neq i} G'_j(\vec{r}) p_j + G'_i(\vec{r})$$

where $G(\vec{r})$ is a (strictly) convex function, $G'(\vec{r})$ is a subgradient of G at $\vec{r}$, and $G'_i(\vec{r})$ is its i-th component.

Section 3.1 of Strictly Proper Scoring Rules, Prediction, and Estimation by Gneiting and Raftery.

Connection to Prediction Market

Who will win the 2020 Democratic presidential nomination?

	Contract	Latest Yes Price		Best Offer			Best Offer
	Elizabeth Warren	45¢	NC	46¢	Buy Yes	Buy No	55¢
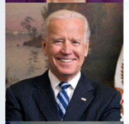	Joe Biden	26¢	2¢↓	27¢	Buy Yes	Buy No	74¢
	Bernie Sanders	10¢	1¢↓	11¢	Buy Yes	Buy No	90¢
	Andrew Yang	9¢	1¢↓	9¢	Buy Yes	Buy No	92¢
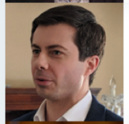	Pete Buttigieg	6¢	1¢↑	7¢	Buy Yes	Buy No	94¢
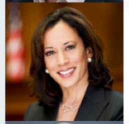	Kamala Harris	4¢	1¢↓	5¢	Buy Yes	Buy No	96¢
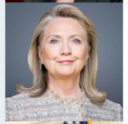	Hillary Clinton	4¢	2¢↓	5¢	Buy Yes	Buy No	96¢
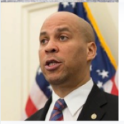	Cory Booker	2¢	1¢↑	2¢	Buy Yes	Buy No	99¢

24 More Contracts

Connection to Prediction Markets

Goal:

Incentivize ***multiple*** agents to share their beliefs, and find a way to ***aggregate*** these beliefs into an unified prediction

1. Could use one scoring rule per agent, but not clear how to aggregate
2. Market itself is an aggregation mechanism (use final price as the prediction). However, standard stock-market-style trading might encounter issues for less popular predictions (market is too *thin*).

Market Scoring Rules

- See Hanson's papers in the optional readings of the Prediction Market lecture
- Intuitions: a "sequentially shared scoring rule"
  - An automatic market maker
  - Market maintains a vector of predictions $\vec{r}^{(t)}$
  - If a trader changes the vector from $\vec{r}^{(t)}$ to $\vec{r}^{(t+1)}$ and the outcome is w_i, the trader obtains reward

$$S\left(\vec{r}^{(t+1)}, w_i\right) - S\left(\vec{r}^{(t)}, w_i\right)$$

- Under some conditions:
  - Agents truthfully report their beliefs
  - The prediction will converge

Market Scoring Rules

- The connection to convex optimization opens up an interesting line of research in the design of efficient market maker…

Oct 7	Application: Prediction Markets Presenter: Ali, Davis, and Kristen	**Required** Results from a Dozen Years of Election Futures Markets Research. Berg et al. 2001. You should first read the first page of the following two-page (not too technical) article for a brief explanation of prediction markets. The Promise of Prediction Markets, K.J. Arrow et. al., Science. 2008 **Optional** Empirical reports: Using Prediction Markets to Track Information Flows: Evidence from Googl Cowgill, Wolfers, and Zitwewitz. 2008. Prediction Markets. Wolfers and Zitzewitz, The Journal of Economic Perspectives 2004. Automatic market makers Logarithmic Market Scoring Rules for Modular Combinatorial Information Aggregation. Hanson. Journal of Prediction Markets 2007. Combinatorial Information Market Design. Hanson. Information Systems Frontier 2003. A blogpost by David Pennock that discuss how to implement market scorin rules as a market maker.

- We won't cover too much on prediction markets. In case you are interested, below are a few more papers to follow up:
  - A New Understanding of Prediction Markets Via No-Regret Learning. Chen and Vaughan. EC 2010.
  - An Optimization-Based Framework for Automated Market-Making. Abernethy, Chen, and Vaughan. EC 2011.
  - and more (the papers authored on these authors)

Very Brief Intro of Peer Prediction

Eliciting Truthful Reports

- Scoring rule relies on the “truth” to be revealed in the future
- What if there is no ground truth (or the ground truth is hard to obtain)
  - Do you like this movie?
  - Peer grading in MOOCs

- Output agreements:
  - Randomly pick two persons
  - If their reports match, reward them 1, otherwise reward 0
  - Truthful reporting is not an equilibrium (you are encouraged to report the majority’s opinion)

Peer Prediction

- How to fix the issue?
  - Assume knowledge about the report distribution, re-weighting the rewards to make sure truthful reporting is a equilibrium

- Drawbacks:
  - Require knowledge of the prior
  - There are usually multiple equilibrium (including naïve bad ones…)

- Still an ongoing research area
  - Some nice theoretical results, however there is little practical success so far

Eliciting Informative Feedback: The Peer-Prediction Method. Miller et al. 2004

Related Course

- The course on peer grading at Northwestern by Jason Hartine.
  - https://sites.northwestern.edu/hartline/eecs-497-peer-grading/