

CSE 417T (Machine Learning): Exam 2 Practice Questions

1. Consider the following dataset:

○	+
+	○

Which of the following classifiers will achieve zero training error on this data set. (There might be more than one.)

- (a) Logistic regression
 - (b) SVM (polynomial kernel)
 - (c) Depth-2 ID3 decision trees
 - (d) 3-NN classifier
 - (e) Neural networks with a single hidden layer
2. Assume we have a decision tree to classify binary vectors with 100 features. Can you specify a 1-Nearest Neighbor model that would result in exactly the same classification of all examples as the decision tree? Explain how, or provide a counter-example.
 3. Suppose you are using depth 1 decision trees (so you can split on one feature) as the weak learners in AdaBoost, but you know that all your features $x_i \in \{\pm 1\}$ (that is, you're solving a Boolean problem, and all features, as well as the class, are either true or false). Show that the final hypothesis output by AdaBoost after T rounds of boosting is a linear classifier.
 4. Using a training data set, you compute a (1-)nearest neighbor classifier and apply it to a held out test data set. You find that this hypothesis achieves a test error of E_{test} . You know that the nearest neighbor model suffers from high variance so you decide to apply bagging. However, you observe that when the number of bags is very large, the test error of the aggregated hypothesis is also exactly E_{test} . Explain why this occurs.
 5. Consider the following classification function. For each point (\vec{x}, y) , where \vec{x} is represented by x_1 and x_2 . The label y is determined by

$$y = \begin{cases} 1 & \text{if } x_2 \geq |x_1| \\ -1 & \text{otherwise} \end{cases}$$

Can we construct a neural network with one hidden layer and with *sign* function as the active function to represent the above function? If yes, draw it. If not, explain why not.

6. Suppose you learn a hard-margin SVM using a training data set with N data points. You observe that there are M support vectors. Explain why the following bound on the leave-one-out cross validation error must be true: $E_{LOOCV} \leq M/N$
7. For linearly separable data, can a small slack penalty term (i.e., C) in soft-margin SVM (without transformations) hurt the training error? Explain your answers.
8. Suppose you apply bagging and boosting to linear classifiers. For which of the two will the output ensemble method also be a linear classifier?
 - ☐ Boosting only.
 - ☐ Bagging only.
 - ☐ Both boosting and bagging.
 - ☐ Neither boosting nor bagging.
9. Since the VC dimension of SVM with Gaussian RBF Kernel is infinite, such an SVM will have worse out-of-sample error than an SVM with polynomial kernel with finite VC dimension.
 - ☐ True
 - ☐ False
10. Consider a data point that is distant from the decision boundary. For logistic regression and SVMs, which ones would be affected by a small perturbation to this data point, assuming it still remains far from the boundary?
 - ☐ Logistic regression only.
 - ☐ SVM only.
 - ☐ Both logistic regression and SVM.
 - ☐ Neither logistic regression nor SVM.
11. What is one reason that you might prefer a single decision tree over a random forest for a particular problem?
 - ☐ Decision trees are more interpretable than random forests
 - ☐ Decision trees are less likely to overfit than random forests
 - ☐ Decision trees have lower variance than random forests
 - ☐ None of the above
12. Which of the following techniques help prevent Neural Networks from overfitting?
 - A. Early stopping
 - B. Dropout
 - C. Data augmentation (adding noises to existing data points)
 - D. All of the above