

CSE 417T

# Introduction to Machine Learning

Lecture 4

Instructor: Chien-Ju (CJ) Ho

# Logistics: Homework 1

- Due: **Feb 14 (Monday), 2022**
  - <http://chienjuho.com/courses/cse417t/hw1.pdf>
  - Intended deadline: Feb 10.
    - Recommend to work on it early to spare time for homework 2
- Two submission links: Report and Code
  - Report: Answer all questions, including the implementation question
    - **Grades are based on the report**
  - Code: Complete and submit **hw1.py** for Problem 2
    - The code will only be used for correctness checking (when in doubts) and plagiarism checking
- Reserve time if you never used Gradescope.
  - Make sure to **specify the pages for each problem**. You **won't get points** otherwise

# Logistics: Office Hours

- Tentative schedule of TA office hours (starting next Monday)

Monday	11:30am (Herbert Zhou)	4pm (Dean Yu)	
Tuesday	1pm (Ziqi Xu)	3:30pm (Neal Huang)	
Wednesday	1pm (Eddie Choi)	4:30pm (Weiwei Ma)	
Thursday	10am (Jackie Zhong)	3pm (Fankun Zeng)	
Friday	8am (Shohaib Shaffiey)	1pm (Yunfan Wang)	7pm (Hao Qin)
Sunday	1pm (Jonathan Ma)		

- 60 minutes per session
- Please follow **Piazza** for additional information
- Recommendation: Try to utilize the office hour early (way ahead of deadlines), you are likely to get more of TAs' time this way

Recap

# Common Notations

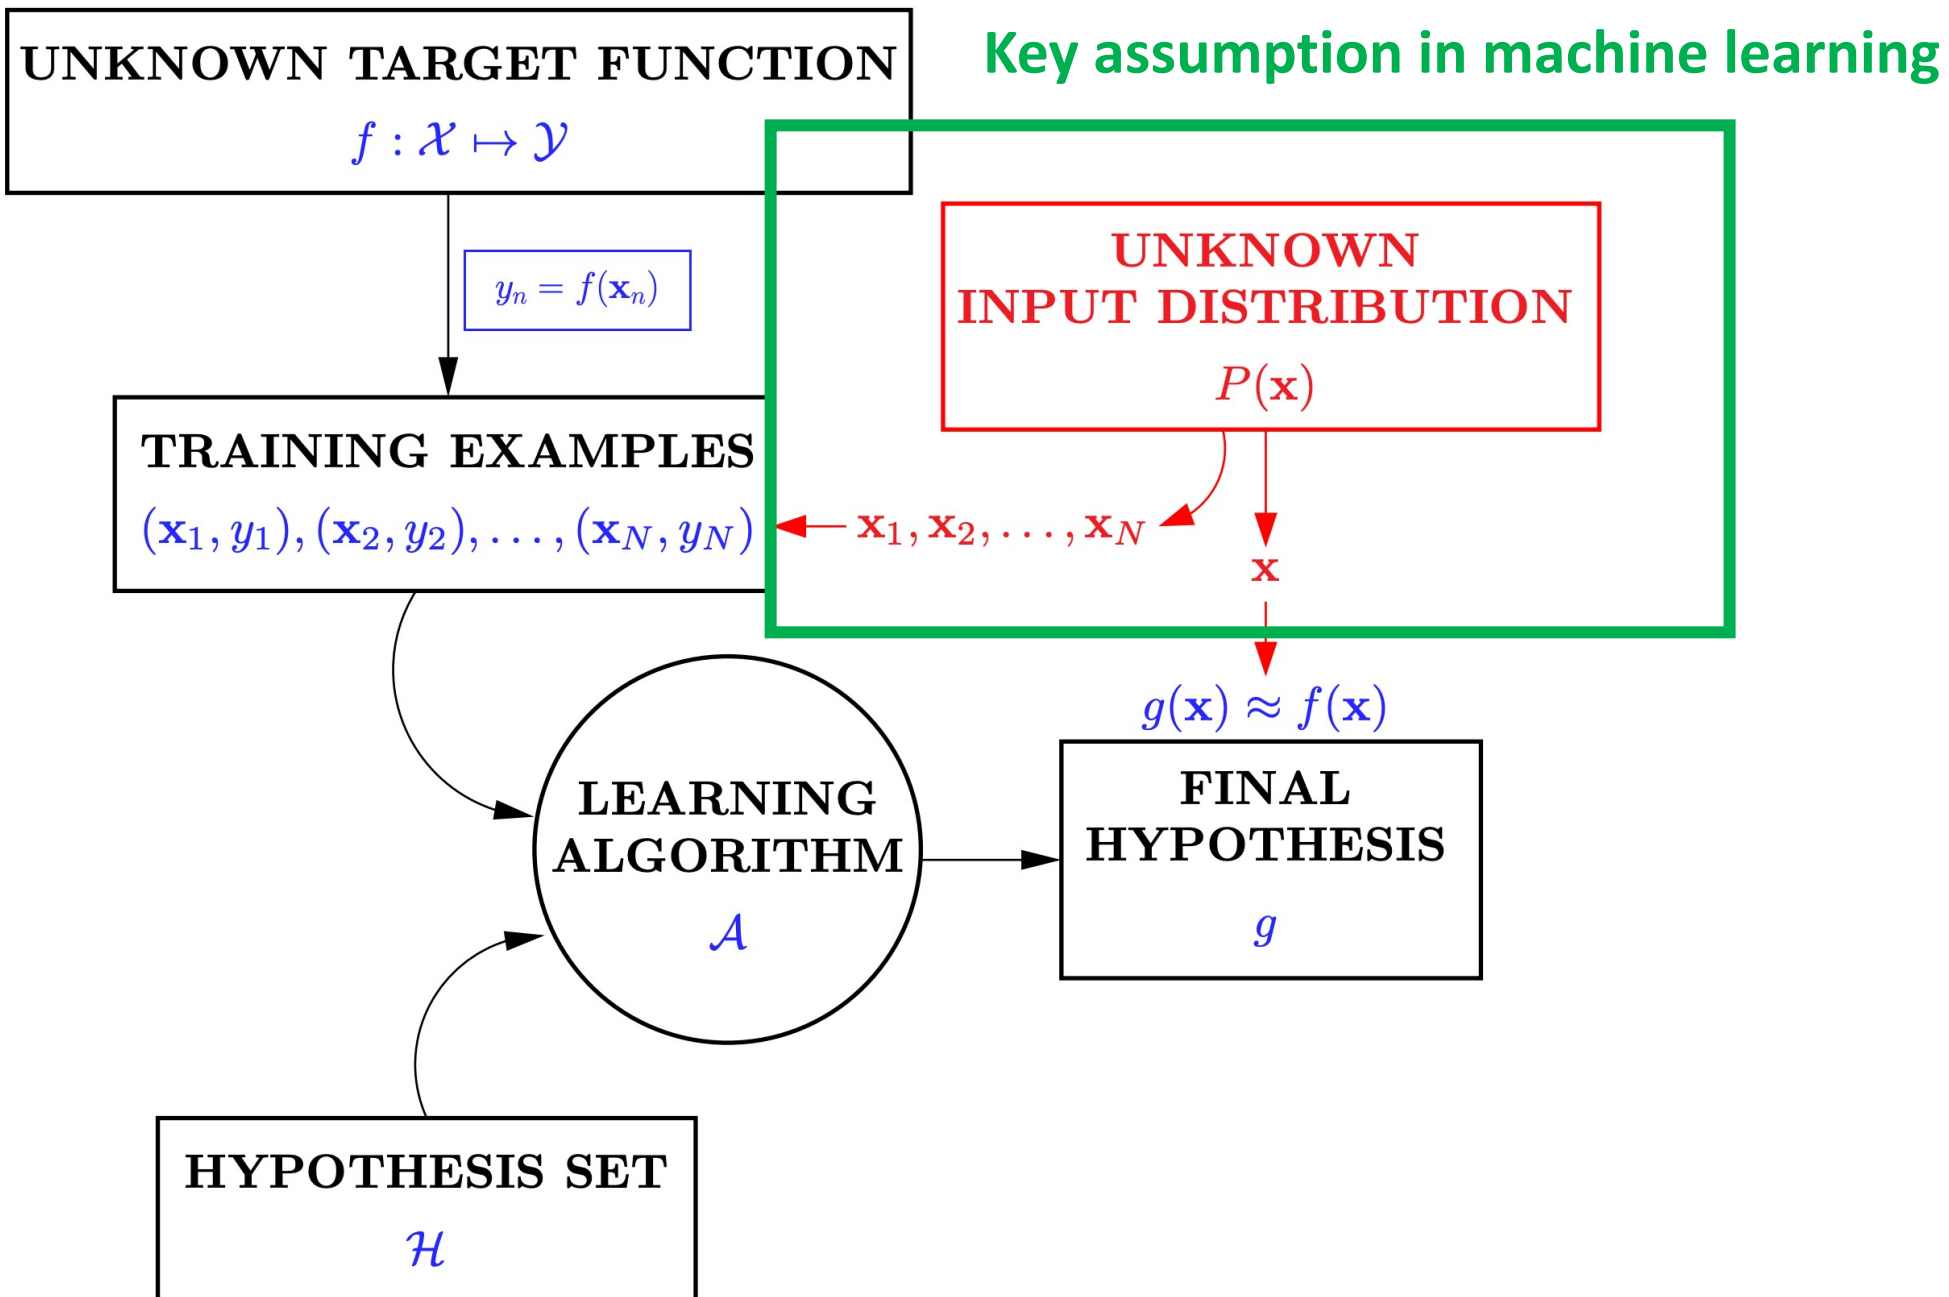
Note that by default,  $\vec{x}$  is a **column** vector.

More formally, we should write  $\vec{x} = \begin{bmatrix} x_0 \\ \vdots \\ x_d \end{bmatrix}$ .

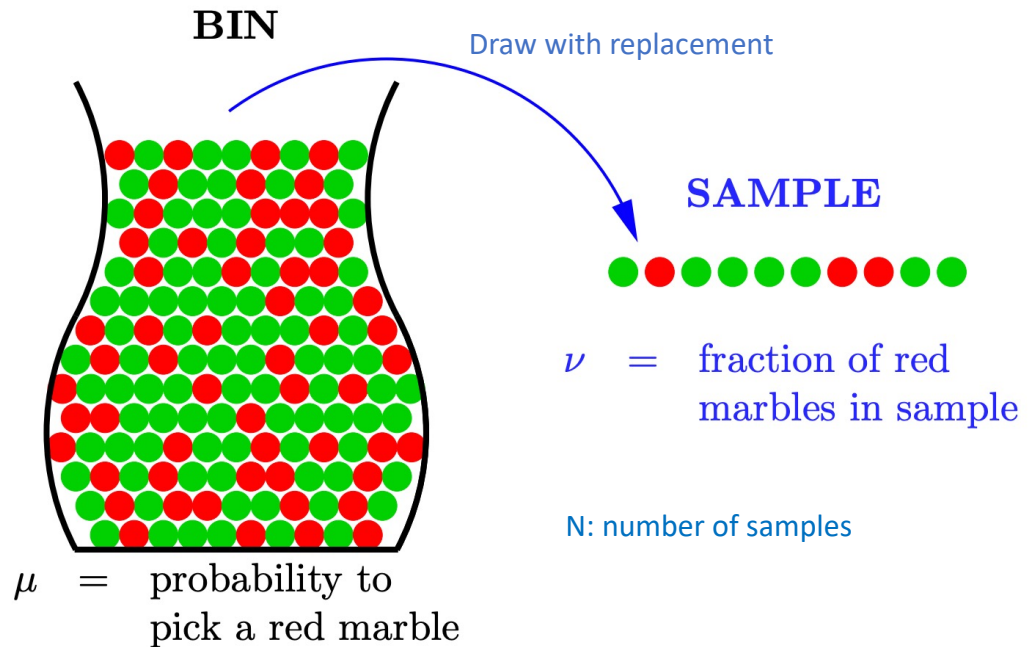
For convenience, I usually write  $\vec{x} = (x_0, \dots, x_d)$ .

- Data point with augmented  $x_0$ :  $\vec{x} = (x_0, \dots, x_d)$ 
  - We often use  $d$  to specify the dimensions of data points
  - We augment  $x_0 = 1$  for each data point (Check Lecture 1 for the reasoning)
- Dataset:  $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$ 
  - We often use  $N$  to specify the number of data points in the dataset
- Hypothesis set  $H$ 
  - We use  $h \in H$  to specify an arbitrary hypothesis
  - We use  $g \in H$  to specify the hypothesis output by the learning algorithm
- Indicator variable:
  - $\mathbb{I}[\text{event}] = \begin{cases} 1 & \text{if event is true} \\ 0 & \text{if event is false} \end{cases}$

Example:  $\mathbb{I}[h(\vec{x}) \neq f(\vec{x})] = \begin{cases} 1 & \text{if } h(\vec{x}) \neq f(\vec{x}) \\ 0 & \text{if } h(\vec{x}) = f(\vec{x}) \end{cases}$



# Hoeffding's Inequality



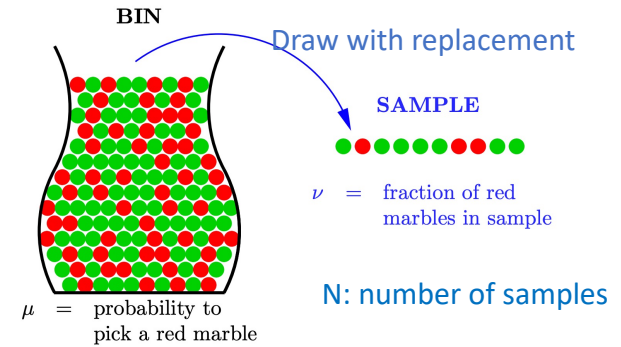
$$\Pr[|\mu - \nu| > \epsilon] \leq 2e^{-2\epsilon^2 N}$$

Define  $\delta = \Pr[|\mu - \nu| > \epsilon]$

- Fix  $\delta$ ,  $\epsilon$  decreases as  $N$  increases
- Fix  $\epsilon$ ,  $\delta$  decreases as  $N$  increases
- Fix  $N$ ,  $\delta$  decreases as  $\epsilon$  increases

Informal intuitions of notations  
 $N$ : # sample  
 $\delta$ : probability of "bad" event  
 $\epsilon$ : error of estimation

# Connection to Learning



- Given dataset  $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$ .

- Fix a hypothesis  $h$

- $E_{in}(h) \stackrel{\text{def}}{=} \frac{1}{N} \sum_{n=1}^N \mathbb{I}[h(\vec{x}_n) \neq f(\vec{x}_n)]$  [In-sample error, analogy to  $\nu$ ]

- $E_{out}(h) \stackrel{\text{def}}{=} \Pr_{\vec{x} \sim P(\vec{x})} [h(\vec{x}) \neq f(\vec{x})]$  [Out-of-sample error, analogy to  $\mu$ ]

- Apply Hoeffding's inequality

$$\Pr[|E_{out}(h) - E_{in}(h)| > \epsilon] \leq 2e^{-2\epsilon^2 N}$$

- This is *verification*, not *learning*



# Connection to “Real” Learning

- Given a **finite** hypothesis set  $H = \{h_1, \dots, h_M\}$
- Apply some learning algorithm on  $D$ , output a  $g \in H$
- What can we say about  $E_{out}(g)$  from  $E_{in}(g)$ ?

$$\Pr[|E_{out}(g) - E_{in}(g)| > \epsilon] \leq 2Me^{-2\epsilon^2 N} \quad \text{for any } \epsilon > 0$$

## Intuitions:

1. Bad event  $B(g) \subseteq B(h_1) \cup B(h_2) \dots \cup B(h_M)$

$g$  is selected within  $\{h_1, \dots, h_M\}$

=> bad event of  $g$  is within the union of the bad events of  $h_1, \dots, h_M$

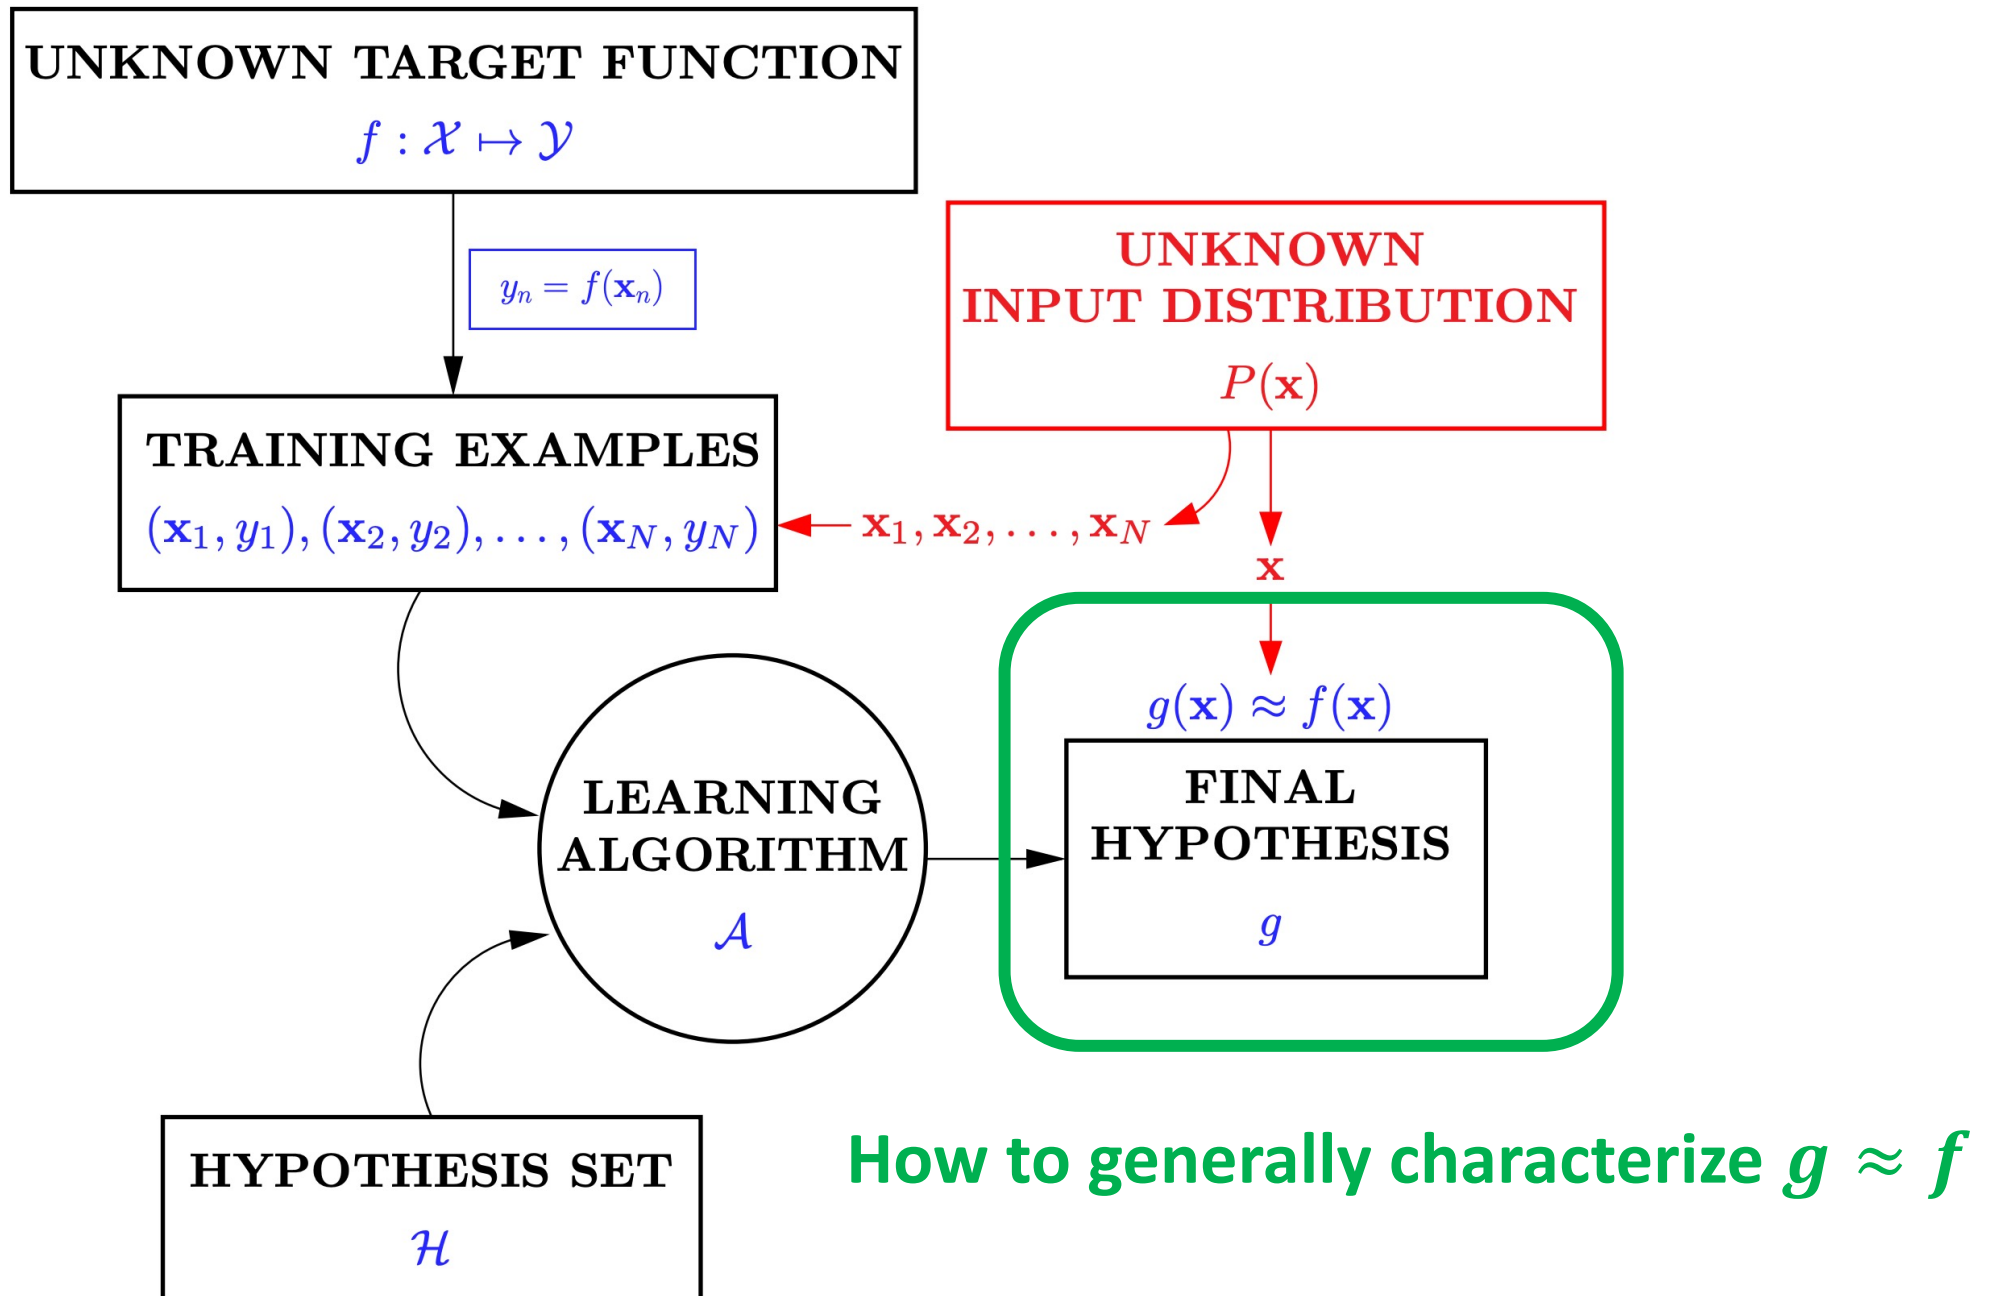
2.  $\Pr[B(g)] \leq \Pr[B(h_1)] + \dots + \Pr[B(h_M)]$

each of the  $\Pr[B(h_m)]$  follows Hoeffding's inequality

# Today's Lecture

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.  
Let me know if you spot errors.

Revisit the learning problem



# Goal: $g \approx f$

- A general approach:
  - Define an error function  $E(h, f)$  that quantify how far away  $h$  is to  $f$
  - choose  $g = \underset{h \in \mathcal{H}}{\operatorname{argmin}} E(h, f)$
- A major component of ML is **optimization**
- $E$  is usually defined in terms of a **pointwise** error function  $e(h(\vec{x}), f(\vec{x}))$ 
  - Binary error (classification):  $e(h(\vec{x}), f(\vec{x})) = \mathbb{I}[h(\vec{x}_n) \neq f(\vec{x}_n)]$
  - Squared error (regression):  $e(h(\vec{x}), f(\vec{x})) = (f(\vec{x}) - h(\vec{x}))^2$

$$E_{in}(h) = \frac{1}{N} \sum_{n=1}^N e(h(\vec{x}_n), f(\vec{x}_n))$$
$$E_{out}(h) = \mathbb{E}_{\vec{x}}[e(h(\vec{x}), f(\vec{x}))]$$

The discussion on the Hoeffding's inequality applies for general (bounded) error functions.

# How to choose the error function?

- Consideration 1: Properties of domain applications
- Example: Fingerprint recognition
  - Input: fingerprints
  - Outputs: whether the person is authorized

		$f(\vec{x})$	
		+1	-1
$h(\vec{x})$	+1	No error	False positive
	-1	False negative	No error

		$f(\vec{x})$	
		+1	-1
$h(\vec{x})$	+1	0	Small
	-1	Large	0

		$f(\vec{x})$	
		+1	-1
$h(\vec{x})$	+1	0	Large
	-1	Small	0

# How to choose the error function?

- Consideration 1: Properties of application problems
- Consideration 2: Computation
  - ML algorithms are essentially performing **optimization** (finding  $g$  with smallest error)

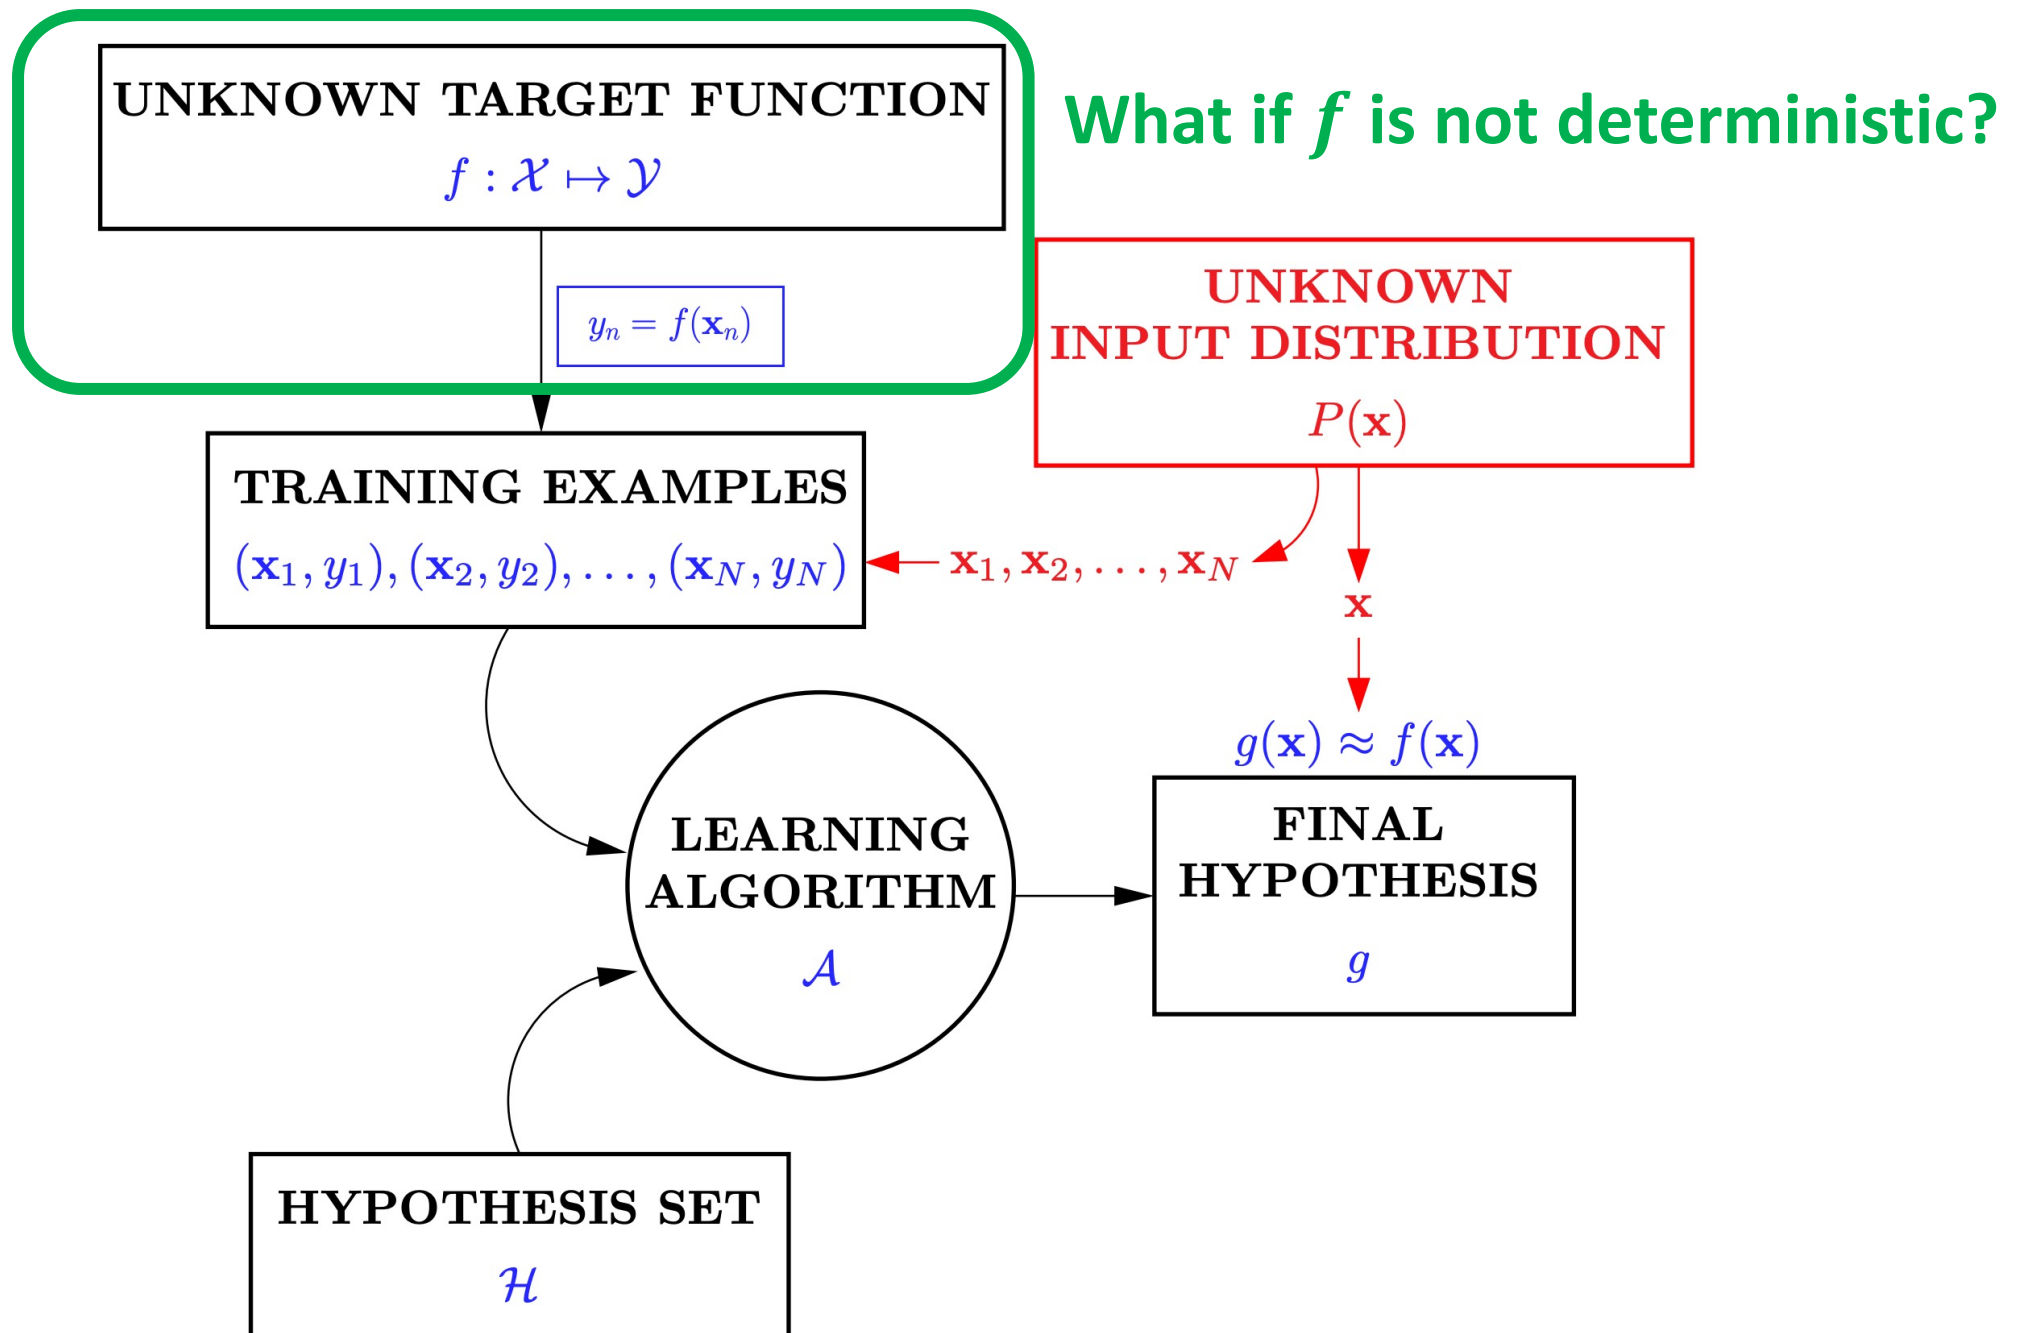
$$g = \operatorname{argmin}_{h \in \mathcal{H}} E(h, f)$$

- Choose the error that is “easier” to optimize
  - e.g., if the error function is continuous, differentiable, and convex, we usually have efficient algorithms

# How to choose the error function?

- Consideration 1: Properties of application problems
- Consideration 2: Computation
- Specifying the error function is part of setting up the learning problem
  - It impacts what you eventually learn



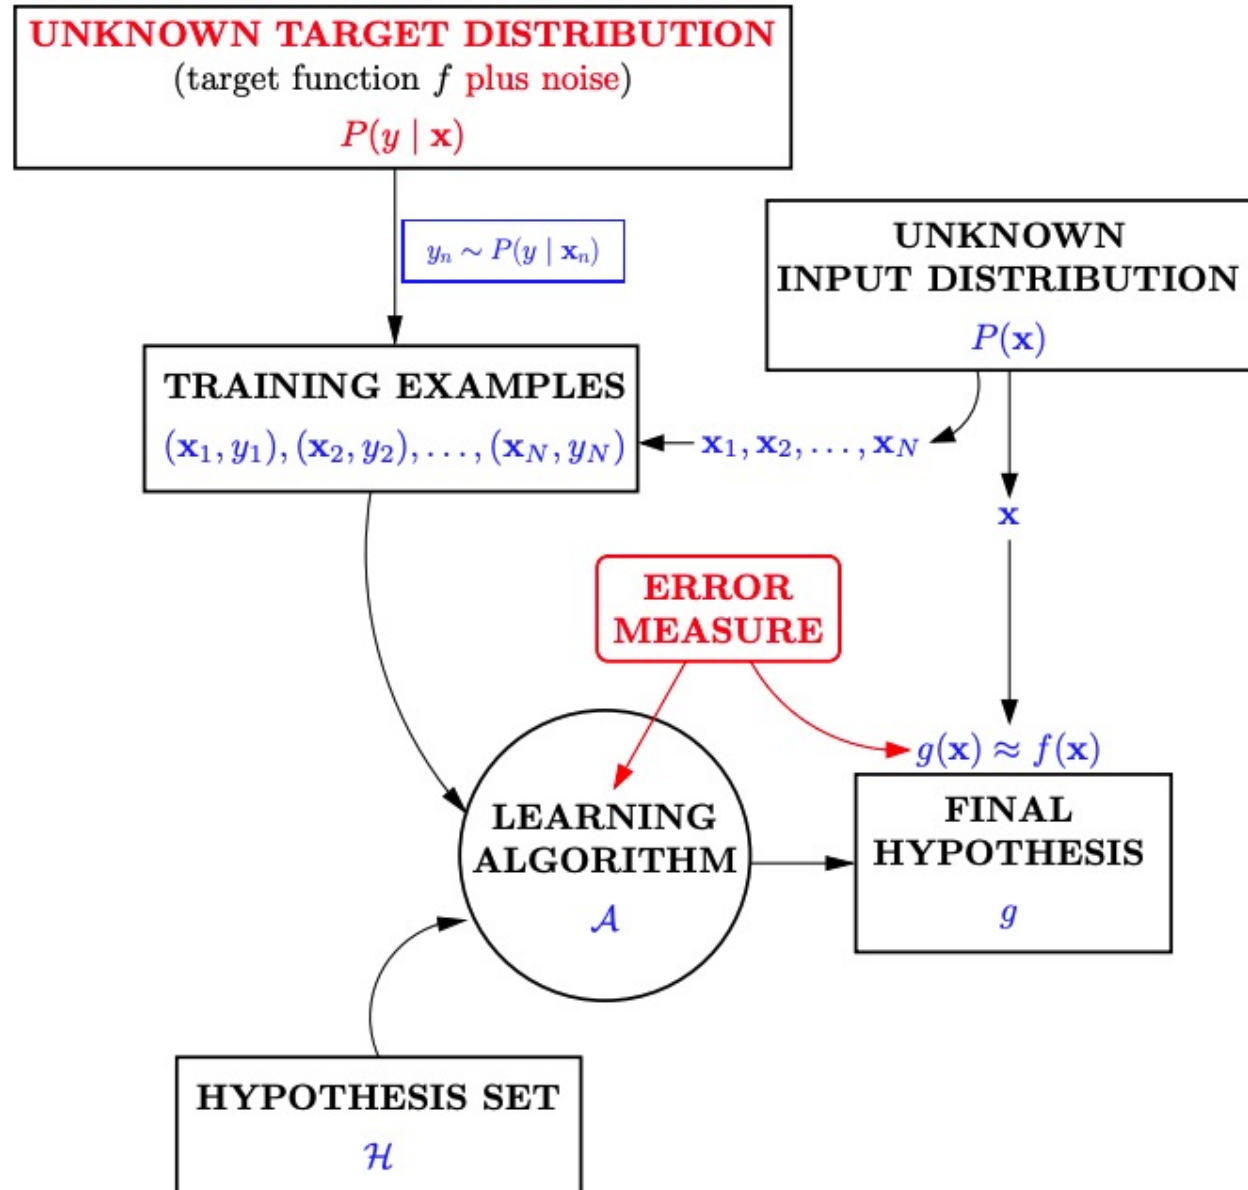


# Noisy Target

- What if there doesn't exist  $f$  such that  $y = f(\vec{x})$ ?
  - $f$  is stochastic instead of deterministic
  - (even if two customers have exactly the same attributes, one might be a good customer for bank, and the other might not be)
- Common approach
  - Instead of a target function, define a target **distribution**
  - Instead of  $y = f(\vec{x})$ ,  $y$  is drawn from a conditional distribution  $P(y|\vec{x})$
  - $y = f(\vec{x}) + \epsilon$ 
    - $f(\vec{x})$  is the mean of the distribution  $\mathbb{E}[y|\vec{x}]$
    - $\epsilon$  is zero-mean noise  $y - \mathbb{E}[y|\vec{x}]$

The discussion on the Hoeffding's inequality applies for noisy targets.

# General Setup of (Supervised) Learning



# Theory of Generalization

# Revisit the “Multi-Hypothesis” Bound

- Given a **finite** hypothesis set  $H = \{h_1, \dots, h_M\}$
- Apply some learning algorithm on  $D$ , output a  $g \in H$
- What can we say about  $E_{out}(g)$  from  $E_{in}(g)$ ?

$$\Pr[|E_{out}(g) - E_{in}(g)| > \epsilon] \leq 2Me^{-2\epsilon^2 N} \quad \text{for any } \epsilon > 0$$

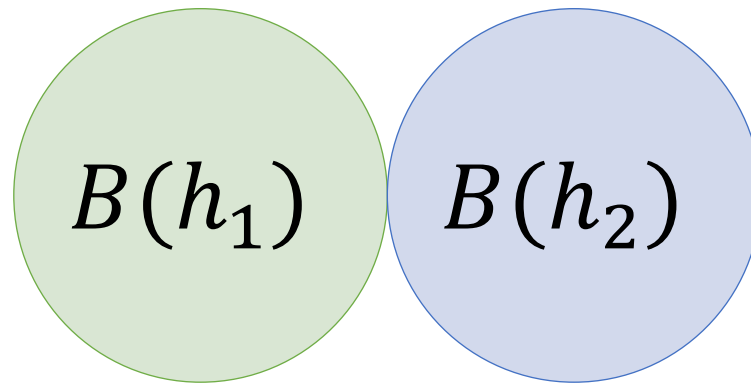
What if  $M$  is infinite?

$Pr[|E_{out}(g) - E_{in}(g)| > \epsilon] \leq 2Me^{-2\epsilon^2 N}$  don't seem to carry any meanings

# Key Intuitions in the Multi-Hypothesis Analysis

- Define "bad event of  $h$ "  $B(h)$  as  $|E_{out}(h) - E_{in}(h)| > \epsilon$
- If  $g$  is selected from  $\{h_1, h_2\}$ 
  - $B(g) \subseteq B(h_1) \cup B(h_2)$
  - $\Pr[B(g)] \leq \Pr[B(h_1) \text{ or } B(h_2)]$

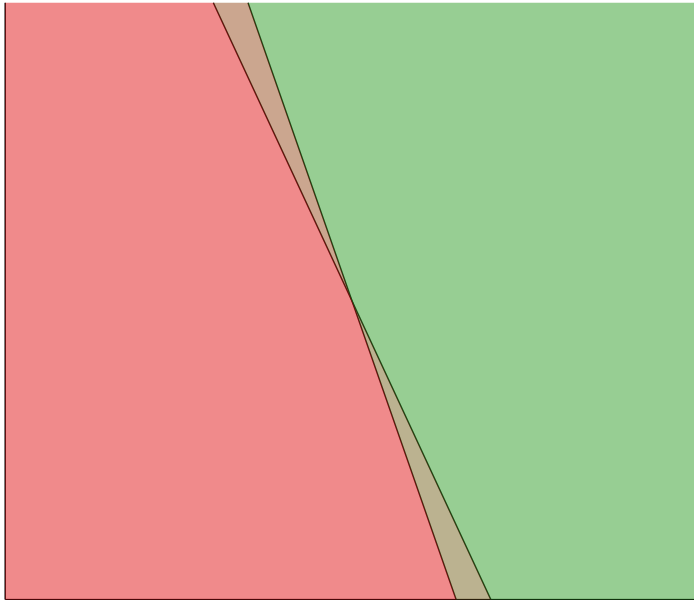
$$\leq \Pr[B(h_1)] + \Pr[B(h_2)] \quad (\text{Union Bound})$$



- Union bound considers the **worst case: Bad events don't overlap**

# Do Bad Events Overlap?

- Oftentimes, they overlap a lot!



The two linear separators on the left make the same predictions for most points.

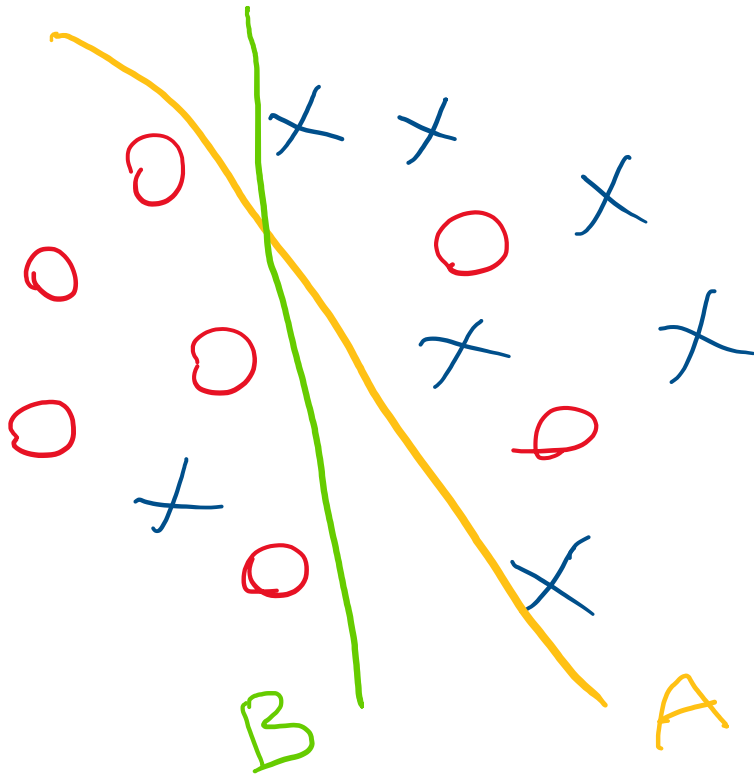
If it's a bad event for one, it's likely to be a bad event for the other.

$$\text{"bad event of } h\text{" } B(h): |E_{out}(h) - E_{in}(h)| > \epsilon$$

Recall: Informally, you can interpret “bad event of  $h$ ” as the event that we draw a “unrepresentative dataset  $D$ ” that makes the in-sample errors of  $h$  to be far away from out-of-sample error of  $h$



# What Can We Do?




For this dataset,  
any difference between **A** and **B**?

For this dataset, probably not.

They make the same predictions for  
every data point in this dataset.

# What Can We Do?

- Let's define “data-dependent” hypothesis, call it **dichotomy**.

 di·chot·o·my  
*/dī'kädəmə/*  
*noun*  
a division or contrast between two things that are or are represented as being opposed or entirely different.  
"a rigid **dichotomy** between science and mysticism"

- A hypothesis  $h: X \rightarrow \{-1, +1\}$
- A dichotomy for a set of data points  $(\vec{x}_1, \dots, \vec{x}_N)$ :
  - Assign either **+1** or **-1** for each of the data points  
(divide the data points into two groups)
- Why dichotomies?
  - It helps us count “effective number of hypothesis” (to replace  $M$ )

# More Formal Definitions

- Dichotomies

- Informally, consider a dichotomy as a “data-dependent” hypothesis
- Characterized by both hypothesis set  $H$  and  $N$  data points  $(\vec{x}_1, \dots, \vec{x}_N)$

$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$

- The set of possible prediction combinations  $h \in H$  can induce on  $\vec{x}_1, \dots, \vec{x}_N$

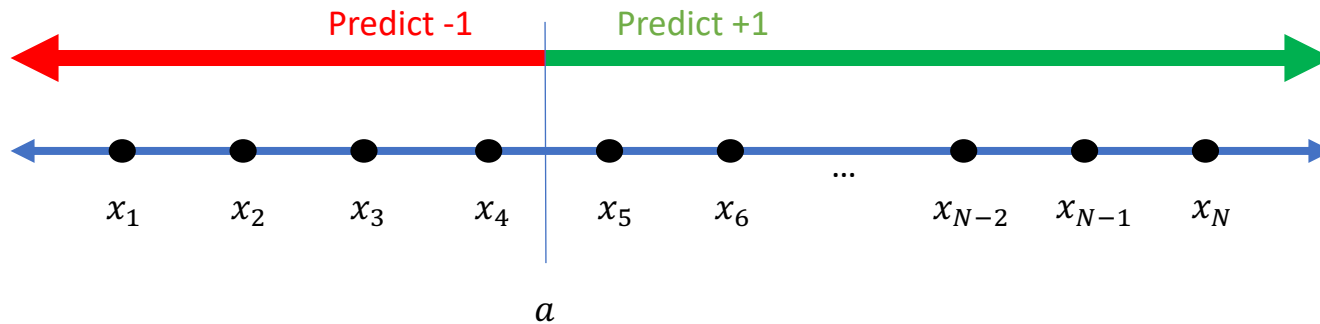
- Growth function

- Largest number of dichotomies  $H$  can induce across all possible data sets of size  $N$

$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

# Example: $H$ = Positive Rays

- Data points are in one-dimensional space
- Positive rays:  $h(x) = \text{sign}(x - a)$



- What is  $H(\vec{x}_1, \dots, \vec{x}_N)$ ?

- What is  $m_H(N)$ ?

## • Dichotomies

- Informally, consider a dichotomy as a “data-dependent” hypothesis
- Characterized by both hypothesis set  $H$  and  $N$  data points  $(\vec{x}_1, \dots, \vec{x}_N)$ 
$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$
- The set of possible prediction combinations  $h \in H$  can induce on  $\vec{x}_1, \dots, \vec{x}_N$

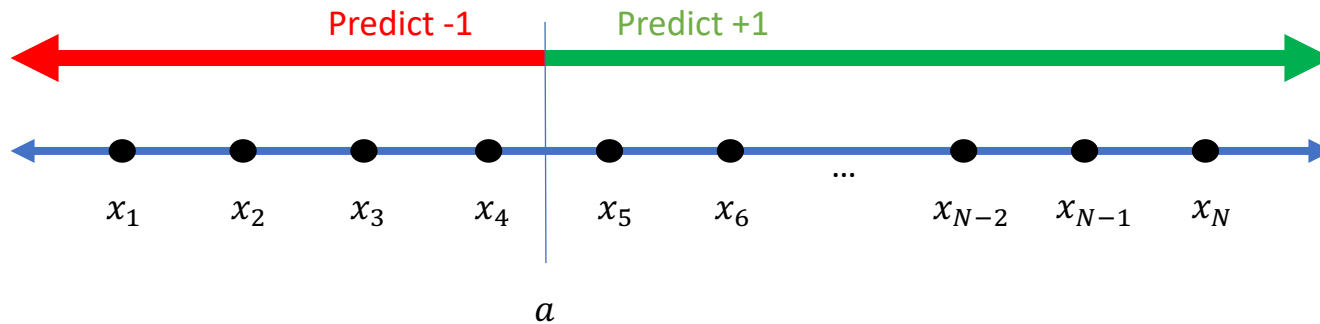
## • Growth function

- Largest number of dichotomies  $H$  can induce across all possible data sets of size  $N$

$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

# Example: $H$ = Positive Rays

- Data points are in one-dimensional space
- Positive rays:  $h(x) = \text{sign}(x - a)$



- What is  $H(\vec{x}_1, \dots, \vec{x}_N)$ ?

$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(+1, +1, \dots, +1), \\ (-1, +1, \dots, +1), \\ \dots \\ (-1, -1, \dots, -1)\}$$

- What is  $m_H(N)$ ?

$$m_H(N) = N + 1$$

## • Dichotomies

- Informally, consider a dichotomy as a “data-dependent” hypothesis
- Characterized by both hypothesis set  $H$  and  $N$  data points  $(\vec{x}_1, \dots, \vec{x}_N)$ 
$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$
- The set of possible prediction combinations  $h \in H$  can induce on  $\vec{x}_1, \dots, \vec{x}_N$

## • Growth function

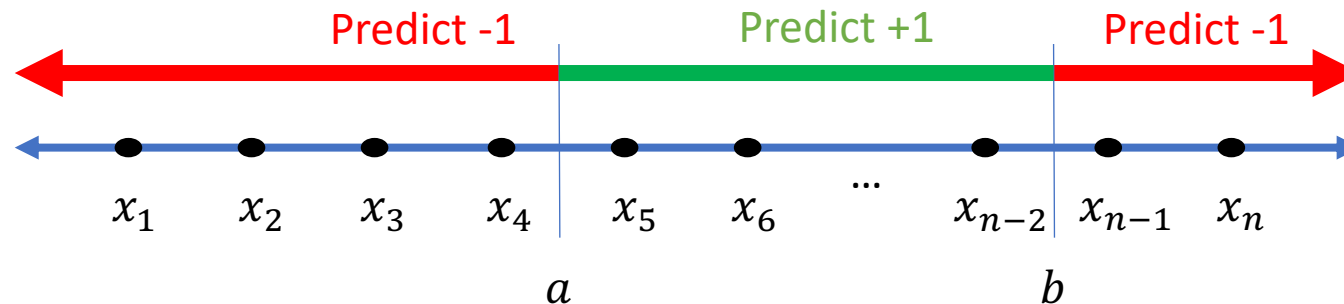
- Largest number of dichotomies  $H$  can induce across all possible data sets of size  $N$

$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

# What is $m_H(N)$ ?

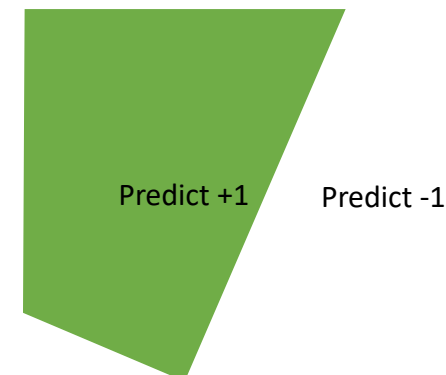
- $H$  = Positive Intervals

- Data points are in one-dimensional space
- Choose two thresholds. Predict +1 within the interval, -1 outside



- $H$  = Convex Sets

- Data points are in 2-dimensional space
- Hypothesis is represented by a convex set



- Dichotomies

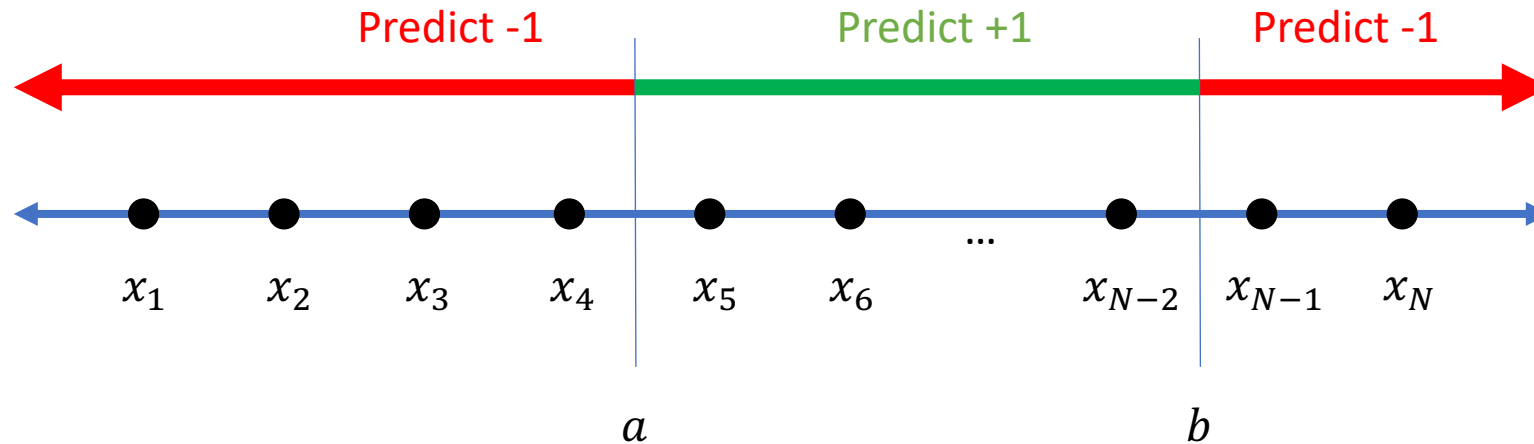
- Informally, consider a dichotomy as a “data-dependent” hypothesis
- Characterized by both hypothesis set  $H$  and  $N$  data points  $(\vec{x}_1, \dots, \vec{x}_N)$ 
$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$
- The set of possible prediction combinations  $h \in H$  can induce on  $\vec{x}_1, \dots, \vec{x}_N$

- Growth function

- Largest number of dichotomies  $H$  can induce across all possible data sets of size  $N$

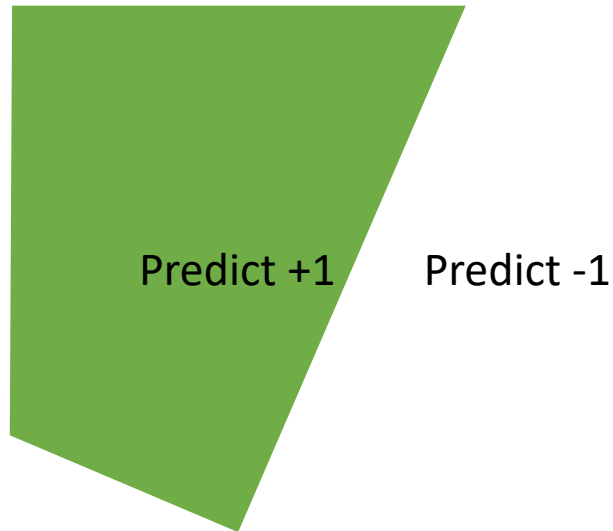
$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

# Example: $H$ = Positive Intervals



- What is  $m_H(N)$ ?
  - $m_H(N) = \binom{N+1}{2} + 1 = \frac{N^2}{2} + \frac{N}{2} + 1$

# Example: $H$ = Convex Sets



- What is  $m_H(N)$ ?
  - $m_H(N) = 2^N$

Note:

$m_H(N) \leq 2^N$  for all  $H$  and all  $N$   
(There are only  $2^N$  possible label combinations for  $N$  points)



# Why Growth Function?

- Growth function  $m_H(N)$

- Largest number of “effective” hypothesis  $H$  can induce on  $N$  data points
- A more precise “complexity” measure for  $H$
- Goal: Replace  $M$  in finite-hypothesis analysis with  $m_H(N)$

- With prob  $1 - \delta$ ,  $E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{1}{2N} \ln \frac{2M}{\delta}}$

- Theorem: VC Inequality (1971)

With prob  $1 - \delta$

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \frac{4m_H(2N)}{\delta}}$$

- Growth function  $m_H(N)$

- Largest number of “effective” hypothesis  $H$  can induce on  $N$  data points
- A more precise “complexity” measure for  $H$
- Goal: Replace  $M$  in finite-hypothesis analysis with  $m_H(N)$ 
  - With prob  $1 - \delta$ ,  $E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{1}{2N} \ln \frac{2M}{\delta}}$

- Theorem: VC Inequality (1971)

With prob  $1 - \delta$

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \frac{4m_H(2N)}{\delta}}$$

# Growth Functions for Other $H$

- $H$  = 2-D Perceptron
  - What is  $m_H(3)$
  - What is  $m_H(4)$

- Dichotomies

- Informally, consider a dichotomy as a “data-dependent” hypothesis
- Characterized by both hypothesis set  $H$  and  $N$  data points  $(\vec{x}_1, \dots, \vec{x}_N)$

$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$

- The set of possible prediction combinations  $h \in H$  can induce on  $\vec{x}_1, \dots, \vec{x}_N$

- Growth function

- Largest number of dichotomies  $H$  can induce across all possible data sets of size  $N$

$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

# Growth Functions for Other $H$

- $H = 2\text{-D Perceptron}$

- What is  $m_H(3)$
- What is  $m_H(4)$

- Dichotomies

- Informally, consider a dichotomy as a “data-dependent” hypothesis
- Characterized by both hypothesis set  $H$  and  $N$  data points  $(\vec{x}_1, \dots, \vec{x}_N)$

$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$

- The set of possible prediction combinations  $h \in H$  can induce on  $\vec{x}_1, \dots, \vec{x}_N$

- Growth function

- Largest number of dichotomies  $H$  can induce across all possible data sets of size  $N$

$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

- Exactly calculating the growth function is generally hard!
- Next lecture
  - Discuss how we can “bound” the growth function
  - Introduce the notion of VC dimension