

CSE 417T

Introduction to Machine Learning

Lecture 5

Instructor: Chien-Ju (CJ) Ho

Logistics: Homework 1

- Due: **September 23 (Friday), 2022**
 - <http://chienjuho.com/courses/cse417t/hw1.pdf>
- Two submission links: Report and Code (The links will be up over the weekend)
 - Report: Answer all questions, including the implementation question
 - **Grades are based on the report**
 - Code: Complete and submit **hw1.py** for Problem 2
 - The code will only be used for correctness checking (when in doubts) and plagiarism checking
- Reserve time if you never used Gradescope.
 - Make sure to **specify the pages for each problem**. You **won't get points** otherwise

Logistics: TA Office Hours

- Tentative schedule of TA office hours

Monday	9:30am Asher Baraban	3pm Qihang Zhao	
Tuesday	10am Di Huang	1pm Andrew Ruttenberg	4pm Quinn Wai Wong
Wednesday	1pm Wenxuan Zhu	3pm William Sepesi	4:30pm Sylvia Tang
Thursday	11:30am Yuan Liu	4pm Elyse Tang	6:30pm Fankun Zen
Friday	11am Riggie Kong	3pm Nan Huang	5:30pm Weiwei Ma
Sunday	Noon Jonathan Ma	1:30pm Kenneth Li	


- 60 minutes per session; In-person office hours are highlighted in orange
- Please follow **Piazza** for additional information (location, zoom link, etc)
- Recommendation: Try to utilize the office hour early (way ahead of deadlines), you are likely to get more of TAs' time this way

Logistics: This Thursday Lecture

- A shortened lecture this Thursday (Sep 15)
 - The lecture starts at 3:10pm
 - I need to attend a meeting that I cannot skip
- As a make-up
 - I'll host an additional one-hour office hour today
 - 5:30pm – 6:30pm at McKelvey 2010A
 - My (tentative) regular office hour
 - 5:30pm – 6:30pm Thursday at McKelvey 2010A

Recap

Hoeffding's Inequality



$$\Pr[|\mu - \nu| > \epsilon] \leq 2e^{-2\epsilon^2 N}$$

Define $\delta = \Pr[|\mu - \nu| > \epsilon]$

- Fix δ , ϵ decreases as N increases
- Fix ϵ , δ decreases as N increases
- Fix N , δ decreases as ϵ increases

Informal intuitions of notations
 N : # sample
 δ : probability of "bad" event
 ϵ : error of estimation

Connection to Learning

- Given dataset $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$
 - $E_{in}(h) \stackrel{\text{def}}{=} \frac{1}{N} \sum_{n=1}^N \mathbb{I}[h(\vec{x}_n) \neq f(\vec{x}_n)]$ [In-sample error, analogy to v]
 - $E_{out}(h) \stackrel{\text{def}}{=} \Pr_{\vec{x} \sim P(\vec{x})} [h(\vec{x}) \neq f(\vec{x})]$ [Out-of-sample error, analogy to μ]

- Learning bounds

- Fixed h (verification)

$$\Pr[|E_{out}(h) - E_{in}(h)| > \epsilon] \leq 2e^{-2\epsilon^2 N}$$

- Finite hypothesis set: learn $g \in \{h_1, \dots, h_M\}$

$$\Pr[|E_{out}(g) - E_{in}(g)| > \epsilon] \leq 2Me^{-2\epsilon^2 N}$$


Dealing with Infinite Hypothesis Set: $M \rightarrow \infty$

- Most of the practical cases involve $M \rightarrow \infty$
- Instead of # hypothesis, counting “effective” # hypothesis
- Dichotomies
 - Informally, consider a dichotomy as “data-dependent” hypothesis
 - Characterized by both H and N data points $(\vec{x}_1, \dots, \vec{x}_N)$
$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$
 - The set of possible prediction combinations $h \in H$ can induce on $\vec{x}_1, \dots, \vec{x}_N$
- Growth function
 - Largest number of dichotomies H can induce across all possible data sets of size N

$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

Examples on Growth Functions

- H = Positive rays
 - $m_H(N) = N + 1$
- H = Positive intervals
 - $m_H(N) = \binom{N+1}{2} + 1 = \frac{N^2}{2} + \frac{N}{2} + 1$
- H = Convex sets
 - $m_H(N) = 2^N$
- For all H and for all N
 - $m_H(N) \leq 2^N$



Why Growth Function?

- Growth function $m_H(N)$
 - Largest number of “effective” hypothesis H can induce on N data points
 - A more precise “complexity” measure for H
 - Goal: Replace M in finite-hypothesis analysis with $m_H(N)$
 - With prob at least $1 - \delta$, $E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{1}{2N} \ln \frac{2M}{\delta}}$

- VC Generalization Bound (VC Inequality, 1971)

With prob at least $1 - \delta$

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \frac{4m_H(2N)}{\delta}}$$

[We will skip the proof of the VC bound. You can find it in the textbook appendix if you are interested.]

Today's Lecture


The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.
Let me know if you spot errors.

Bounding Growth Function

- What we know so far
 - H = Positive rays: $m_H(N) = N + 1$
 - H = Positive intervals: $m_H(N) = \binom{N+1}{2} + 1$
 - H = Convex sets: $m_H(N) = 2^N$

- What about H = 2-D Perceptron?

- $m_H(3) = 8$
- $m_H(4) = 14$
- $m_H(5) = ?$



- Generally hard to write down the growth function exactly
 - Goal: “bound” the growth function using some proxy

Bounding Growth Function

- More definitions....
 - Shatter:
 - H **shatters** $(\vec{x}_1, \dots, \vec{x}_N)$ if $|H(\vec{x}_1, \dots, \vec{x}_N)| = 2^N$
 - H can induce all label combinations for $(\vec{x}_1, \dots, \vec{x}_N)$
 - Break point
 - k is a **break point** for H if no data set of size k can be shattered by H

- A peek at the key result (take this as a fact for now)
 - If there are no break points for H , $m_H(N) = 2^N$
 - If k is a break point for H , $m_H(N)$ is polynomial in N .

In particular, $m_H(N) = O(N^{k-1})$ 

A bit more accurately:

- $m_H(N) \leq \sum_{i=1}^{k-1} \binom{N}{i}$, or
- $m_H(N) \leq N^{k-1} + 1$

Practice

- Dichotomies

- Informally, consider a dichotomy as a “data-dependent” hypothesis
- Characterized by both hypothesis set H and N data points $(\vec{x}_1, \dots, \vec{x}_N)$

$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$

- The set of possible prediction combinations $h \in H$ can induce on $\vec{x}_1, \dots, \vec{x}_N$

- Growth function

- Largest number of dichotomies H can induce across all possible data sets of size N

$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

- Shatter:


- H **shatters** $(\vec{x}_1, \dots, \vec{x}_N)$ if $|H(\vec{x}_1, \dots, \vec{x}_N)| = 2^N$
- H can induce all label combinations for $(\vec{x}_1, \dots, \vec{x}_N)$

- Break point


- k is a **break point** for H if no data set of size k can be shattered by H

- What are the break points for


1. Positive Rays




2. Positive Intervals



3. Convex Sets



4. 2-D Perceptron



Practice

- Dichotomies
 - Informally, consider a dichotomy as a “data-dependent” hypothesis
 - Characterized by both hypothesis set H and N data points $(\vec{x}_1, \dots, \vec{x}_N)$

$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$
 - The set of possible prediction combinations $h \in H$ can induce on $\vec{x}_1, \dots, \vec{x}_N$
- Growth function
 - Largest number of dichotomies H can induce across all possible data sets of size N

$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

- Shatter:
 - H **shatters** $(\vec{x}_1, \dots, \vec{x}_N)$ if $|H(\vec{x}_1, \dots, \vec{x}_N)| = 2^N$
 - H can induce all label combinations for $(\vec{x}_1, \dots, \vec{x}_N)$
- Break point
 - k is a **break point** for H if no data set of size k can be shattered by H

	$m_H(N)$					
$m_H(N)$	N=1	N=2	N=3	N=4	N=5	Break Points
$N + 1$						Positive Rays
$\frac{N^2}{2} + \frac{N}{2} + 1$						Positive Intervals
2^N						Convex Sets
						2D Perceptron

Practice

- Dichotomies
 - Informally, consider a dichotomy as a “data-dependent” hypothesis
 - Characterized by both hypothesis set H and N data points $(\vec{x}_1, \dots, \vec{x}_N)$

$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$
 - The set of possible prediction combinations $h \in H$ can induce on $\vec{x}_1, \dots, \vec{x}_N$
- Growth function
 - Largest number of dichotomies H can induce across all possible data sets of size N

$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

- Shatter:
 - H **shatters** $(\vec{x}_1, \dots, \vec{x}_N)$ if $|H(\vec{x}_1, \dots, \vec{x}_N)| = 2^N$
 - H can induce all label combinations for $(\vec{x}_1, \dots, \vec{x}_N)$
- Break point
 - k is a **break point** for H if no data set of size k can be shattered by H

		$m_H(N)$					
$m_H(N)$		N=1	N=2	N=3	N=4	N=5	Break Points
$N + 1$	Positive Rays	2	3	4	5	6	$k = 2,3,4, \dots$
$\frac{N^2}{2} + \frac{N}{2} + 1$	Positive Intervals						
2^N	Convex Sets						
	2D Perceptron						

Practice

- Dichotomies
 - Informally, consider a dichotomy as a “data-dependent” hypothesis
 - Characterized by both hypothesis set H and N data points $(\vec{x}_1, \dots, \vec{x}_N)$
$$H(\vec{x}_1, \dots, \vec{x}_N) = \{(h(\vec{x}_1), \dots, h(\vec{x}_N)) | h \in H\}$$
 - The set of possible prediction combinations $h \in H$ can induce on $\vec{x}_1, \dots, \vec{x}_N$
- Growth function
 - Largest number of dichotomies H can induce across all possible data sets of size N
$$m_H(N) = \max_{(\vec{x}_1, \dots, \vec{x}_N)} |H(\vec{x}_1, \dots, \vec{x}_N)|$$

- Shatter:
 - H **shatters** $(\vec{x}_1, \dots, \vec{x}_N)$ if $|H(\vec{x}_1, \dots, \vec{x}_N)| = 2^N$
 - H can induce all label combinations for $(\vec{x}_1, \dots, \vec{x}_N)$
- Break point
 - k is a **break point** for H if no data set of size k can be shattered by H

		$m_H(N)$					
$m_H(N)$		N=1	N=2	N=3	N=4	N=5	Break Points
$N + 1$	Positive Rays	2	3	4	5	6	$k = 2, 3, 4, \dots$
$\frac{N^2}{2} + \frac{N}{2} + 1$	Positive Intervals	2	4	7	11	16	$k = 3, 4, 5, \dots$
2^N	Convex Sets	2	4	8	16	32	None
	2D Perceptron	2	4	8	14	?	$k = 4, 5, 6, \dots$

Why Break Points?

- Theorem statement (Again, take it as a fact for now)
 - If there is no break point for H , then $m_H(N) = 2^N$ for all N .
 - If k is a break point for H , i.e., if $m_H(k) < 2^k$ for some value k , then
$$m_H(N) \leq \sum_{i=0}^{k-1} \binom{N}{i}$$
- The above theorem can be rephrased below:
 - If there is no break point for H , then $m_H(N) = 2^N$ for all N .
 - If k is a break point for H , the following statements are true
 - $m_H(N) \leq N^{k-1} + 1$ [Can be proven using induction. See LFD Problem 2.5]
 - $m_H(N) = O(N^{k-1})$
 - $m_H(N)$ is polynomial in N
- We can “bound” the growth function without knowing it exactly.
 - Find break point!

Recall the VC Generalization Bound

- VC Generalization Bound

With prob $1 - \delta$

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \frac{4m_H(2N)}{\delta}}$$

- In the following discussion, we treat δ as a constant [i.e., with high probability, the following is true]

$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{\frac{1}{N} \ln m_H(N)}\right)$$

[For example, we can set δ to be a small constant, say 0.01. Then every time we wrote the above inequality, we mean that it is true with probability at least 99%.]

Applying Break Points in VC Bound

- VC Bound:

$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{\frac{1}{N} \ln m_H(N)}\right)$$



- Rephrase the above theorem

- If there is no break point for H , then $m_H(N) = 2^N$ for all N .
- If k is a break point for H , the following statements are true
 - $m_H(N) \leq N^{k-1} + 1$ [Can be proven using induction. See LFD Problem 2.5]
 - $m_H(N) = O(N^{k-1})$
 - $m_H(N)$ is polynomial in N

- If there are no break point ($m_H(N) = 2^N$)

$$E_{out}(g) \leq E_{in}(g) + O(1)$$

(This implies that we can't infer E_{out} from E_{in} even when $N \rightarrow \infty$)

- If k is a break point for H , i.e., $m_H(N) = O(N^{k-1})$

$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{(k-1) \frac{\ln N}{N}}\right)$$

H is Either Good or Bad

- Rephrase the above theorem
 - If there is no break point for H , then $m_H(N) = 2^N$ for all N .
 - If k is a break point for H , the following statements are true
 - $m_H(N) \leq N^{k-1} + 1$ [Can be proven using induction. See LFD Problem 2.5]
 - $m_H(N) = O(N^{k-1})$
 - $m_H(N)$ is polynomial in N


- The growth function of H is either one of the two
 - Without break points, $m_H(N) = 2^N$
 - With some break point, $m_H(N)$ is polynomial in N (it can be bounded more tightly using the theorem)
 - There is nothing in between!

- **Bad** hypothesis set

$$E_{out}(g) \leq E_{in}(g) + O(1)$$

- **Good** hypothesis set $m_H(N) = O(N^{k-1})$

$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{(k-1)\frac{\ln N}{N}}\right)$$



VC Dimension

- VC Dimension of H : $d_{vc}(H)$ or d_{vc}
 - The VC dimension of H is the **largest N such that $m_H(N) = 2^N$** .
 - $d_{vc}(H) = \infty$ if $m_H(N) = 2^N$ for all N .
 - Or, let k^* be the smallest break point for H , the VC dimension of H is $k^* - 1$

	$m_H(N)$					Break Points	VC Dimension
	N=1	N=2	N=3	N=4	N=5		
Positive Rays	2	3	4	5	6	$k = 2, 3, 4, \dots$	
Positive Intervals	2	4	7	11	16	$k = 3, 4, 5, \dots$	
Convex Sets	2	4	8	16	32	None	
2D Perceptron	2	4	8	14	?	$k = 4, 5, 6, \dots$	

VC Dimension

- VC Dimension of H : $d_{vc}(H)$ or d_{vc}
 - The VC dimension of H is the **largest N such that $m_H(N) = 2^N$** .
 - $d_{vc}(H) = \infty$ if $m_H(N) = 2^N$ for all N .
 - Or, let k^* be the smallest break point for H , the VC dimension of H is $k^* - 1$

	$m_H(N)$					Break Points	VC Dimension
	N=1	N=2	N=3	N=4	N=5		
Positive Rays	2	3	4	5	6	$k = 2, 3, 4, \dots$	1
Positive Intervals	2	4	7	11	16	$k = 3, 4, 5, \dots$	2
Convex Sets	2	4	8	16	32	None	∞
2D Perceptron	2	4	8	14	?	$k = 4, 5, 6, \dots$	3

VC Dimension

- VC Dimension of H : $d_{vc}(H)$ or d_{vc}

- The VC dimension of H is the **largest N such that $m_H(N) = 2^N$** .

- $d_{vc}(H) = \infty$ if $m_H(N) = 2^N$ for all N .

- Or, let k^* be the smallest break point for H , the VC dimension of H is $k^* - 1$

- Plug the definition into VC Generalization Bound

$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{vc} \frac{\ln N}{N}}\right)$$

- If there are no break point ($m_H(N) = 2^N$)

$$E_{out}(g) \leq E_{in}(g) + O(1)$$

(This implies that we can't infer E_{out} from E_{in} even when $N \rightarrow \infty$)

- If k is a break point for H , i.e., $m_H(N) = O(N^{k-1})$

$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{(k-1) \frac{\ln N}{N}}\right)$$

Discussion on the VC Theory

*All models are wrong
but some are useful*



George E.P. Box

Discussion on the VC Theory

- VC Bound

$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{VC} \frac{\ln N}{N}}\right)$$

- Built on top of the i.i.d. data assumption
- The bound is “loose”
 - Depends only on H and N
 - The analysis is loose in many places
- However, it qualitatively characterizes the practice reasonably well
 - (the bound is roughly equally loose for every H)


$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{VC} \frac{\ln N}{N}}\right)$$

- Goal of learning: Minimize $E_{out}(g)$
- How to achieve that
 - Minimize $E_{in}(g)$
 - Choose a hypothesis set with large d_{VC} (complex hypothesis likely fit data better)
 - Minimize **generalization error**
 - Choose a hypothesis with small d_{VC}
 - Have a lot of data points to train on (N is large)
- Think about the high-level tradeoff of choosing d_{VC} and its dependency on N

Discussion on the VC Theory

- It establishes the feasibility of learning for infinite hypothesis set
- It provides nice intuitions on what's happening underneath ML
 - A single parameter to characterize complexity of H


$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{VC} \frac{\ln N}{N}}\right)$$



Discussion on the VC Theory

- It establishes the feasibility of learning for infinite hypothesis set.
- It provides nice intuitions on what's happening underneath ML.
 - A single parameter to characterize complexity of H

$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{VC} \frac{\ln N}{N}}\right)$$



Sample Complexity

- Sample complexity:
 - Analogy to time/space complexity
 - How many data points do we need to achieve generalization error less than ϵ with prob $1 - \delta$?

- Recall the (full) VC Bound:

$$\text{With prob at least } 1 - \delta, E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \frac{4((2N)^{d_{vc}} + 1)}{\delta}}$$

- How to determine the sample complexity?

- Set $\sqrt{\frac{8}{N} \ln \frac{4((2N)^{d_{vc}} + 1)}{\delta}} \leq \epsilon$
- We get $N \geq \frac{8}{\epsilon^2} \ln \left(\frac{4(1 + (2N)^{d_{vc}})}{\delta} \right)$

- $N \propto 1/\epsilon^2$
- $N = O(d_{vc} \ln N)$
 - Empirically, people roughly $N \propto d_{vc}$

Test Set

- Goal of learning: Minimize $E_{out}(g)$
- Can we estimate E_{out} directly?
 - Reserve a test set (D_{test}) before learning
 - Ensure D_{test} is **not used at all** in any way for learning
 - For D_{test} , g is a “fixed” hypothesis and standard Hoeffding’s inequality is valid
 - Let $E_{test}(g)$ be the error in the test set

$$P\{|E_{test}(g) - E_{out}(g)| > \epsilon\} \leq 2e^{-2\epsilon^2 N_{test}} \text{ where } N_{test} = |D_{test}|$$

Test Set

- Test set is great: we can obtain an unbiased estimate of E_{out}
- At what cost?
 - We have a finite amount of data
 - Data points in test set cannot be involved in learning at all
 - More points in test set
 - Better estimate of E_{out}
 - Less data points in training set -> often leads to worse learned hypothesis
- Practical rule of thumb (i.e., a common heuristic, not really a gold rule)
 - 80% for training, 20% for testing

Proof: Bounding Growth Functions

Recall: Theorem in Bounding Growth Function

- Theorem statement:
 - If there is no break point for H , then $m_H(N) = 2^N$ for all N .
 - If k is a break point for H , i.e., if $m_H(k) < 2^k$ for some value k , then

$$m_H(N) \leq \sum_{i=0}^{k-1} \binom{N}{i}$$

- You were asked to take this as a fact
- Will provide proof sketch now

Proof Sketch

[See LFD Section 2.1.2 for the formal proof]

[Safe to Skip] (This proof won't appear in exams/homework)

[Safe to Skip]

Key Intuitions

- When there exist a break point k
 - No datasets of size k can be shattered
 - It also imposes strong constraints on dataset of size $k' > k$
 - No subset of data with size k can be shattered
- This leads to the bound $m_H(N) = O(N^{k-1})$

[Safe to Skip]

Proof Intuitions

- Max # dichotomies you can list on **2 points** when **no 2 points can be shattered**

\vec{x}_1	\vec{x}_2
+1	+1
+1	-1
-1	+1

[Safe to Skip]

Proof Intuitions

- Max # dichotomies you can list on **4 points** when **no 2 points can be shattered**

\vec{x}_1	\vec{x}_2	\vec{x}_3	\vec{x}_4
+1	+1	+1	+1
+1	+1	+1	-1
+1	+1	-1	+1
+1	-1	+1	+1
-1	+1	+1	+1

Can you add an additional dichotomy?

[Safe to Skip]

Proof Intuitions

- How “no 2 points are shattered” impacts the scenario with **4 points**?

\vec{x}_1	\vec{x}_2	\vec{x}_3	\vec{x}_4
+1	+1	+1	+1
+1	+1	+1	-1
+1	+1	-1	+1
+1	-1	+1	+1
-1	+1	+1	+1

$(\vec{x}_1, \vec{x}_2, \vec{x}_3)$ appear twice, with different \vec{x}_4

No 1 point is shattered

$(\vec{x}_1, \vec{x}_2, \vec{x}_3)$ appear once (including one in each of the pair above)

No 2 points are shattered

[Safe to Skip]

Proof Intuitions

- Max # dichotomies you can list on **4 points** when **no 2 points can be shattered**

No 1 point is shattered

\vec{x}_1	\vec{x}_2	\vec{x}_3	\vec{x}_4
+1	+1	+1	+1
+1	+1	+1	-1
+1	+1	-1	+1
+1	-1	+1	+1
-1	+1	+1	+1

No 2 points are shattered

$B(N, k)$: max # dichotomies on N points when no k points are shattered

A recursive definition:

$$B(N, k) \leq B(N - 1, k) + B(N - 1, k - 1)$$

Sauer's Lemma: $B(N, k) \leq \sum_{i=0}^{k-1} \binom{N}{i}$

Can be proved by induction

$B(N, k) \leq \sum_{i=0}^{k-1} \binom{N}{i}$ is the bound of $m_H(N)$ for H with break point k

Summary: Bounding Growth Functions

- Theorem statement:

- If there is no break point for H , then $m_H(N) = 2^N$ for all N .
- If k is a break point for H , i.e., if $m_H(k) < 2^k$ for some value k , then

$$m_H(N) \leq \sum_{i=0}^{k-1} \binom{N}{i}$$

- Rephrase the 2nd statement of the above theorem

- If k is a break point for H , the following statements are true
 - $m_H(N) \leq N^{k-1} + 1$ [Can be proven using induction from above. See LFD Problem 2.5]
 - $m_H(N) = O(N^{k-1})$
 - $m_H(N)$ is polynomial in N

- If d_{vc} is the VC dimension of H , then

- $m_H(N) \leq \sum_{i=0}^{d_{vc}} \binom{N}{i}$
- $m_H(N) \leq N^{d_{vc}} + 1$
- $m_H(N) = O(N^{d_{vc}})$

If d_{vc} is the VC dimension of H ,
 $d_{vc} + 1$ is a break point for H

Summary: Vapnik–Chervonenkis (VC) Bound

- VC Generalization Bound

With prob at least $1 - \delta$

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \frac{4m_H(2N)}{\delta}}$$

- Let d_{vc} be the VC dimension of H , we have $m_H(N) \leq N^{d_{vc}} + 1$.

With prob at least $1 - \delta$

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \frac{4((2N)^{d_{vc}} + 1)}{\delta}}$$

- If we treat δ as a constant, then we can say, with high probability


$$E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{vc} \frac{\ln N}{N}}\right)$$

Bias-Variance Decomposition

Another theory of generalization

Real-Value Target and Squared Error

- So far, we focus on binary target function and binary error
 - Binary target function $f(\vec{x}) \in \{-1, 1\}$
 - Binary error $e(h(\vec{x}), f(\vec{x})) = \mathbb{I}[h(\vec{x}) \neq f(\vec{x})]$
- Real-value functions [“**regression**”] and squared error?
 - Real-value target function $f(\vec{x}) \in \mathbb{R}$
 - Square error $e(h(\vec{x}), f(\vec{x})) = (h(\vec{x}) - f(\vec{x}))^2$



Real-Value Target and Squared Error

- Real-value functions [called "**regression**"] and squared error
 - Real-value target function $f(\vec{x}) \in \mathbb{R}$
 - Squared error $e(h(\vec{x}), f(\vec{x})) = (h(\vec{x}) - f(\vec{x}))^2$
- Errors:
 - In-sample error: $E_{in}(g) = \frac{1}{N} \sum_{n=1}^N e(h(\vec{x}_n), f(\vec{x}_n)) = \frac{1}{N} \sum_{n=1}^N (h(\vec{x}_n) - f(\vec{x}_n))^2$
 - Out-of-sample error: $E_{out}(g) = \mathbb{E}_{\vec{x}}[e(h(\vec{x}), f(\vec{x}))] = \mathbb{E}_{\vec{x}}[(g(\vec{x}) - f(\vec{x}))^2]$
- Theory of generalization: What can we say about $E_{out}(g)$?

- Note that g is learned by some algorithm on the dataset D
 - We'll make the dependency on D explicit and write it as $g^{(D)}$ here.
 - [In VC theory, we consider the worst-case D through the definition of growth function $m_H(N)$]

- $E_{out}(g^{(D)}) = \mathbb{E}_{\vec{x}}[(g^{(D)}(\vec{x}) - f(\vec{x}))^2]$

- $\mathbb{E}_D[E_{out}(g^{(D)})]$

$$= \mathbb{E}_D \left[\mathbb{E}_{\vec{x}} \left[(g^{(D)}(\vec{x}) - f(\vec{x}))^2 \right] \right]$$

$$= \mathbb{E}_{\vec{x}} \left[\mathbb{E}_D \left[(g^{(D)}(\vec{x}) - \bar{g}(\vec{x}) + \bar{g}(\vec{x}) - f(\vec{x}))^2 \right] \right]$$

$$= \mathbb{E}_{\vec{x}} \left[\mathbb{E}_D \left[(g^{(D)}(\vec{x}) - \bar{g}(\vec{x}) + \bar{g}(\vec{x}) - f(\vec{x}))^2 \right] \right]$$

$$= \mathbb{E}_{\vec{x}} \left[\mathbb{E}_D \left[(g^{(D)}(\vec{x}) - \bar{g}(\vec{x}))^2 + (\bar{g}(\vec{x}) - f(\vec{x}))^2 + 2(g^{(D)}(\vec{x}) - \bar{g}(\vec{x}))(\bar{g}(\vec{x}) - f(\vec{x})) \right] \right]$$

- Note that $\mathbb{E}_D \left[(g^{(D)}(\vec{x}) - \bar{g}(\vec{x}))(\bar{g}(\vec{x}) - f(\vec{x})) \right] = (\bar{g}(\vec{x}) - f(\vec{x})) \mathbb{E}_D \left[(g^{(D)}(\vec{x}) - \bar{g}(\vec{x})) \right] = 0$

Define “expected” hypothesis
 $\bar{g}(\vec{x}) = \mathbb{E}_D[g^{(D)}(\vec{x})]$

$$\bar{g}(\vec{x}) = \mathbb{E}_D[g^{(D)}(\vec{x})]$$

Finishing Up

- $$\begin{aligned} & \mathbb{E}_D[E_{out}(g^{(D)})] \\ &= \mathbb{E}_{\vec{x}} \left[\mathbb{E}_D \left[\left(g^{(D)}(\vec{x}) - \bar{g}(\vec{x}) \right)^2 + \left(\bar{g}(\vec{x}) - f(\vec{x}) \right)^2 \right] \right] \\ &= \mathbb{E}_{\vec{x}} \left[\mathbb{E}_D \left[\left(g^{(D)}(\vec{x}) - \bar{g}(\vec{x}) \right)^2 \right] \right] + \mathbb{E}_{\vec{x}} \left[\left(\bar{g}(\vec{x}) - f(\vec{x}) \right)^2 \right] \\ &= \mathbb{E}_{\vec{x}} [\text{Variance of } g^{(D)}(\vec{x}) + \text{Bias of } \bar{g}(\vec{x})] \\ &= \text{Variance} + \text{Bias} \end{aligned}$$

X : a random variable
 μ : the mean of X

Variance of X :
 $Var(X) = \mathbb{E}[(X - \mu)^2]$

- Bias-Variance Decomposition

Discussion

$$\bullet \mathbb{E}_D[E_{out}(g^{(D)})] = \mathbb{E}_{\vec{x}} \left[\overset{\text{Bias}(\vec{x})}{\left(\bar{g}(\vec{x}) - f(\vec{x}) \right)^2} \right] + \mathbb{E}_{\vec{x}} \left[\overset{\text{Var}(\vec{x})}{\mathbb{E}_D \left[\left(g^{(D)}(\vec{x}) - \bar{g}(\vec{x}) \right)^2 \right]} \right]$$

- This is a **conceptual** decomposition
 - Both \bar{g} and f are unknown
 - We can't really calculate bias and variance in practice
- However, it provides a conceptual guideline in decreasing E_{out}