

CSE 417T

Introduction to Machine Learning

Lecture 10

Instructor: Chien-Ju (CJ) Ho

Logistics

- Homework 2: due on **Feb 24 (Thursday)**
- Exam 1: **Mar 10 (Thursday)**
 - Topics: LFD Chapters 1 to 5
 - Covid-permitting
 - Timed exam (75 min) during lecture time in the classroom
 - Closed-book exam with 2 letter-size cheat sheets allowed (4 pages in total)
 - No format limitations (it can be typed, written, or a combination)
 - Mar 8 (Tuesday) will be a review lecture
- Homework 3 will be posted next week
 - Expected due: **Mar 5 (Saturday)**
 - A bit shorter amount of time than usual, but the goal is for you to read the topics that Exam 1 covers

Recap

VC Dimension of d-dimension Perceptron

- Claim:
 - The VC Dimension of d-dim perceptron is $d + 1$
- How to prove it?
 1. Show that the VC dimension of d-dim perceptron $\geq d + 1$
 2. Show that the VC dimension of d-dim perceptron $\leq d + 1$

- To prove $d_{vc}(H) \geq d + 1$, what do we need to prove?

There is a set of $d + 1$ points that can be shattered by H

Proof Sketch:

1. Let's construct a dataset of $d + 1$ points: $X = \begin{bmatrix} \vec{x}_1^T \\ \vdots \\ \vec{x}_{d+1}^T \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & \dots & 0 & 0 \\ 1 & 0 & 0 & \dots & 0 & 1 \\ 1 & 0 & 0 & \dots & 1 & 0 \\ \dots & \dots & \dots & \dots & \dots & \dots \\ 1 & 1 & 0 & \dots & 0 & 0 \end{bmatrix}$; It's easy to check that X^{-1} exist
2. For any possible dichotomy \vec{y} , there exists a \vec{w} such that $X\vec{w} = \vec{y}$, i.e., $\vec{w} = X^{-1}\vec{y}$
3. Therefore, d-dim perceptron can shatter X

- To prove $d_{vc}(H) \leq d + 1$, what do we need to prove?

Every set of $d + 2$ points cannot be shattered by H

Proof Sketch:

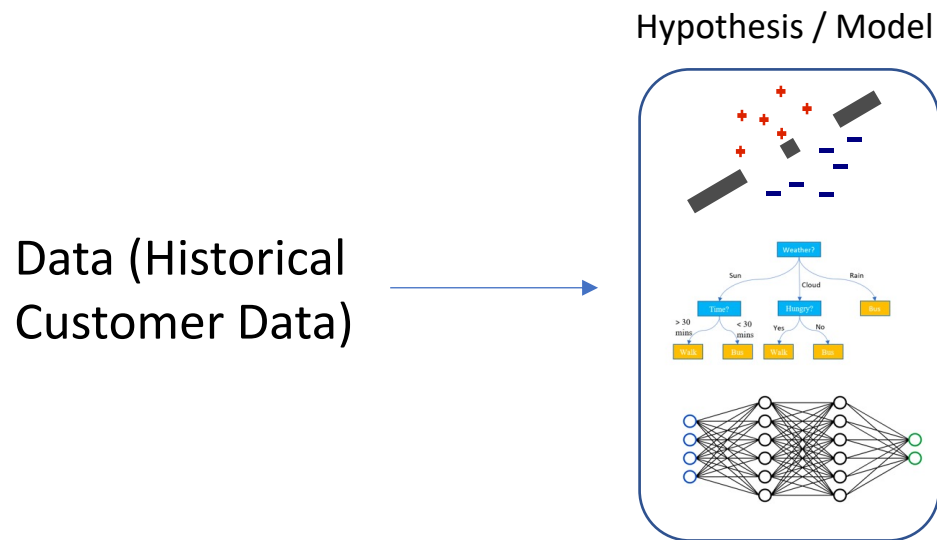
1. For every set of $d + 2$ points (in $d+1$ dimensions), there exists a point that can be written as linear combinations of the others.
2. Denote the point \vec{x}_{d+2} , we have $\vec{x}_{d+2} = \sum_{i=1}^{d+1} a_i \vec{x}_i$
3. Consider the dichotomy $(y_1, \dots, y_{d+2}) = (\text{sign}(a_1), \dots, \text{sign}(a_{d+1}), -1)$, we can show that no linear separator can generate this dichotomy (think about why).
4. Therefore, for every set of $d + 2$ points, there exist a dichotomy that H cannot shatter.

VC “Dimension”

- Degrees of freedom for your hypothesis in H
- (effective) # of parameters that control the hypothesis
- Examples:
 - d-dim perceptron: h is represented by (w_0, \dots, w_d) ; $d_{vc} = d + 1$
 - Positive rays: h is represented by a threshold; $d_{vc} = 1$
 - Positive or negative rays: h is represented by a threshold and a direction; $d_{vc} = 2$
 - Positive intervals: h is represented by two thresholds; $d_{vc} = 2$
 - Positive or negative intervals: h is represented by two thresholds and a direction; $d_{vc} = 3$
- Effective # parameters: An “approximation” for VC dimension

What We Have Taught So Far

- The explanation of “machine learning” from the first lecture



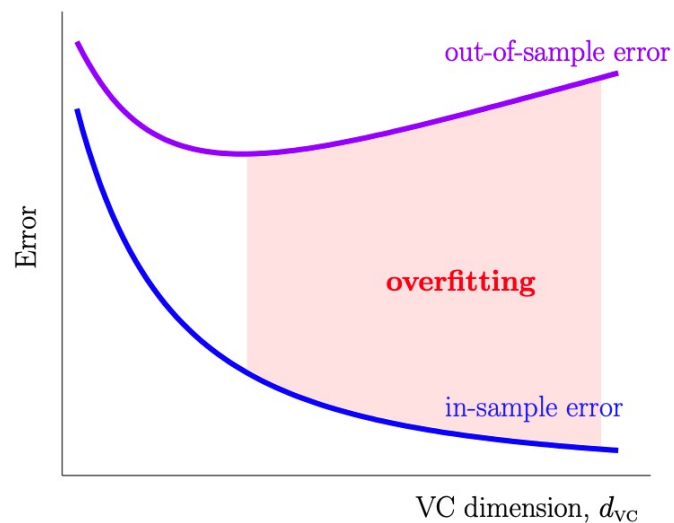
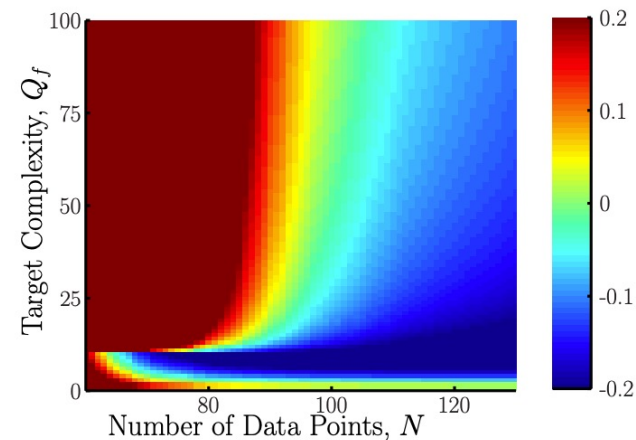

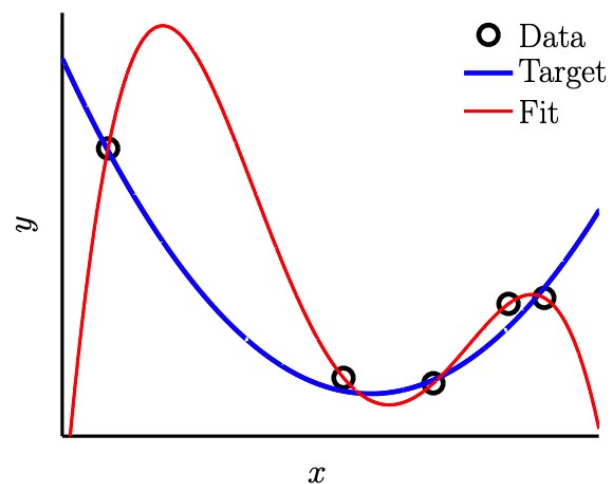
Find a hypothesis that “fits” the data
(The process requires a lot of computation)

Our progress so far

- Generalization of learning
 - What to say about $E_{out}(g)$ from $E_{in}(g)$
- How to find g
 - Using linear models as examples
 - Focus on $g = \operatorname{argmin}_{h \in H} E_{in}(g)$

Seems to make sense, but...

Overfitting



Number of data points \uparrow	Overfitting \downarrow
Noise \uparrow	Overfitting \uparrow
Target complexity \uparrow	Overfitting \uparrow

Today's Lecture

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.
Let me know if you spot errors.

Overfitting and Its Cures

- Overfitting

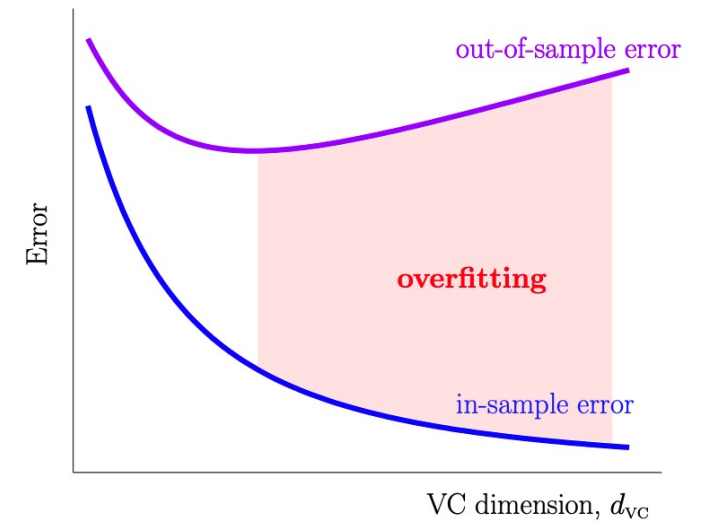
- Fitting the data more than is warranted
- Fitting the noise instead of the pattern of the data
- Decreasing E_{in} but getting larger E_{out}
- When H is too strong, but N is not large enough

- Regularization

- Intuition: Constrain H to make overfitting less likely to happen
- (Topic of this lecture)

- Validation

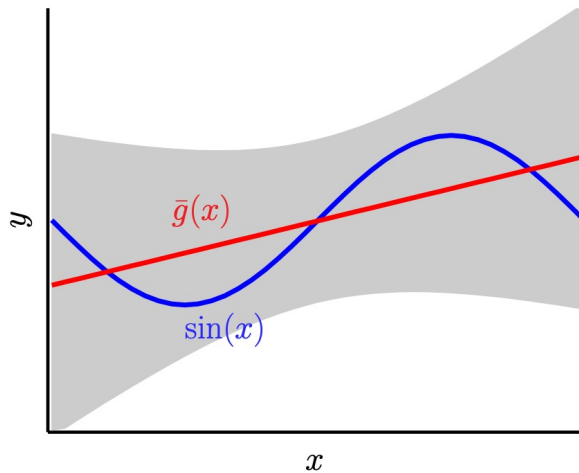
- Intuition: Reserve data to estimate E_{out}
- (Focus of next lecture)



Regularization (Constraining H)

- Informal example:

- Regression; $f = \sin(\pi x)$; $H = \{h(x) = ax + b\}$; $N = 2$

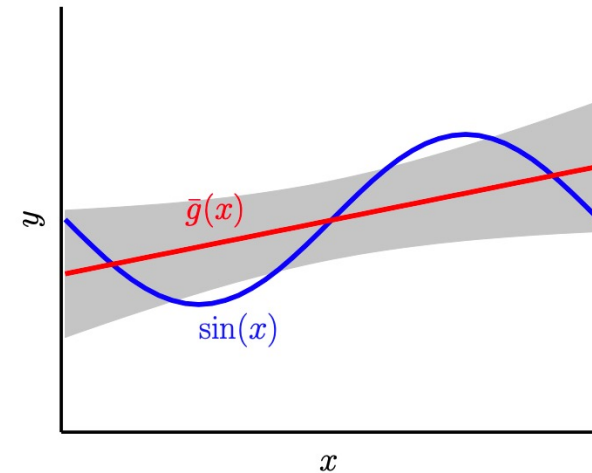


no regularization

bias = 0.21

var = 1.69

Regularization:
Constrain the hypothesis set to
avoid large a and b



regularization

bias = 0.23

var = 0.33

How to do this in a principled way?

Hard Constraints

- We have seen hard constraints already

$$H_2 = \{h(x) = w_0 + w_1x + w_2x^2\}$$

$$H_{10} = \{h(x) = w_0 + w_1x + w_2x^2 + \cdots + w_{10}x^{10}\}$$

- H_2 can be written as constrained H_{10}

$$H_2 = \{h \in H_{10} \text{ and } w_3 = w_4 = \cdots = w_{10} = 0\}$$



Constraints

Soft-Order Constraints

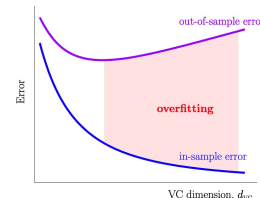
Hard constraints

$$H_2 = \{h \in H_{10} \text{ and } w_3 = w_4 = \dots = w_{10} = 0\}$$

- Instead of setting the weights to 0

$$H(C) = \left\{ h \in H_Q \text{ and } \sum_{q=0}^Q w_q^2 \leq C \right\}$$
$$= \{h \in H_Q \text{ and } \vec{w}^T \vec{w} \leq C\}$$

- Observations
 - When $C \rightarrow \infty$, $H(C) = H_Q$
 - When $C_1 \leq C_2$, $H(C_1) \subseteq H(C_2)$ and therefore $d_{vc}(H(C_1)) \leq d_{vc}(H(C_2))$
 - A smoother way to tune the complexity of hypothesis set



Soft-Order Constraints

$$H(\mathcal{C}) = \{h \in H_Q \text{ and } \vec{w}^T \vec{w} \leq \mathcal{C}\}$$

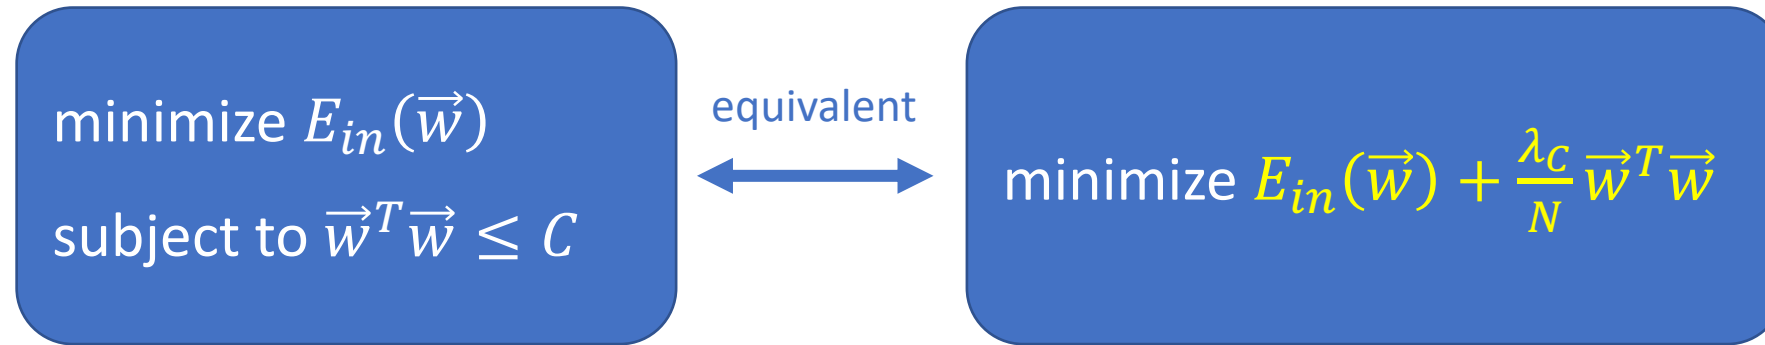
- Two main questions
 - How do we choose \mathcal{C}
 - Model selection: The same question as selecting H
 - The focus of the next lecture
 - How do we perform learning, i.e., find a $g \in H(\mathcal{C})$ such that $g \approx f$
 - Solve the following **constrained optimization** problem

minimize $E_{in}(\vec{w})$

subject to $\vec{w}^T \vec{w} \leq \mathcal{C}$

Constrained to Unconstrained Optimization

- Constrained optimization \Leftrightarrow Unconstrained optimization



- Why the above is true?
 - Will talk about how to utilize Lagrangian relaxation to get this in the 2nd half of the semester
 - For now, let's think about it graphically

minimize $E_{in}(\vec{w})$ subject to $\vec{w}^T \vec{w} \leq C$

- Notations

- \vec{w}_{lin} : the solution for $\min E_{in}(\vec{w})$
- \vec{w}_{reg} : the solution for $\min E_{in}(\vec{w})$ subject to $\vec{w}^T \vec{w} \leq C$

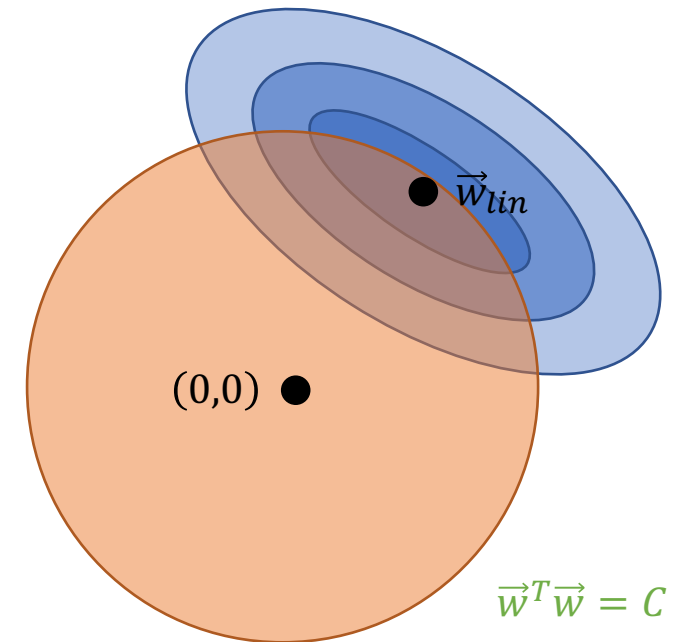
minimize $E_{in}(\vec{w})$ subject to $\vec{w}^T \vec{w} \leq C$

- Notations

- \vec{w}_{lin} : the solution for $\min E_{in}(\vec{w})$
- \vec{w}_{reg} : the solution for $\min E_{in}(\vec{w})$ subject to $\vec{w}^T \vec{w} \leq C$

- When C is large enough, i.e., $\vec{w}_{lin}^T \vec{w}_{lin} \leq C$
 - $\vec{w}_{reg} = \vec{w}_{lin}$

- When C is not large enough, i.e., $\vec{w}_{lin}^T \vec{w}_{lin} > C$
 - $\vec{w}_{reg}^T \vec{w}_{reg} = C$



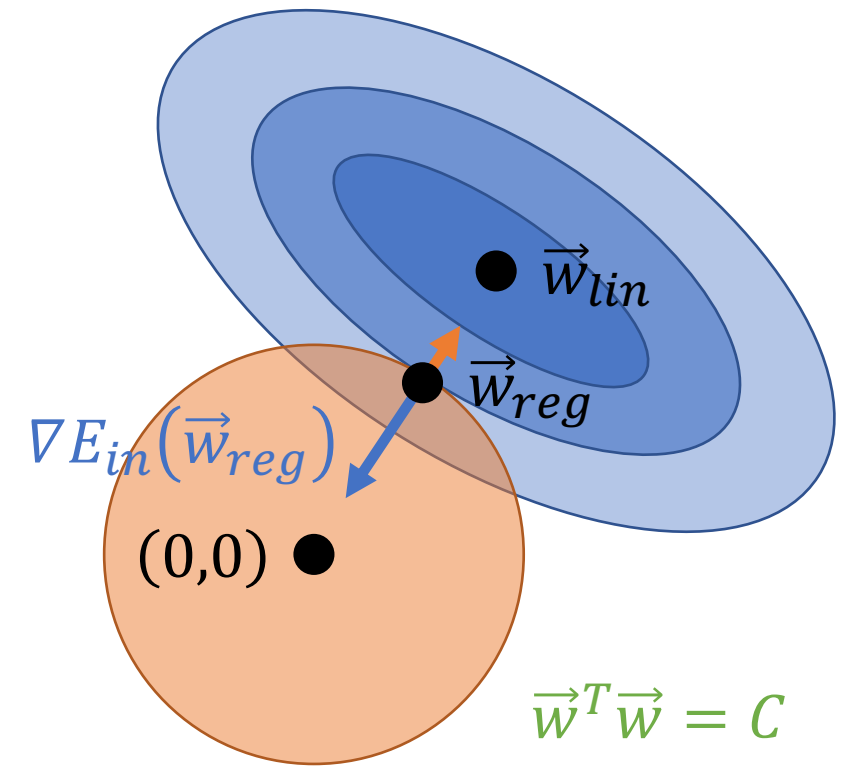
minimize $E_{in}(\vec{w})$ subject to $\vec{w}^T \vec{w} \leq C$

- When C is not large enough

- \vec{w}_{lin} : the solution for $\min E_{in}(\vec{w})$
- \vec{w}_{reg} : the solution for $\min E_{in}(\vec{w})$ subject to $\vec{w}^T \vec{w} \leq C$

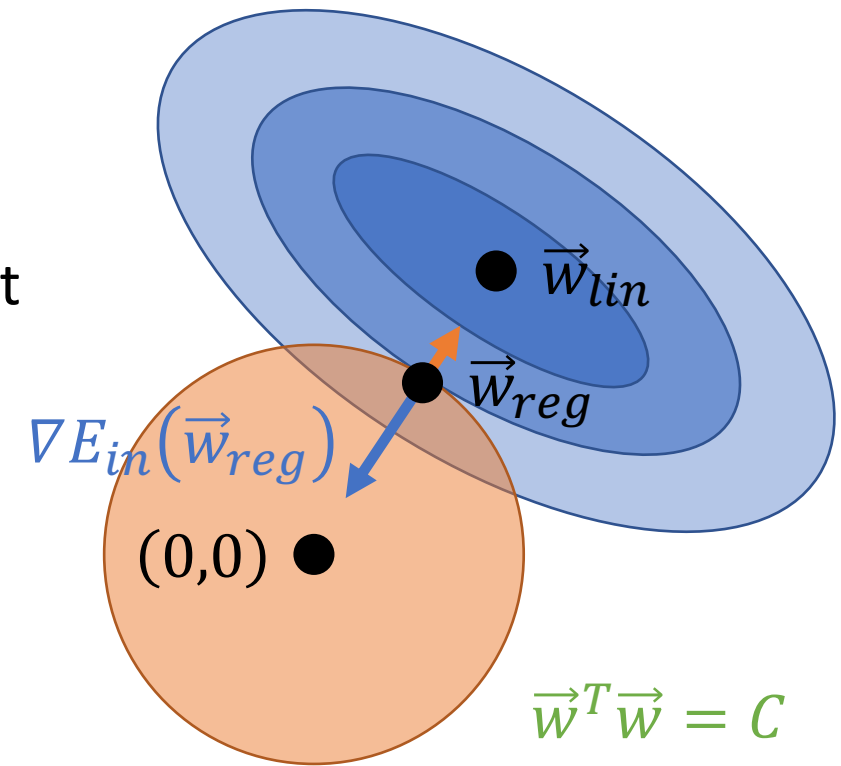
minimize $E_{in}(\vec{w})$ subject to $\vec{w}^T \vec{w} \leq C$

- When C is not large enough
 - Using graphical arguments
 - $\vec{w}_{reg} \propto -\nabla_{\vec{w}} E_{in}(\vec{w}_{reg})$



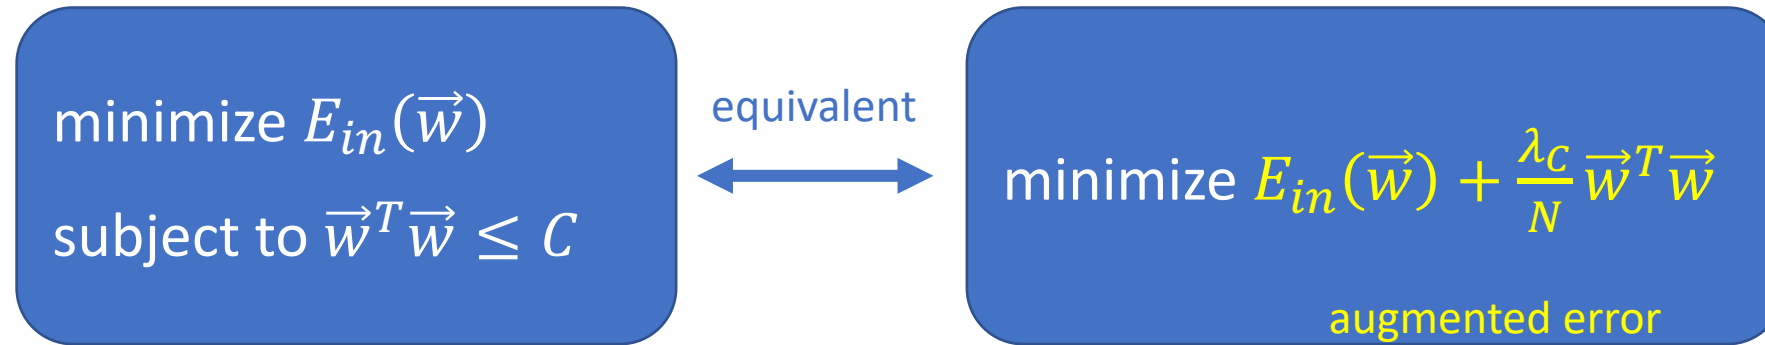
minimize $E_{in}(\vec{w})$ subject to $\vec{w}^T \vec{w} \leq C$

- When C is not large enough
 - Using graphical arguments
 - $\vec{w}_{reg} \propto -\nabla_{\vec{w}} E_{in}(\vec{w}_{reg})$
 - That is, we can find some constant $\lambda_C \geq 0$ such that
 - $\nabla_{\vec{w}} E_{in}(\vec{w}_{reg}) = -\frac{2\lambda_C}{N} \vec{w}_{reg}$
- Therefore,
 - $\nabla_{\vec{w}} \left(E_{in}(\vec{w}_{reg}) + \frac{\lambda_C}{N} \vec{w}_{reg}^T \vec{w}_{reg} \right) = 0$
- This implies, \vec{w}_{reg} is the solution for
 - minimize $E_{in}(\vec{w}) + \frac{\lambda_C}{N} \vec{w}^T \vec{w}$



Constrained to Unconstrained Optimization

- Constrained optimization \Leftrightarrow Unconstrained optimization



- Interpretations of regularization
 - Constraining H (by adding constraints)
 - Adding penalty to complex hypothesis in augmented errors

Augmented Error

- Define augmented error

- $E_{aug}(\vec{w}) = E_{in}(\vec{w}) + \frac{\lambda_C}{N} \vec{w}^T \vec{w}$

- Algorithm: Find $\vec{w}^* = \operatorname{argmin} E_{aug}(\vec{w})$

$\vec{w}^T \vec{w}$: weight decay

- A bit more discussion

- When $C \rightarrow \infty$, $\lambda_C = 0$
 - Smaller C (stronger constraints)
 - \Rightarrow larger λ_C
 - \Rightarrow smaller H
 - \Rightarrow stronger regularization
 - Use λ_C to tune the level of regularization

Side note:

You will see people/us interchangeably use λ_C and $\frac{\lambda_C}{N}$ to be the constant, depending on whether the dependency on N is emphasized.

General Form of Regularization

$$E_{aug}(h, \lambda, \Omega) = E_{in}(\vec{w}) + \frac{\lambda}{N} \Omega(h)$$

- Key parameters
 - Ω : Regularizer
 - λ : Amount of regularization
- Does the form look familiar: VC Theory
 - $E_{out}(g) \leq E_{in}(g) + O\left(\sqrt{d_{vc} \frac{\ln N}{N}}\right)$
- If we pick the right Ω , E_{aug} can be a better proxy for E_{out}

How to Pick the Right Ω

- No definite answer, but generally
 - We like to pick Ω that leads to “smoother” hypothesis
 - Overfitting is due to noise
 - Informally, noise is usually “high frequency”
 - We prefer Ω that makes the optimization easier (e.g., convex/differentiable)
 - Similar to picking the error measure
 - We might have some other objective in mind
 - Ex: L-1 regularizer leads to weight vectors with more 0s
 - $E_{aug}(\vec{w}) = E_{in}(\vec{w}) + \lambda \|\vec{w}\|_1 = E_{in}(\vec{w}) + \lambda \sum_i |w_i|$
- What if we pick the wrong Ω (Think about weight growth)
 - We might still fix it by picking the right λ – using validation in the next lecture

More Discussion on Regularization

Why $\vec{w}^T \vec{w}$ is Called Weight Decay

- Run gradient descent on $E_{aug}(\vec{w}) = E_{in}(\vec{w}) + \lambda_c \vec{w}^T \vec{w}$
- The update rule would be

$$\vec{w}(t+1) \leftarrow \vec{w}(t) - \eta \nabla_{\vec{w}} E_{aug}(\vec{w}(t))$$

$$\Rightarrow \vec{w}(t+1) \leftarrow (1 - 2\eta\lambda_c) \vec{w}(t) - \eta \nabla_{\vec{w}} E_{in}(\vec{w}(t))$$

We are **decaying** the weights first, then do the update

Linear Regression with Weight Decay

- $E_{aug}(\vec{w}) = E_{in}(w) + \frac{\lambda_C}{N} \vec{w}^T \vec{w} = \frac{1}{N} \|X\vec{w} - \vec{y}\|^2 + \frac{\lambda_C}{N} \vec{w}^T \vec{w}$
- Solve $\nabla_{\vec{w}} E_{aug}(\vec{w})|_{\vec{w}=\vec{w}_{reg}} = 0$, we get
 - $\frac{2}{N} (X^T X \vec{w}_{reg} - X^T \vec{y} + \lambda_C \vec{w}_{reg}) = 0$
 - $(X^T X + \lambda_C I) \vec{w}_{reg} = X^T \vec{y}$
 - $\vec{w}_{reg} = (X^T X + \lambda_C I)^{-1} X^T \vec{y}$

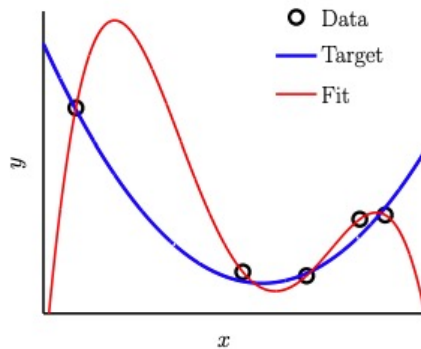
Notation: I is an identity matrix: only the elements in the diagonals are 1, and all others are 0.

This is called “Ridge Regression” in statistics.

Effect of Regularization (Different λ)

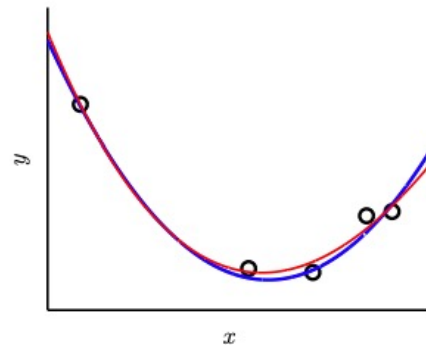
- Minimizing $E_{in}(\vec{w}) + \frac{\lambda}{N} \vec{w}^T \vec{w}$ with different λ

$$\lambda = 0$$

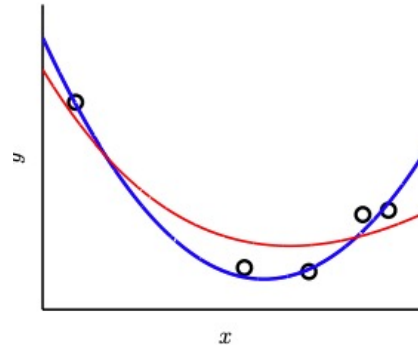


Overfitting

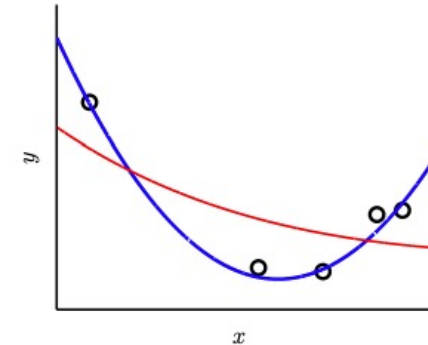
$$\lambda = 0.0001$$



$$\lambda = 0.01$$

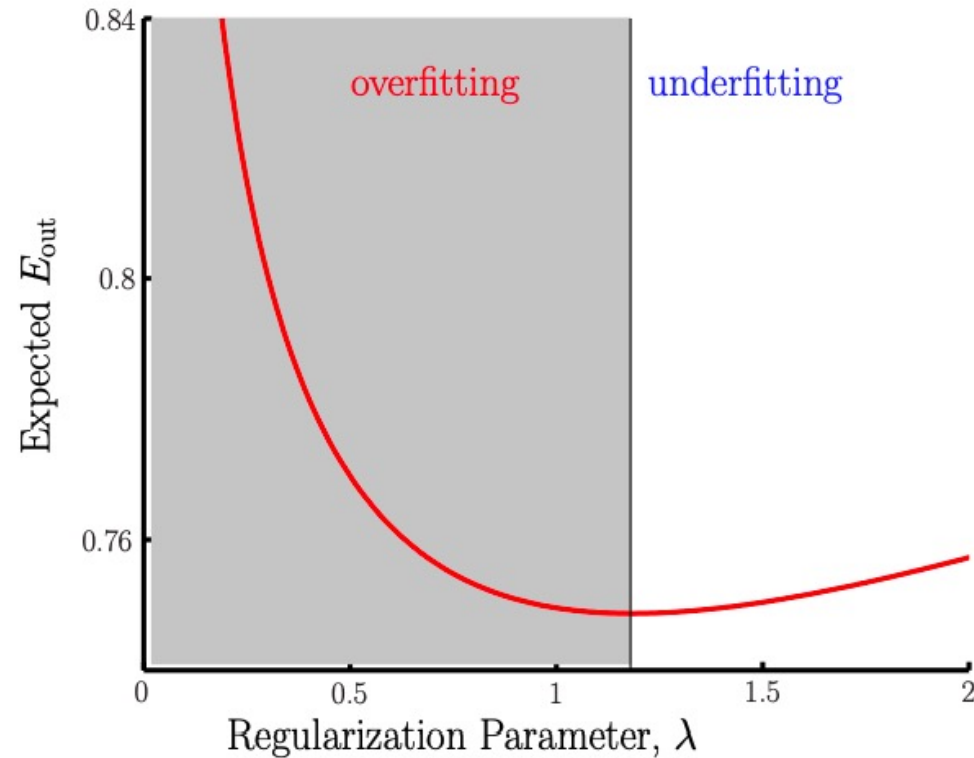


$$\lambda = 1$$

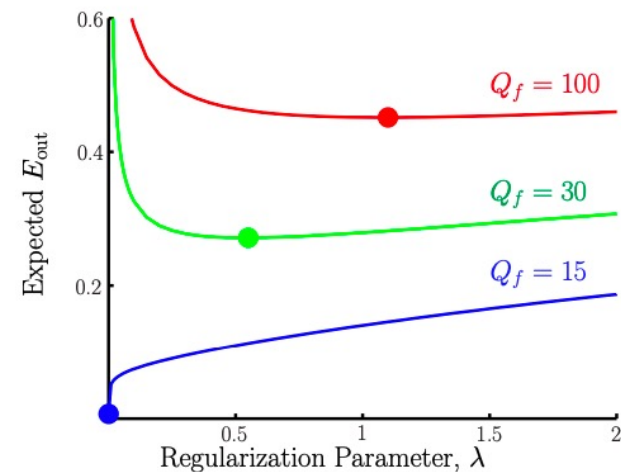
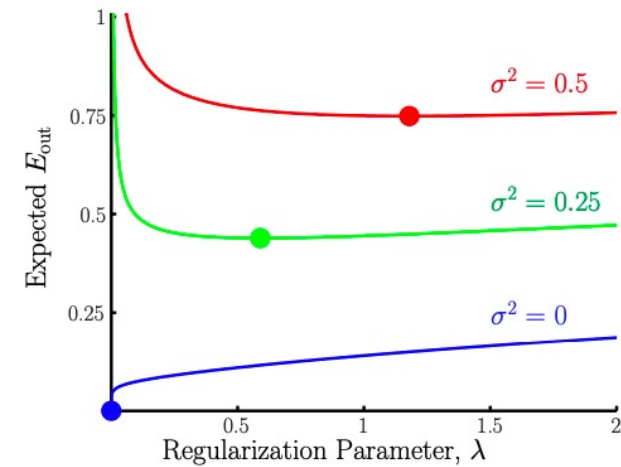


Underfitting

Overfitting and Underfitting

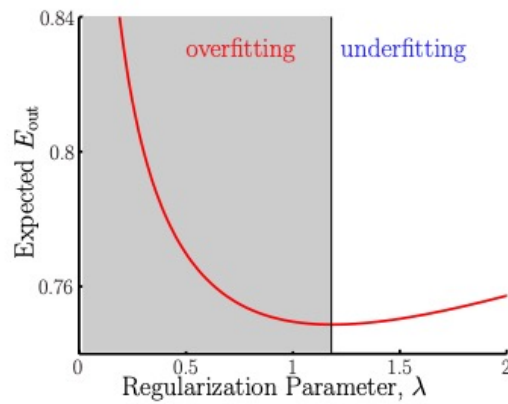


Need to pick the right λ :
Using validation: Focus of next lecture



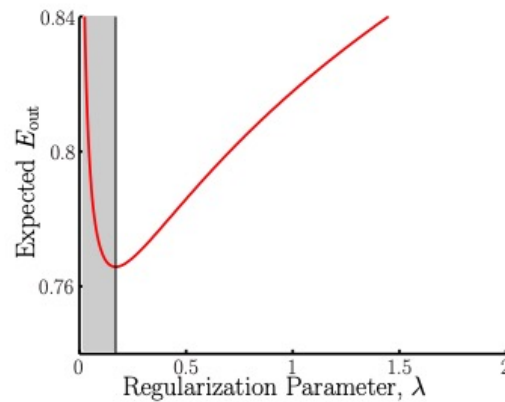
Variations on Weight Decay (Different Ω)

Uniform Weight Decay



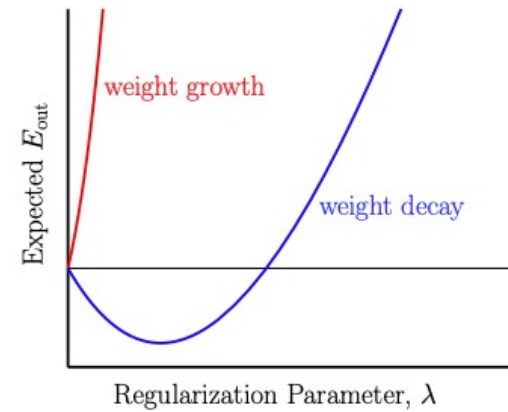
$$\sum_{q=0}^Q w_q^2$$

Low Order Fit



$$\sum_{q=0}^Q q w_q^2$$

Weight Growth!



$$\sum_{q=0}^Q \frac{1}{w_q^2}$$

How to Pick the Right Ω

- As discussed earlier
 - Intuition: pick Ω that leads to “smoother” hypothesis
 - Overfitting is due to noise
 - Informally, noise is generally “high frequency”
 - Computation: prefer Ω that makes the optimization easier (e.g., convex/differentiable)
 - Similar to picking the error measure
 - We might have some other objective in mind
 - Ex: L-1 regularizer leads to weight vectors with more 0s
- What if we pick the **wrong Ω** (weight growth)
 - We might still fix it by picking the right λ – using **validation**

Summarizing Regularization

- Regularization is **everywhere** in machine learning
- Two main ways of thinking about regularization
 - **Constraining H** to make overfitting less likely to happen
 - Will discuss more regularization methods in the 2nd half of the semester
 - Pruning for decision trees, early stopping / dropout for neural networks, etc
 - Define **augmented error** E_{aug} to better approximate E_{out}
 - $E_{aug}(h, \lambda, \Omega) = E_{in}(\vec{w}) + \frac{\lambda}{N} \Omega(h)$
- We show the **equivalence** of the two for weight decay
 - The conceptual equivalence is general with Lagrangian relaxation (will cover later in the semester)

Validation

Prevent Overfitting

$$E_{out}(g) = E_{in}(g) + \text{overfit penalty}$$

- Regularization
 - Choose a regularizer Ω to approximate the penalty
- Validation
 - Directly estimate E_{out} (The real goal of learning is to minimize E_{out})

Review: Test Set

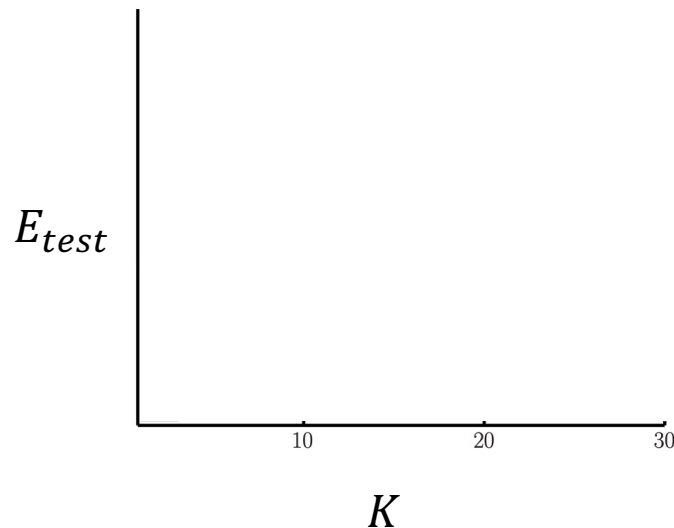
- Out-of-sample error $E_{out}(g) = \mathbb{E}_{\vec{x}}[e(g(\vec{x}), y)]$
 - Key: \vec{x} is **out of sample**
- Test set $D_{test} = \{(\vec{x}_1, y_1), \dots, (\vec{x}_K, y_K)\}$
 - Reserve K data points used to estimate E_{out}
 - **None** of the data points in **test set** can be **involved in training**
- Using the data in test set to estimate E_{out}
 - Since all data points in D_{test} are **out of sample**

Test Set

- Test set $D_{test} = \{(\vec{x}_1, y_1), \dots, (\vec{x}_K, y_K)\}$
- For a g learned using **only training set**
- Let $E_{test}(g) = \frac{1}{K} \sum_{k=1}^K e(g(\vec{x}_k), y_k)$
 - $E_{test}(g)$ is an **unbiased** estimate of $E_{out}(g)$
 - $\mathbb{E}[E_{test}(g)] = \frac{1}{K} \sum_{k=1}^K \mathbb{E}[e(g(\vec{x}_k), y_k)] = E_{out}(g)$
 - **Single hypothesis** Hoeffding bound applies
 - $E_{out}(g) \leq E_{test}(g) + O\left(\sqrt{\frac{1}{K}}\right)$

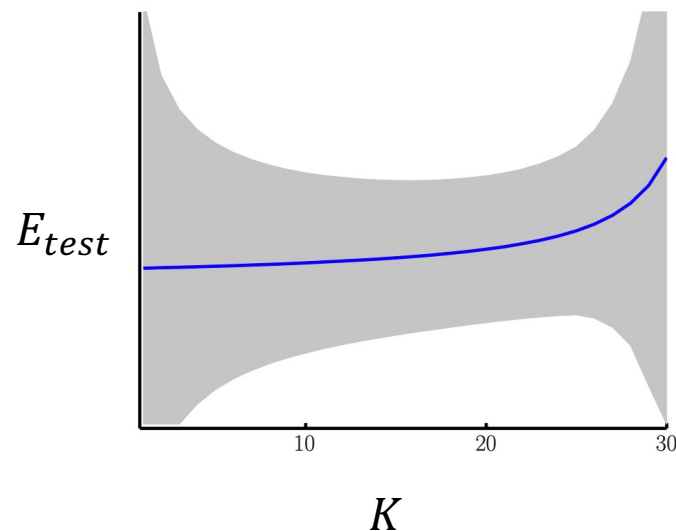
Where are Test Set From?

- Given a data set D of N points
 - $D = D_{train} \cup D_{test}$
 - Reserving K points for test set means we only have $N - K$ points for training
- Effect of the choice of K



Where are Test Set From?

- Given a data set D of N points
 - $D = D_{train} \cup D_{test}$
 - Reserving K points for test set means we only have $N - K$ points for training
- Effect of the choice of K



Rule of Thumb: $K^* = \frac{N}{5}$

Utilizing the Whole D

- Process:
 - $D = D_{train} \cup D_{test}$ where $|D_{test}| = K, |D_{train}| = N - K$
 - Learn some hypothesis g^- using only D_{train}
 - Estimate $E_{out}(g^-)$ using D_{test}
 - Let g be the hypothesis that would be learned using D
- Generally (informally, not theoretically proven)
 - Training on more data leads to better learned hypothesis
 - $E_{out}(g) \leq E_{out}(g^-)$

Validation: Beyond Test Set

- What if we want to estimate E_{out} multiple times?
- Model selection:
 - Should I use linear models or decision trees?
 - Should I set the regularization parameter λ to 0.1, 0.01, or 0.001?
 - A model with different λ can be considered as different model
- Validation set
 - $D = D_{train} \cup D_{val}$
 - Key difference: We need to **account for** the multiple usage of D_{val}