

CSE 417T

Introduction to Machine Learning

Lecture 13

Instructor: Chien-Ju (CJ) Ho

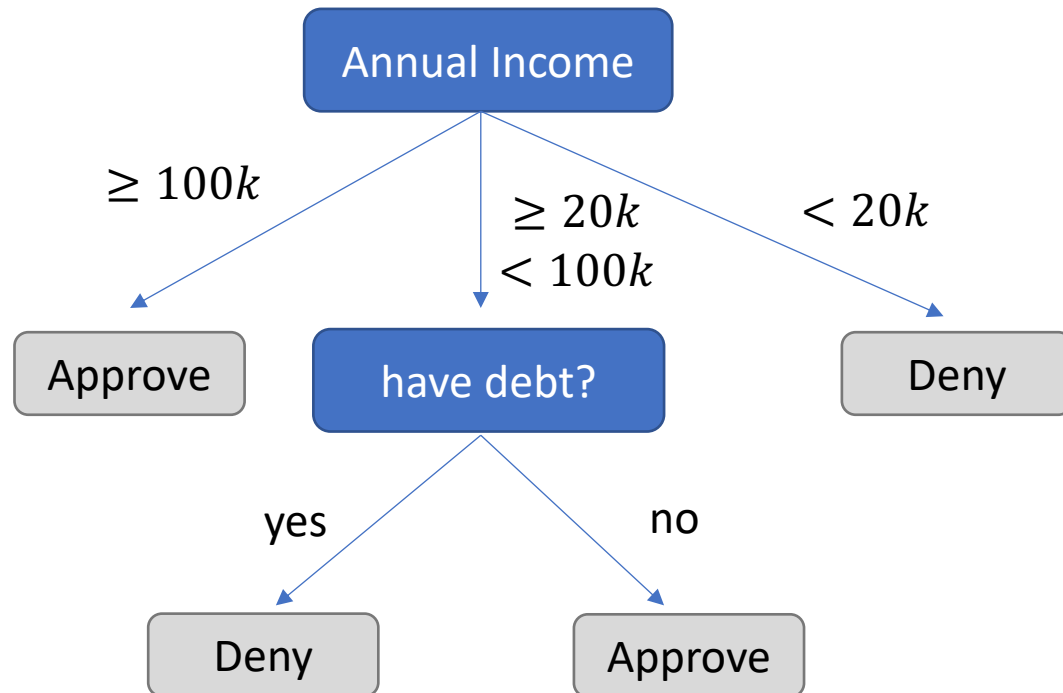
Logistics

- Homework 3: due on **Oct 19 (Wednesday)**
 - Track your own late-day usages
 - Assignments over-using late days won't be graded
- Exam 1: **October 27 (Thursday)**
 - Topics: LFD Chapters 1 to 5
 - Timed exam (75 min) during lecture time
 - Location TBD
 - Closed-book exam with 2 letter-size cheat sheets allowed (4 pages in total)
 - No format limitations (it can be typed, written, or a combination)
- October 25 (Tuesday) will be a review session

Recap

Decision Tree

Decision Tree Hypothesis



Credit Card Approval Example

- Pros
 - Easy to interpret (interpretability is getting attention and is important in many domains)
 - Can handle multi-type data (Numerical, categorical. ...)
 - Easy to implement (Bunch of if-else rules)
- Cons
 - Generally speaking, **bad generalization**
 - VC dimension is infinity
 - High variance (small change of data leads to very different hypothesis)
 - Easily overfit
- Why we care?
 - One of the classical model
 - Building block for other models (e.g., random forest)


Learning Decision Tree from Data

- Given dataset D , how to learn a decision tree hypothesis?

| x_1 | x_2 | x_3 | y |
|-------|-------|-------|-----|
| +1 | +1 | +1 | +1 |
| +1 | +1 | -1 | +1 |
| +1 | -1 | +1 | +1 |
| +1 | -1 | -1 | +1 |
| -1 | +1 | +1 | +1 |
| -1 | +1 | -1 | +1 |
| -1 | -1 | +1 | -1 |
| -1 | -1 | -1 | -1 |

- Potential approach:
 - Empirical risk minimization
 - Find $g = \operatorname{argmin}_{h \in H} E_{in}(h)$

- Multiple decision trees with zero E_{in}



How to avoid overfitting?

Learning Decision Tree from Data

- Conceptual intuition to deal with overfitting
 - Regularization: **Constrain H**
- Informally,

$$\begin{array}{ll} \text{minimize} & E_{in}(\vec{w}) \\ \text{subject to} & \text{size}(\text{tree}) \leq C \end{array}$$
- This optimization is generally computationally intractable.
- Most decision tree learning algorithms rely on **heuristics** to approximate the goal.

Greedy-Based Decision Tree Algorithm

- Greedily, recursively, choose the next feature to split
- DecisionTreeLearn(D): Input a dataset D , output a decision tree hypothesis
 - Create a root node
 - If **termination conditions** are met
 - return a single node tree with **leaf prediction** based on D
 - Else: Greedily find a feature A to split according to **split criteria**
 - For each possible value v_i of A
 - Let D_i be the dataset containing data with value v_i for feature A
 - Create a subtree DecisionTreeLearn(D_i) that being the child of root
- Most decision tree learning algorithms follow this template, but with different choices of **heuristics**

ID3: Using Information Gain as Selection Criteria

- Information gain of choosing feature A to split
 - $Gain(D, A) = H(D) - \sum_i \frac{|D_i|}{|D|} H(D_i)$ [The amount of decrease in entropy]
- ID3: Choose the split that maximize $Gain(D, A)$

Notations:

$H(D)$: Entropy of D

$|D|$ is the number of points in D

DecisionTreeLearn(D)

Create a root node r

If **termination conditions** are met

return a single node tree with **leaf prediction** based on

Else: Greedily find a feature A to split according to **split criteria**

For each possible value v_i of A

Let D_i be the dataset containing data with value v_i for feature A

Create a subtree DecisionTreeLearn(D_i) that being the child of root r

- ID3 termination conditions
 - If all labels are the same
 - If all features are the same
 - If dataset is empty
- ID3 leaf predictions
 - Most common labels (majority voting)
- ID3 split criteria
 - Information gain

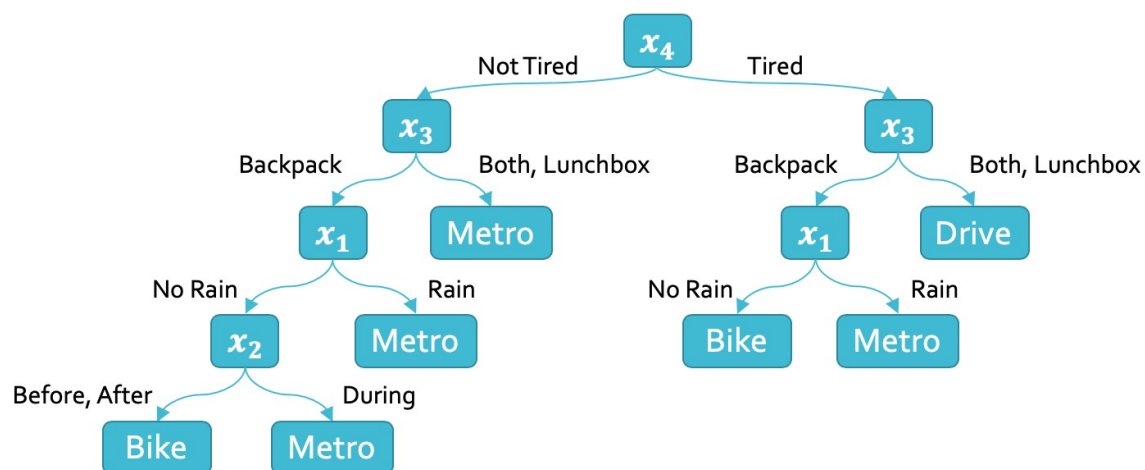
Illustration of "High Variance" of Decision Trees

| x_1 | x_2 | x_3 | x_4 | y |
|---------|--------|----------|-----------|-------|
| Rain | Before | Both | Tired | Drive |
| Rain | During | Both | Not Tired | Metro |
| Rain | During | Both | Tired | Drive |
| Rain | After | Backpack | Not Tired | Metro |
| Rain | After | Backpack | Tired | Metro |
| Rain | After | Lunchbox | Tired | Drive |
| No Rain | Before | Backpack | Tired | Bike |
| No Rain | Before | Lunchbox | Not Tired | Metro |
| No Rain | Before | Lunchbox | Tired | Drive |
| No Rain | During | Backpack | Not Tired | Metro |
| No Rain | During | Both | Tired | Drive |
| No Rain | After | Backpack | Not Tired | Bike |
| No Rain | After | Backpack | Tired | Bike |
| No Rain | After | Both | Not Tired | Metro |
| No Rain | After | Both | Tired | Drive |
| No Rain | After | Lunchbox | Not Tired | Metro |

| x_1 | x_2 | x_3 | x_4 | y |
|---------|--------|----------|-----------|-------|
| Rain | Before | Both | Tired | Drive |
| Rain | During | Both | Not Tired | Drive |
| Rain | During | Both | Tired | Drive |
| Rain | After | Backpack | Not Tired | Metro |
| Rain | After | Backpack | Tired | Metro |
| Rain | After | Lunchbox | Tired | Drive |
| No Rain | Before | Backpack | Tired | Bike |
| No Rain | Before | Lunchbox | Not Tired | Metro |
| No Rain | Before | Lunchbox | Tired | Drive |
| No Rain | During | Backpack | Not Tired | Metro |
| No Rain | During | Both | Tired | Drive |
| No Rain | After | Backpack | Not Tired | Bike |
| No Rain | After | Backpack | Tired | Bike |
| No Rain | After | Both | Not Tired | Metro |
| No Rain | After | Both | Tired | Drive |
| No Rain | After | Lunchbox | Not Tired | Metro |


Illustration of "High Variance" of Decision Trees

| x_1 | x_2 | x_3 | x_4 | y |
|-------|--------|-------|-----------|-------|
| Rain | Before | Both | Tired | Drive |
| Rain | During | Both | Not Tired | Metro |



| | | | | |
|---------|-------|----------|-----------|-------|
| No Rain | After | Backpack | Not Tired | Bike |
| No Rain | After | Backpack | Tired | Bike |
| No Rain | After | Both | Not Tired | Metro |
| No Rain | After | Both | Tired | Drive |
| No Rain | After | Lunchbox | Not Tired | Metro |

| x_1 | x_2 | x_3 | x_4 | y |
|-------|--------|-------|-----------|-------|
| Rain | Before | Both | Tired | Drive |
| Rain | During | Both | Not Tired | Drive |



| | | | | |
|---------|-------|----------|-----------|-------|
| No Rain | After | Backpack | Not Tired | Bike |
| No Rain | After | Backpack | Tired | Bike |
| No Rain | After | Both | Not Tired | Metro |
| No Rain | After | Both | Tired | Drive |
| No Rain | After | Lunchbox | Not Tired | Metro |

High variance: A small deviation of data would lead to very different learned hypothesis

Decision Tree Hypothesis



Credit Card Approval Example

- Pros
 - Easy to interpret (interpretability is getting attention and is important in some domains)
 - Can handle multi-type data (Numerical, categorical. ...)
 - Easy to implement (Bunch of if-else rules)
- Cons
 - Generally speaking, **bad generalization**
 - VC dimension is infinity
 - **High variance** (small change of data leads to very different hypothesis)
 - Easily overfit
- Why we care?
 - One of the classical model
 - **Building block for other models**

Today's Lecture

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.
Let me know if you spot errors.

Ensemble Learning

Ensemble Learning

- Assume we are given a set of learned hypothesis
 - g_1, g_2, \dots, g_M
- What can we do?
 - Use validation to pick the best one
 - What if all of them are not good enough
- Can we **aggregate** them?

An Anecdote about Aggregation



- At a 1906 country fair, ~800 people participate in a contest to guess the weight of an ox.
- Reward is given to the person with the closest guess.
- The average guess is 1,197lbs.
The true answer is 1,198lbs.

How to do Aggregation

- Given a set of **weak learners** g_1, \dots, g_M , how to output a **stronger learner** that performs better?
- Uniform aggregation
 - Regression (average): $\bar{g}(\vec{x}) = \frac{1}{M} \sum_{m=1}^M g_m(\vec{x})$
 - Classification (majority vote): $\bar{g}(\vec{x}) = \text{sign} \left(\frac{1}{M} \sum_{m=1}^M g_m(\vec{x}) \right)$
- Weighted aggregation
 - Regression (average): $\bar{g}(\vec{x}) = \frac{1}{M} \sum_{m=1}^M \alpha_m g_m(\vec{x})$
 - Classification (majority vote): $\bar{g}(\vec{x}) = \text{sign} \left(\frac{1}{M} \sum_{m=1}^M \alpha_m g_m(\vec{x}) \right)$
- Stacking (won't talk about this in this course)
 - Take the prediction of g_1 to g_m as input features, train another model on top of that

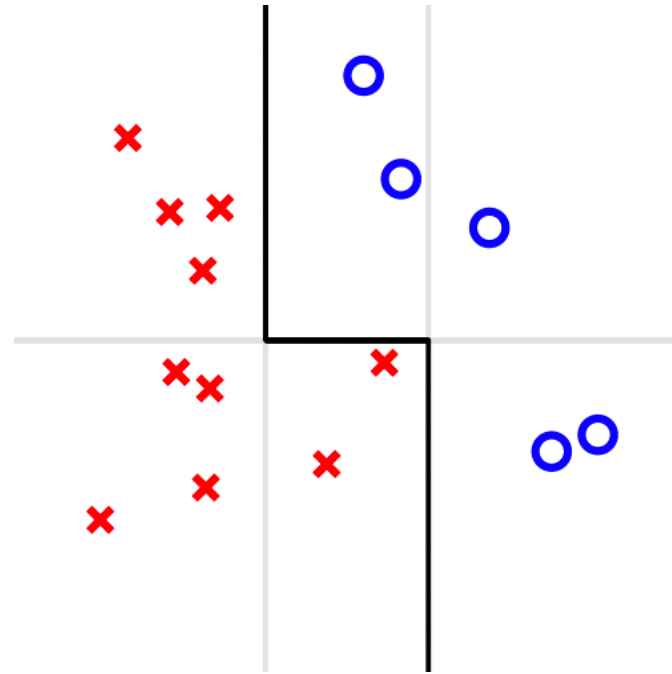
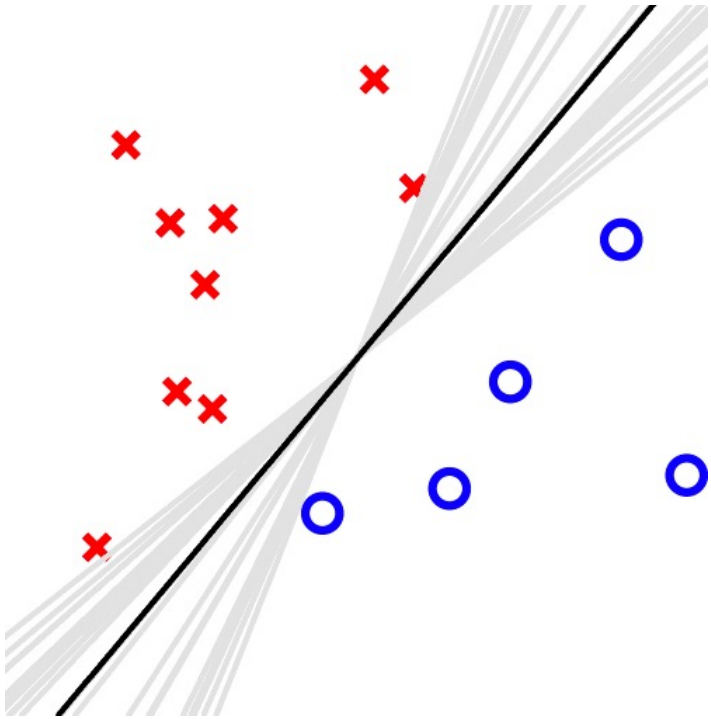
Mathematically, majority voting and average is similar with +1/-1 labels

Is Aggregation a Good Idea?

- Some illustrative examples

Is Aggregation a Good Idea?

- Some illustrative examples



Is Aggregation a Good Idea?

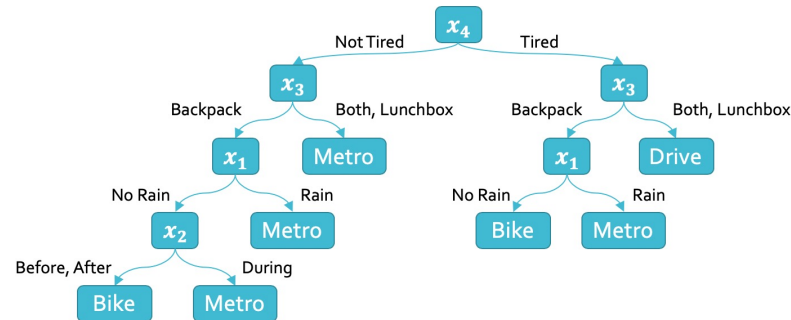
- Maybe
 - If the hypothesis is **diverse**, and “on average” they seem good
 - (If you take humans as weak learners, this is almost democracy)
- Question:
 - How do we **find** a set of hypothesis that are **diverse** and “on average” good
 - How do we **aggregate** the set of hypothesis
- Ensemble learning
 - Bagging – Random Forest (This lecture)
 - Boosting – AdaBoost (Next lecture)

Diverse Weak Learners

- One common way to construct weak learners is via **decision trees**


- Fully grown decision trees

- High variance
- Low bias



- Decision stump (One-depth decision trees, split on only one attribute)

- Low variance
- High bias




Bagging

Bootstrapped Aggregating


(Using randomization to construct diverse weak learners)

Review: Bias-Variance Decomposition


- f : sine function, $H: h(x) = ax + b, N=2$



For each dataset,
learn a hypothesis.



Draw many datasets,
learn many hypothesis



Take average.

- Observations

- The variance of each learned hypothesis is high
- The variance of “average” of them ($\bar{g}(\vec{x})$) is lower

- Can we apply similar intuitions?

- Generate a lot of **high-variance but low bias** weak learners
- Aggregate them using uniform aggregation

We only have one
dataset in practice!

Bootstrapping

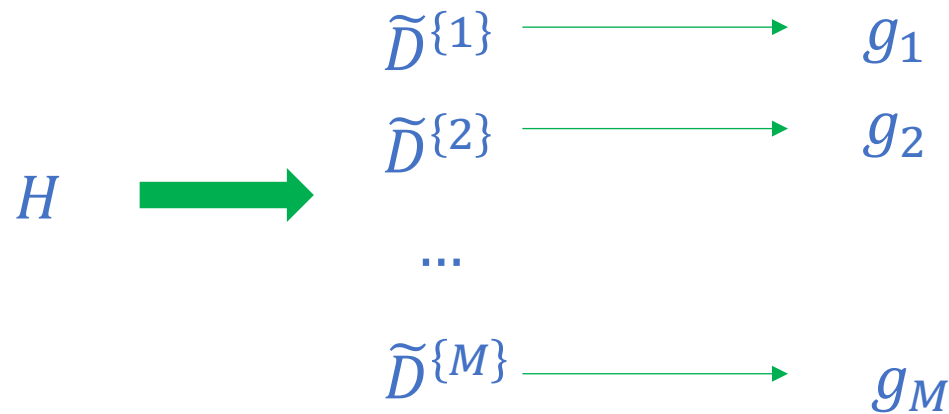
- Intuition:
 - Use the dataset D we have to approximate the data distribution
 - Sample (with replacement) from D to create bootstrapped datasets
- Bootstrapping:
 - Let $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$ be the dataset we have
 - Repeatedly uniformly sample N points from D with replacement
 - The number of sampled points doesn't have to be N , but it's a reasonable/common choice.
 - Obtain many bootstrapped datasets
 - $\tilde{D}^{(1)} = \{(x_1, y_1), (x_1, y_1), (x_4, y_4), \dots\}$
 - \dots
 - $\tilde{D}^{(M)}$

Bagging - Bootstrapped Aggregating

- Bootstrap M datasets $\{\tilde{D}^{\{m\}}\}$ and learn a hypothesis from each of them
- Aggregate the learned hypothesis

Bagging - Bootstrapped Aggregating

- Bootstrap M datasets $\{\tilde{D}^{(m)}\}$ and learn a hypothesis from each of them



- Aggregate the learned hypothesis

$$G(\vec{x}) = \bar{g}(\vec{x}) = \text{sign} \left(\frac{1}{M} \sum_{m=1}^M g_m(\vec{x}) \right) \quad (\text{assume we are doing classification})$$

Why/When Bagging Might Be Helpful?

- What we know from statistics
 - Consider M independent random variables x_1, x_2, \dots, x_M , each with variance σ^2
 - The variance of $\frac{1}{M} \sum_{m=1}^M x_m$ is $\frac{\sigma^2}{M}$
- If we have “weak learners” that have **high variance** but **low bias**
 - Bagging helps **reduce the variance** and **maintain low bias**
 - From bias-variance decomposition, this implies a strong learner

Break and Question

Exercise:

Given a dataset D with N points.

Consider we bootstrap a dataset $\tilde{D}^{(m)}$ by sampling N points with replacement from D , what's the probability that a given point (\vec{x}_n, y_n) is not in $\tilde{D}^{(m)}$?

Out-Of-Bag (OOB) Error

Probability for a Point to be Out of Bag

- Consider we bootstrap a dataset $\tilde{D}^{(m)}$ by sampling N points from D , what's the probability that a given point (\vec{x}_n, y_n) is not in $\tilde{D}^{(m)}$.

$$\begin{aligned} & \left(1 - \frac{1}{N}\right)^N \\ &= \left(\frac{1}{1 + \frac{1}{N-1}}\right)^N \\ &\approx \frac{1}{e} \approx 0.36 \text{ when } N \rightarrow \infty \end{aligned}$$

When N is large, for each bootstrapped dataset $\tilde{D}^{(m)}$, a significant proportion of points in D is not included.

- A point that is not in $\tilde{D}^{(m)}$ is not involved in training g_m
 - Can we utilize it to **validate** the performance of g_m ?
 - Yes, but we care about the overall performance, not just the performance of g_m ...

Out-Of-Bag (OOB) Error

| | $\tilde{D}^{(1)}$ | $\tilde{D}^{(2)}$ | $\tilde{D}^{(3)}$ | $\tilde{D}^{(4)}$ | ... |
|--------------------|-------------------|-------------------|-------------------|-------------------|-----|
| (\vec{x}_1, y_1) | Yes | Yes | No | No | ... |
| (\vec{x}_2, y_2) | Yes | No | No | No | ... |
| ... | ... | ... | ... | ... | ... |
| (\vec{x}_N, y_N) | No | Yes | Yes | Yes | ... |

Whether a point is in a bootstrapped dataset

- Recall that we learn g_1, \dots, g_M using $\tilde{D}^{(1)}, \dots, \tilde{D}^{(M)}$
- Which set of hypothesis can (\vec{x}_1, y_1) be used for **validation**?

Out-Of-Bag (OOB) Error

| | $\tilde{D}^{(1)}$ | $\tilde{D}^{(2)}$ | $\tilde{D}^{(3)}$ | $\tilde{D}^{(4)}$ | ... |
|--------------------|-------------------|-------------------|-------------------|-------------------|-----|
| (\vec{x}_1, y_1) | Yes | Yes | No | No | ... |
| (\vec{x}_2, y_2) | Yes | No | No | No | ... |
| ... | ... | ... | ... | ... | ... |
| (\vec{x}_N, y_N) | No | Yes | Yes | Yes | ... |

Whether a point is in a bootstrapped dataset

- G_n^- : the aggregation of hypothesis that \vec{x}_n is OOB of

- $G_1^- = \text{aggregate}(g_3, g_4, \dots)$
- $G_2^- = \text{aggregate}(g_2, g_3, g_4, \dots)$
- $G_N^- = \text{aggregate}(g_1, \dots)$

Aggregate:
Majority voting for classification
Average for regression

- OOB Error

- $E_{OOB}(G) = \frac{1}{N} \sum_{n=1}^N \text{error}(G_n^-(\vec{x}_n), y_n)$

Error:
Binary error for classification
Squared error for regression

Out-Of-Bag (OOB) Error

$$E_{OOB}(G) = \frac{1}{N} \sum_{n=1}^N \text{error}(G_n^-(\vec{x}_n), y_n)$$

- Bagging provided an **intrinsic** mechanism for us to perform validation
 - We don't need to split the dataset into training and validation
- Practical issues (you might face this in HW4)
 - What if some \vec{x}_n appears in all bootstrapped datasets?
 - The probability of this happening is small when the number of bags M is large
 - Let S be the set of points that is out of bag for at least one bootstrapped dataset
 - $E_{OOB}(G) = \frac{1}{|S|} \sum_{(\vec{x}_n, y_n) \in S} \text{error}(G_n^-(\vec{x}_n), y_n)$

Random Forest

What We Have Learned

Bagging:

A method to generate and aggregate many **high-variance** weak learners into a stronger one.



Decision tree:

Various nice properties
Bad generalization
- Due to **high variance**



Random Forest:

1. Construct many random trees
2. Aggregate the random trees

Random Forest

- Construct many random trees
 - Bootstrapping datasets and learn a **max-depth tree** for each of them
 - Other randomizations (not required in HW4)
 - When choosing split features, choose from a random subset (instead of all features)
 - Randomly project features (similar to non-linear transformation) for each tree
- Aggregate the random trees
 - Classification: Majority vote $\bar{g}(\vec{x}) = \text{sign} \left(\frac{1}{M} \sum_{m=1}^M g_m(\vec{x}) \right)$
 - Regression: Average $\bar{g}(\vec{x}) = \frac{1}{M} \sum_{m=1}^M g_m(\vec{x})$

Questions?

- Note that in HW4, you will be asked to implement Bagging Decision Tree and calculate the OOB errors.
- Make sure you know the definitions/algorithm well.

Brief Discussion on Feature Importance

- Not all features are equally important
 - Some features could be **redundant** -- (birth year, age)
 - Some features might be **irrelevant** -- feature: name, label: prob of heart attack
- How do we know which features are more important?
 - Linear models:
 - The **size of the weight** is a proxy for feature importance
 - Applying L1 regularization is one way to reduce the number of features
 - Decision tree:
 - The feature **closer to the root** is probably more important
 - Random forest:
 - **Average “information gain”** of all trees is a proxy for feature importance
- See LFD e-Chap 9.2 for more discussion on feature selection

Boosting

Ensemble Learning

- Goal: Utilize a set of **weak learners** to obtain a **strong learner**
- Format of ensemble learning
 - **Construct** many **diverse** weak learners
 - **Aggregate** the weak learners

Bagging:

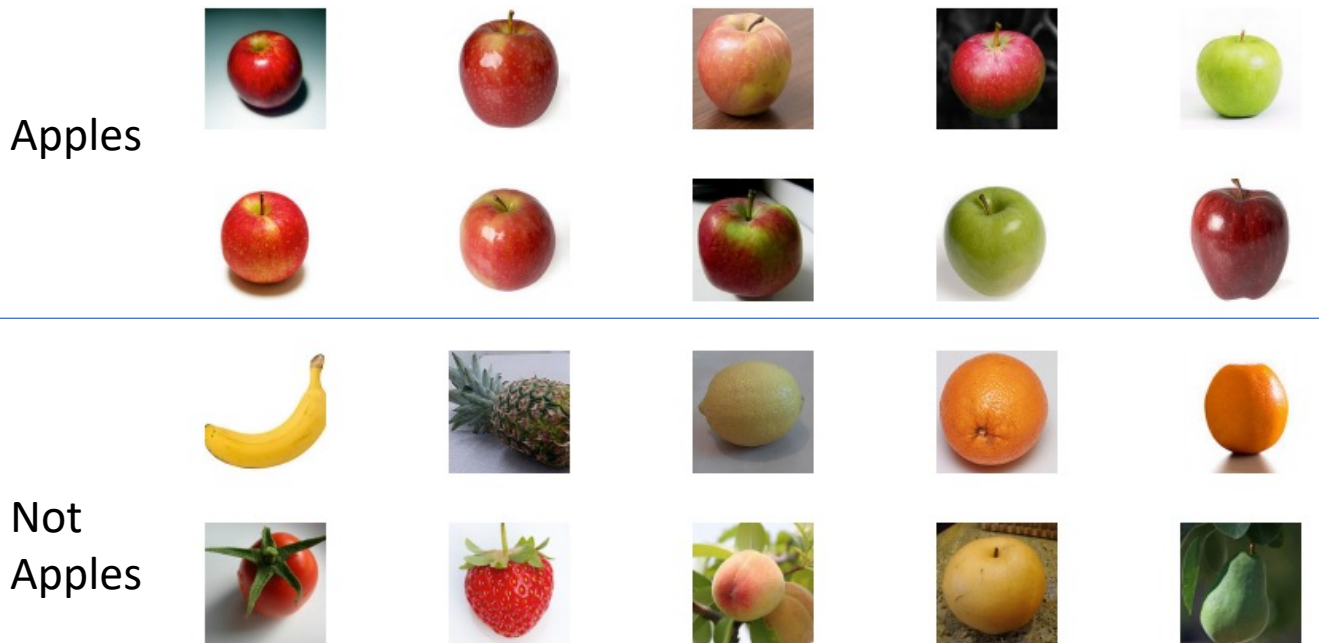
- Construct diverse weak learners
 - (**Simultaneously**) bootstrapping datasets
 - Train weak learners on them
- Aggregate the weak learners
 - **Uniform** aggregation

Boosting

- Construct diverse weak learners
 - **Adaptively** generating datasets
 - Train weak learners on them
- Aggregate the weak learners
 - **Weighted** aggregation

Informal Intuitions about Boosting

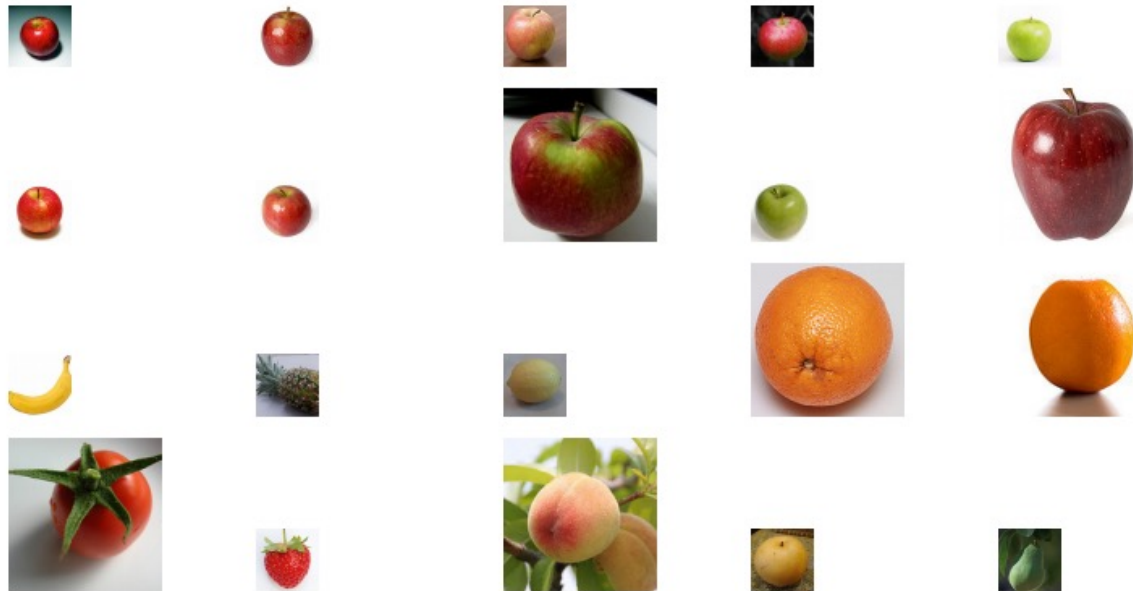
- Example : Teach a class of kids to identify apples from data



- Alice: Apples are **circular**
- Teacher:
Circular is a good feature, but using this feature might make some mistakes
Let me **highlight** the mistakes.
 - Make correct images smaller
 - Make incorrect images larger

Informal Intuitions about Boosting

- Example : Teach a class of kids to identify apples from data



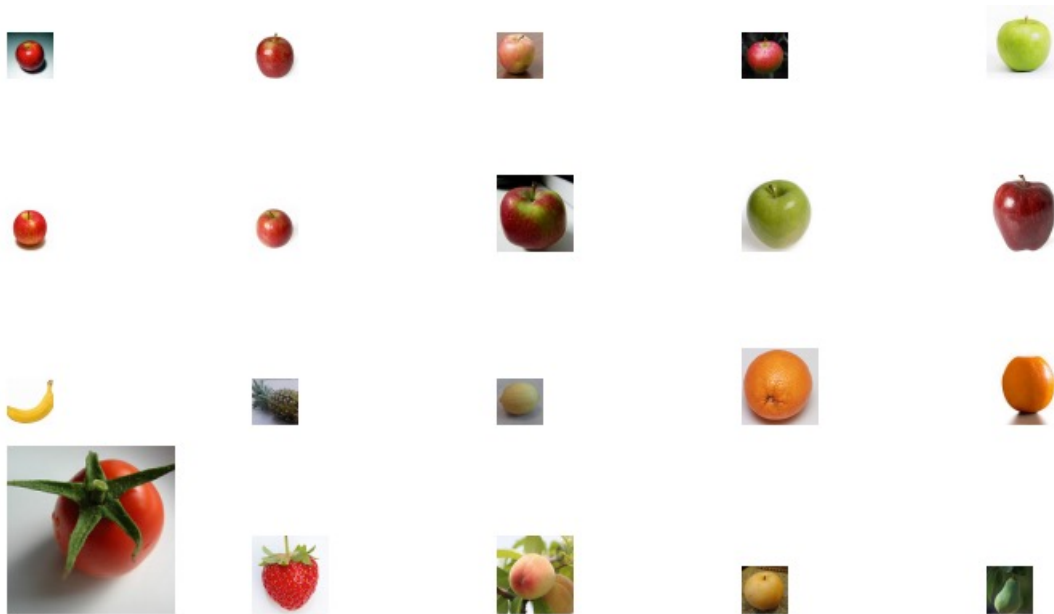
- Alice: Apples are **circular**
- Bob: Apples are **red**
- Teacher:
Yes, many apples are red but it could still make mistakes.

Let me **highlight** the mistakes.

- Make correct images smaller
- Make incorrect images larger

Informal Intuitions about Boosting

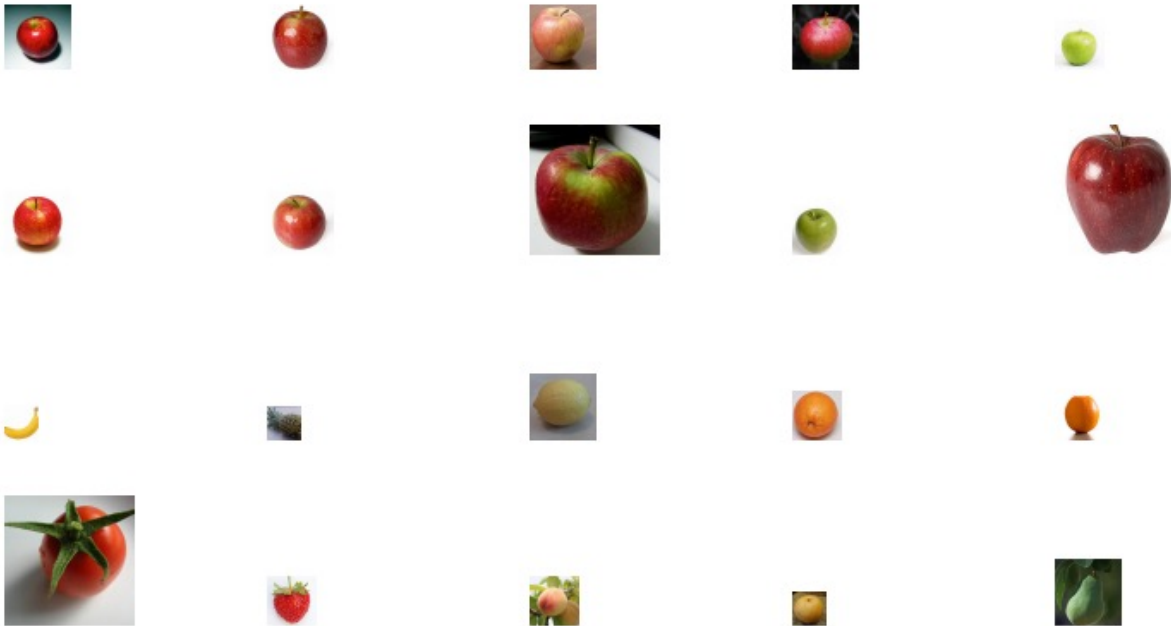
- Example : Teach a class of kids to identify apples from data



- Alice: Apples are **circular**
- Bob: Apples are **red**
- Charlie: Apples could be **green**

Informal Intuitions about Boosting

- Example : Teach a class of kids to identify apples from data



- Alice: Apples are **circular**
- Bob: Apples are **red**
- Charlie: Apples could be **green**
- David: Apples have **stems** at the top
- Class: Apples are **somewhat circular, somewhat red, possibly green, and may have stems at the top**

Informal Intuitions about Boosting

- Example : Teach a class of kids to identify apples from data



Key steps of this process:

- Learn a **simple** hypothesis for each dataset
- Iteratively update the dataset to focus on what we got wrong (i.e., create **diversity**)
- **Aggregate** the learned simple hypothesis



- Alice: Apples are **circular**
- Bob: Apples are **red**
- Charlie: Apples could be **green**
- David: Apples have **stems** at the top
- Class: Apples are **somewhat circular, somewhat red, possibly green, and may have stems at the top**

Outline of a Boosting Algorithm

- Initialize D_1 (usually the same as the initial dataset D)
- For $t = 1$ to T
 - Learn g_t from D_t
 - Reweight the distribution and obtain D_{t+1} based on g_t and D_t
- Output **weighted**-aggregate(g_1, \dots, g_T)
 - Classification: $G(\vec{x}) = \bar{g}(\vec{x}) = \text{sign}\left(\frac{1}{T} \sum_{t=1}^T \alpha_t g_t(\vec{x})\right)$

Questions

How to learn g_t from D_t

How to reweight the distribution and obtain D_{t+1}

How to perform weighted aggregation

Discussion on Re-weighted D_t (What does re-weighting mean?)

- Original Dataset $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$
- Notation of D_t
 - $D_t(n)$ is the weight/probability of data point (\vec{x}_n, y_n) in D_t
 - $\sum_{n=1}^N D_t(n) = 1$
- What is $E_{in}(h)$ on D_t ? (Expressed as $E_{in}^{(D_t)}(h)$)
 - Re-sample dataset (noisier)
 - Re-sample the dataset from D according to distribution D_t
 - Calculate E_{in} on the re-sampled dataset as usual
 - Calculate weighted error
 - $E_{in}^{(D_t)}(h) = \sum_{n=1}^N D_t(n) \text{error}(h(\vec{x}_n), y_n)$

When $D_t(n) = 1/N$. This reduces to standard definition of E_{in} .