

CSE 417T

# Introduction to Machine Learning

Lecture 14

Instructor: Chien-Ju (CJ) Ho

# Logistics

- Homework 1 Returned
  - Regrade requests till this Saturday
  - Please be concise and polite
- Homework 3: Due **Mar 5 (Sat)**
  - Keep track of your own late-day usages
- Exam 1: **Mar 10 (Thursday)**
  - Topics: LFD Chapters 1 to 5
  - Covid-permitting
    - Timed exam (75 min) during lecture time in the classroom
    - Closed-book exam with 2 letter-size cheat sheets allowed (4 pages in total)
      - No format limitations (it can be typed, written, or a combination)
  - Mar 8 (Tuesday) will be a review lecture

Recap

# Decision Tree Hypothesis



Credit Card Approval Example

- Pros
  - Easy to interpret (interpretability is getting attention and is important in some domains)
  - Can handle multi-type data (Numerical, categorical. ...)
  - Easy to implement (Bunch of if-else rules)
- Cons
  - Generally speaking, **bad generalization**
  - VC dimension is infinity
  - High variance (small change of data leads to very different hypothesis)
  - Easily overfit
- Why we care?
  - One of the classical model
  - Building block for other models (e.g., random forest)

# ID3: Using Information Gain as Selection Criteria

- Information gain of choosing feature  $A$  to split
  - $Gain(D, A) = H(D) - \sum_i \frac{|D_i|}{|D|} H(D_i)$  [The amount of decrease in entropy]
- ID3: Choose the split that maximize  $Gain(D, A)$

Notations:

$H(D)$ : Entropy of  $D$

$|D|$  is the number of points in  $D$

DecisionTreeLearn( $D$ )

Create a root node  $r$

If **termination conditions** are met

return a single node tree with **leaf prediction** based on

Else: Greedily find a feature  $A$  to split according to **split criteria**

For each possible value  $v_i$  of  $A$

Let  $D_i$  be the dataset containing data with value  $v_i$  for feature  $A$

Create a subtree DecisionTreeLearn( $D_i$ ) that being the child of root  $r$

- ID3 termination conditions
  - If all labels are the same
  - If all features are the same
  - If dataset is empty
- ID3 leaf predictions
  - Most common labels (majority voting)
- ID3 split criteria
  - Information gain

# Ensemble Learning


- Goal: Utilize a set of **weak learners** to obtain a **strong learner**.
- Format of ensemble learning
  - **Construct** many **diverse** weak learners
  - **Aggregate** the weak learners

## Bagging

- Construct diverse weak learners
  - (**Simultaneously**) bootstrap datasets
  - Train weak learners on them
- Aggregate the weak learners
  - **Uniform** aggregation

# Bagging - Bootstrapped Aggregating

- Bootstrap  $M$  datasets  $\{\tilde{D}^{\{m\}}\}$  (Sample with replacement from  $D$ )
- Learn a hypothesis from each of them



- Aggregate the learned hypothesis (assume we are doing classification)

$$G(\vec{x}) = \bar{g}(\vec{x}) = \text{sign}\left(\frac{1}{M} \sum_{m=1}^M g_m(\vec{x})\right)$$

# Out-Of-Bag (OOB) Error

	$\tilde{D}^{(1)}$	$\tilde{D}^{(2)}$	$\tilde{D}^{(3)}$	$\tilde{D}^{(4)}$	...
$(\vec{x}_1, y_1)$	Yes	Yes	No	No	...
$(\vec{x}_2, y_2)$	Yes	No	No	No	...
...	...	...	...	...	...
$(\vec{x}_N, y_N)$	No	Yes	Yes	Yes	...

Whether a point is in a bootstrapped dataset

- $G_n^-$ : the aggregation of hypothesis that  $\vec{x}_n$  is OOB of

- $G_1^- = \text{aggregate}(g_3, g_4, \dots)$
- $G_2^- = \text{aggregate}(g_2, g_3, g_4, \dots)$
- $G_N^- = \text{aggregate}(g_1, \dots)$

Aggregate:  
Majority voting for classification  
Average for regression

- OOB Error

- $E_{OOB}(G) = \frac{1}{N} \sum_{n=1}^N \text{error}(G_n^-(\vec{x}_n), y_n)$

Error:  
Binary error for classification  
Squared error for regression



# Out-Of-Bag (OOB) Error

$$E_{OOB}(G) = \frac{1}{N} \sum_{n=1}^N \text{error}(G_n^-(\vec{x}_n), y_n)$$

- Bagging provided an **intrinsic** mechanism for us to perform validation
  - We don't need to split the dataset into training and validation
- Practical issues (you might face this in HW4)
  - What if some  $\vec{x}_n$  appears in all bootstrapped datasets?
    - The probability of this happening is small when the number of bags  $M$  is large
  - Let  $S$  be the set of points that is out of bag for at least one bootstrapped dataset
    - $E_{OOB}(G) = \frac{1}{|S|} \sum_{(\vec{x}_n, y_n) \in S} \text{error}(G_n^-(\vec{x}_n), y_n)$

# Random Forest

- Construct many random trees
  - Bootstrapping datasets and learn a **max-depth tree** for each of them
  - Other randomizations (not required in HW4)
    - When choosing split features, choose from a random subset (instead of all features)
    - Randomly project features (similar to non-linear transformation) for each tree
- Aggregate the random trees
  - Classification: Majority vote  $\bar{g}(\vec{x}) = \text{sign} \left( \frac{1}{M} \sum_{m=1}^M g_m(\vec{x}) \right)$
  - Regression: Average  $\bar{g}(\vec{x}) = \frac{1}{M} \sum_{m=1}^M g_m(\vec{x})$

# Today's Lecture

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.  
Let me know if you spot errors.

# Boosting

# Ensemble Learning

- Goal: Utilize a set of **weak learners** to obtain a **strong learner**.
- Format of ensemble learning
  - **Construct** many **diverse** weak learners
  - **Aggregate** the weak learners

## Bagging:


- Construct diverse weak learners
  - (**Simultaneously**) bootstrapping datasets
  - Train weak learners on them
- Aggregate the weak learners
  - **Uniform** aggregation

## Boosting

- Construct diverse weak learners
  - **Adaptively** generating datasets
  - Train weak learners on them
- Aggregate the weak learners
  - **Weighted** aggregation

# Informal Intuitions about Boosting


- Example : Teach a class of kids to identify apples from data



- Alice: Apples are **circular**
- Teacher:  
Circular is a good feature, but using this feature might make some mistakes  
Let me **highlight** the mistakes.
  - Make correct images smaller
  - Make incorrect images larger

# Informal Intuitions about Boosting


- Example : Teach a class of kids to identify apples from data



- Alice: Apples are **circular**
- Bob: Apples are **red**

# Informal Intuitions about Boosting

- Example : Teach a class of kids to identify apples from data




- Alice: Apples are **circular**
- Bob: Apples are **red**
- Charlie: Apples could be **green**



# Informal Intuitions about Boosting

- Example : Teach a class of kids to identify apples from data



- Alice: Apples are **circular**
- Bob: Apples are **red**
- Charlie: Apples could be **green**
- David: Apples have **stems** at the top
- Class: Apples are **somewhat circular, somewhat red, possibly green, and may have stems at the top**

# Informal Intuitions about Boosting

- Example : Teach a class of kids to identify apples from data



Key steps of this process:

- Learn a **simple** hypothesis for each dataset
- Iteratively update the dataset to focus on what we got wrong (i.e., create **diversity**)
- **Aggregate** the learned simple hypothesis



- Alice: Apples are **circular**
- Bob: Apples are **red**
- Charlie: Apples could be **green**
- David: Apples have **stems** at the top
- Class: Apples are **somewhat circular, somewhat red, possibly green, and may have stems at the top**

# Outline of a Boosting Algorithm

- Initialize  $D_1$  (usually the same as the initial dataset  $D$ )
- For  $t = 1$  to  $T$ 
  - Learn  $g_t$  from  $D_t$
  - Reweight the distribution and obtain  $D_{t+1}$  based on  $g_t$  and  $D_t$
- Output **weighted**-aggregate( $g_1, \dots, g_T$ )
  - Classification:  $G(\vec{x}) = \bar{g}(\vec{x}) = \text{sign}\left(\frac{1}{T} \sum_{t=1}^T \alpha_t g_t(\vec{x})\right)$

## Questions

How to learn  $g_t$  from  $D_t$

How to reweight the distribution and obtain  $D_{t+1}$

How to perform weighted aggregation

# Discussion on Re-weighted $D_t$ (What does re-weighting mean?)

- Original Dataset  $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$
- Notation of  $D_t$ 
  - $D_t(n)$  is the weight/probability of data point  $(\vec{x}_n, y_n)$  in  $D_t$
  - $\sum_{n=1}^N D_t(n) = 1$
- What is  $E_{in}(h)$  on  $D_t$ ? (Expressed as  $E_{in}^{(D_t)}(h)$ )
  - Re-sample dataset (noisier)
    - Re-sample the dataset from  $D$  according to distribution  $D_t$
    - Calculate  $E_{in}$  on the re-sampled dataset as usual
  - Calculate weighted error
    - $E_{in}^{(D_t)}(h) = \sum_{n=1}^N D_t(n) \text{error}(h(\vec{x}_n), y_n)$

When  $D_t(n) = 1/N$ . This reduces to standard definition of  $E_{in}$ .

# AdaBoost – Adaptive Boosting

How to learn  $g_t$  from  $D_t$

How to reweight the distribution and obtain  $D_{t+1}$

How to perform weighted aggregation

[AdaBoost focuses on **classification** problem]

# Boosting Background

- A theoretical question asked by Kearns and Valiant
  - Whether a “weak” learning algorithm which performs just slightly better than random guessing in the PAC model can be “boosted” into an arbitrarily accurate “strong” learning algorithm
- AdaBoost
  - The first adaptive boosting algorithm that
    - has nice theoretical guarantees
    - successfully deployed in real-world applications

# What Does AdaBoost Do?

## Outline of a Boosting Algorithm

Initialize  $D_1$  (usually the same as the initial dataset  $D$ )

For  $t = 1$  to  $T$

    Learn  $g_t$  from  $D_t$

    Reweight the distribution and obtain  $D_{t+1}$  based on  $g_t$  and  $D_t$


Output **weighted-aggregate**( $g_1, \dots, g_T$ )

Classification:  $G(\vec{x}) = \bar{g}(\vec{x}) = \text{sign}\left(\frac{1}{T} \sum_{t=1}^T \alpha_t g_t(\vec{x})\right)$

- Will discuss the following for AdaBoost
  1. How to learn  $g_t$  from  $D_t$
  2. How to reweight the distribution and obtain  $D_{t+1}$
  3. How to perform weighted aggregation

# 1. Learn a Weak Learner $g_t$ from $D_t$

- AdaBoost uses *simple* weak learners
  - low variance, high bias
  - Decision stump (one-level decision tree) is one good option



- How to learn  $g_t$  from  $D_t$ 
  - Find the decision stump that
    - Minimizes  $E_{in}^{(D_t)}$
    - Maximize (weighted) information gain (you can call decision tree library directly)



## 2. How to Reweight $D_{t+1}$

- We want to make  $g_{t+1}$  (learned from  $D_{t+1}$ ) to be **diverse** from  $g_t$ 
  - **Increase** the weights of points that  $g_t$  makes **wrong** predictions
  - **Decrease** the weights of points that  $g_t$  makes **correct** predictions
- Define a parameter  $\gamma > 1$ 
  - If  $g_t$  makes **wrong** predictions on  $\vec{x}_n$ 
    - $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) \cdot \gamma$  (increase the weight)
  - If  $g_t$  makes **correct** predictions on  $\vec{x}_n$ 
    - $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) / \gamma$  (decrease the weight)
- Goal:
  - Choose  $\gamma$  such that  $E_{in}^{(D_{t+1})}(g_t) = 0.5$
  - Since  $g_{t+1}$  minimizes  $E_{in}^{(D_{t+1})} \Rightarrow g_t$  and  $g_{t+1}$  are **diverse**

$Z_t$ : normalization constant  
We need to ensure  
 $\sum_{n=1}^N D_{t+1}(n) = 1$

Choose  $\gamma$  such that  $E_{in}^{(D_{t+1})}(g_t) = 0.5$

Math derivations in the next few slides

- Define  $\epsilon_t = E_{in}^{(D_t)}(g_t) = \sum_{n=1}^N D_t(n) \mathbb{I}[g_t(\vec{x}_n) \neq y_n]$ 
  - Weighted in-sample error of  $g_t$  on  $D_t$
  - $\epsilon_t < 0.5$  (requirement of weak learners)

We consider the case weak learners are better than random guessing:  
 $\epsilon_t < 0.5$

- Goal: Want to make  $E_{in}^{(D_{t+1})}(g_t) = 0.5$

$$E_{in}^{(D_t)}(h) = \sum_{n=1}^N D_t(n) \text{error}(h(\vec{x}_n), y_n)$$

- If  $g_t$  makes **wrong** predictions on  $\vec{x}_n$ 
  - $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) \cdot \gamma$  (increase the weight)
- If  $g_t$  makes **correct** predictions on  $\vec{x}_n$ 
  - $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) / \gamma$  (decrease the weight)

- Define  $\epsilon_t = E_{in}^{(D_t)}(g_t) = \sum_{n=1}^N D_t(n) \mathbb{I}[g_t(\vec{x}_n) \neq y_n]$ 
  - Weighted in-sample error of  $g_t$  on  $D_t$
  - $\epsilon_t < 0.5$  (requirement of weak learners)

We consider the case weak learners are better than random guessing:  
 $\epsilon_t < 0.5$

- Goal: Want to make  $E_{in}^{(D_{t+1})}(g_t) = 0.5$

$$\begin{aligned} E_{in}^{(D_{t+1})}(g_t) &= \sum_{n=1}^N D_{t+1}(n) \mathbb{I}[g_t(\vec{x}_n) \neq y_n] \\ &= \sum_{n=1}^N \frac{1}{Z_t} D_t(n) \gamma \mathbb{I}[g_t(\vec{x}_n) \neq y_n] \\ &= \frac{\gamma}{Z_t} \sum_{n=1}^N D_t(n) \mathbb{I}[g_t(\vec{x}_n) \neq y_n] = \frac{\gamma}{Z_t} \epsilon_t \end{aligned}$$

$$E_{in}^{(D_t)}(h) = \sum_{n=1}^N D_t(n) \text{error}(h(\vec{x}_n), y_n)$$

- If  $g_t$  makes **wrong** predictions on  $\vec{x}_n$ 
  - $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) \cdot \gamma$  (increase the weight)
- If  $g_t$  makes **correct** predictions on  $\vec{x}_n$ 
  - $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) / \gamma$  (decrease the weight)

- Remember  $Z_t$  is the normalization constant

$$\begin{aligned} Z_t &= \sum_{n=1}^N D_t(n) \gamma \mathbb{I}[g_t(\vec{x}_n) \neq y_n] + \sum_{n=1}^N D_t(n) \frac{1}{\gamma} \mathbb{I}[g_t(\vec{x}_n) = y_n] \\ &= \gamma \epsilon_t + \frac{1}{\gamma} (1 - \epsilon_t) \end{aligned}$$

- Want to make  $E_{in}^{(D_{t+1})}(g_t) = 0.5$ 
  - $E_{in}^{(D_{t+1})}(g_t) = \frac{\gamma}{Z_t} \epsilon_t$
  - $Z_t = \gamma \epsilon_t + \frac{1}{\gamma} (1 - \epsilon_t)$

- Want to make  $E_{in}^{(D_{t+1})}(g_t) = 0.5$

- $E_{in}^{(D_{t+1})}(g_t) = \frac{\gamma}{Z_t} \epsilon_t$

- $Z_t = \gamma \epsilon_t + \frac{1}{\gamma} (1 - \epsilon_t)$

- $\frac{\gamma \epsilon_t}{\gamma \epsilon_t + (1 - \epsilon_t)/\gamma} = 0.5 \Rightarrow \frac{1 - \epsilon_t}{\gamma} = \gamma \epsilon_t \Rightarrow \gamma = \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}}$

Both  $g_t(\vec{x}_n)$  and  $y_n$  are either +1 or -1  
 If  $g_t(\vec{x}_n) \neq y_n$ ,  $g_t(\vec{x}_n)y_n = -1$   
 If  $g_t(\vec{x}_n) = y_n$ ,  $g_t(\vec{x}_n)y_n = 1$

- The rule for reweighting

- If  $g_t(\vec{x}_n) \neq y_n$ , then  $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) \left( \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \right)$

- If  $g_t(\vec{x}_n) = y_n$ , then  $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) \left( \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \right)^{-1}$

- Want to make  $E_{in}^{(D_{t+1})}(g_t) = 0.5$

- $E_{in}^{(D_{t+1})}(g_t) = \frac{\gamma}{Z_t} \epsilon_t$

- $Z_t = \gamma \epsilon_t + \frac{1}{\gamma} (1 - \epsilon_t)$

- $\frac{\gamma \epsilon_t}{\gamma \epsilon_t + (1 - \epsilon_t)/\gamma} = 0.5 \Rightarrow \frac{1 - \epsilon_t}{\gamma} = \gamma \epsilon_t \Rightarrow \gamma = \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}}$

- The rule for reweighting

- If  $g_t(\vec{x}_n) \neq y_n$ , then  $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) \left( \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \right) = \frac{1}{Z_t} D_t(n) \left( \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \right)^{-g_t(\vec{x}_n) y_n}$

- If  $g_t(\vec{x}_n) = y_n$ , then  $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) \left( \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \right)^{-1} = \frac{1}{Z_t} D_t(n) \left( \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \right)^{-g_t(\vec{x}_n) y_n}$

- Reweight rule:  $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) \left( \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \right)^{-g_t(\vec{x}_n) y_n}$

## 2. How to Reweight $D_{t+1}$ : Summary

- Reweight rule:

- $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) \left( \sqrt{\frac{1-\epsilon_t}{\epsilon_t}} \right)^{-g_t(\vec{x}_n)y_n}$

- A bit more manipulations (the reason will be clear later)

- Define  $\alpha_t = \frac{1}{2} \ln \left( \frac{1-\epsilon_t}{\epsilon_t} \right)$

- $e^{\alpha_t} = \sqrt{\frac{1-\epsilon_t}{\epsilon_t}}$

- Final reweight rule:  $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) e^{-\alpha_t g_t(\vec{x}_n)y_n}$



# 3. How to Aggregate Weak Learners

- Intuition:
  - We want to put more weights on better weak learners
  - $\epsilon_t = E_{in}^{(D_t)}(g_t)$  is a proxy on how well  $g_t$  performs (smaller  $\epsilon_t \Rightarrow$  better  $g_t$ )
  - Recall that  $\alpha_t = \frac{1}{2} \ln \left( \frac{1-\epsilon_t}{\epsilon_t} \right)$ 
    - Better  $g_t$ , smaller  $\epsilon_t$ , higher  $\alpha_t$
    - When  $\epsilon_t = 0.5$ ,  $\alpha_t = 0$  (random guessing leads to 0 weights)
    - When  $\epsilon_t = 0$ ,  $\alpha_t = \infty$  (if a feature perfectly classifies the data, use it as our final hypothesis)
- Aggregation rule
  - $G(\vec{x}) = \text{sign}(\sum_{t=1}^T \alpha_t g_t(\vec{x}))$

# AdaBoost Algorithm

- Given  $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$
- Initialize  $D_1(n) = 1/N$  for all  $n = 1, \dots, N$
- For  $t = 1, \dots, T$ 
  - Learn  $g_t$  from  $D_t$  (using decision stumps)
  - Calculate  $\epsilon_t = E_{in}^{(D_t)}(g_t)$
  - Set  $\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$
  - Update  $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) e^{-\alpha_t y_n g_t(\vec{x}_n)}$
- Output  $G(\vec{x}) = \text{sign}(\sum_{t=1}^T \alpha_t g_t(\vec{x}))$

# Theoretical Properties of AdaBoost


- See [Freund & Schapire's Tutorial](#) for more discussion
- The training error of AdaBoost converges fast
  - Let  $\gamma_t = \frac{1}{2} - \epsilon_t$  (how good each weak learner is better than random guessing)
  - $E_{in} \leq e^{-2 \sum_{t=1}^T \gamma_t^2}$
- Generalization error
  - VC analysis gives us  $E_{out} \leq E_{in} + \tilde{O}\left(\sqrt{\frac{T d_{vc}}{m}}\right)$
  - It seems as  $T$  goes large, overfitting could happen
  - Empirically, AdaBoost is relatively robust to overfitting
  - There are some more delicate analysis using the idea of **margins** to explain why

$d_{vc}$  is the VC dimension  
of the weak learner


# AdaBoost in Action

# AdaBoost in Action

- A toy example (by Yoav Freund Rob Schapire)
- Weak learner: decision stump (one-level decision tree)




# Round 1




$$\epsilon_1 = 0.30$$
$$\alpha_1 = 0.42$$

$D_2$




# Round 2




$$\epsilon_2 = 0.21$$
$$\alpha_2 = 0.65$$

$D_3$




# Round 3



$$\epsilon_3 = 0.14$$


$$\alpha_3 = 0.92$$






$$\varepsilon_1=0.30$$

$$\alpha_1=0.42$$



$$\varepsilon_2=0.21$$

$$\alpha_2=0.65$$



$$\varepsilon_3=0.14$$

$$\alpha_3=0.92$$


$H_{\text{final}}$

$$= \text{sign} \left( \begin{array}{c} 0.42 \\ \text{Diagram 1} \end{array} + 0.65 \begin{array}{c} \text{Diagram 2} \\ 0.65 \end{array} + 0.92 \begin{array}{c} \text{Diagram 3} \\ 0.92 \end{array} \right) = \text{Diagram 4}$$




# Practical Success of AdaBoost

# Viola-Jones Face Detection (2001)


- First real-time object detection framework
- Paul Viola and Michael Jones. Rapid object detection using a boosted cascade of simple features. CVPR 2001.



# Weak Learners (Haar wavelet features)



# Weak Learners (Haar wavelet features)



- Each hypothesis is very weak.
- There are many possible features.
  - For a 24x24 detection region, more than 160,000 features
- AdaBoost!
  - Training is slow
  - Testing is fast
    - (inherent feature selection)

# Brief Discussion on Gradient Boosting

Gradient boosting is **safe to skip** for Exam 2

# Look at the AdaBoost Algorithm Again

Given  $D = \{(\vec{x}_1, y_1), \dots, (\vec{x}_N, y_N)\}$   
Initialize  $D_1(n) = 1/N$  for all  $n = 1, \dots, N$   
For  $t = 1, \dots, T$   
    Learn  $g_t$  from  $D_t$  (using decision stumps)  
    Calculate  $\epsilon_t = E_{in}^{(D_t)}(g_t)$   
    Set  $\alpha_t = \frac{1}{2} \ln \left( \frac{1-\epsilon_t}{\epsilon_t} \right)$   
    Update  $D_{t+1}(n) = \frac{1}{Z_t} D_t(n) e^{-\alpha_t y_n g_t(\vec{x}_n)}$   
Output  $G(\vec{x}) = \text{sign}(\sum_{t=1}^T \alpha_t g_t(\vec{x}))$



Initialize  $G(\vec{x}) = 0$   
For  $t = 1, \dots, T$   
     $G(\vec{x}) \leftarrow G(\vec{x}) + \alpha_t g_t(\vec{x})$   
Output  $\text{sign}(G(\vec{x}))$

- The format is similar to **gradient descent**!
  - If we consider the space of the weak learners (i.e.,  $g_t(\vec{x})$ ) as the space of “weights”
  - This observation leads to a general class of boosting algorithms: **gradient boosting**
  - XGBoost is one implementation of gradient boosting that is popular in practice
  - See CASI 17.4 and the reference in CASI P.350 for more discussion

[Safe to Skip]

# Gradient Boosting


Initialize  $G(\vec{x}) = 0$

For  $t = 1, \dots, T$

$$G(\vec{x}) \leftarrow G(\vec{x}) + \alpha_t g_T(\vec{x})$$

Output  $\text{sign}(G(\vec{x}))$

- AdaBoost is a special case of Gradient Boosting
  - minimizing the exponential **loss** ( $e_{\text{exp}}(h(\vec{x}), y) = e^{-yh(\vec{x})}$ )
  - using decision stump as the **weak learners**



- $e_{\text{exp}}$  is a **surrogate loss function** of the binary classification error we care about
  - Minimizing an alternative error (loss function) is a common trick in ML, especially when the target loss function is hard to optimize.
  - There are some theoretical discussions on when doing this makes sense (“calibration”: whether minimizing the surrogate is consistent with minimizing the original loss).

[Safe to Skip]