

CSE 417T

# Introduction to Machine Learning

Lecture 16

Instructor: Chien-Ju (CJ) Ho

# Logistics

- Homework 4 is due November 14 (Monday)
- Keep track of your own late days
  - Gradescope doesn't allow separate deadlines
  - Your submissions **won't be graded** if you exceed the late-day limit

Recap


# Nearest Neighbor

- Predict  $\vec{x}$  according to its nearest neighbor
  - Given  $D = \{(\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_N, y_N)\}; y_n \in \{+1, -1\}$
  - Let  $\vec{x}_{[1]}$  be  $\vec{x}$ 's nearest neighbor, i.e., the closest point to  $\vec{x}$  in  $D$
  - Let  $y_{[i]}(\vec{x})$  or  $y_{[i]}$  be the label of  $\vec{x}_{[i]}$
- Nearest neighbor hypothesis

$$g(\vec{x}) = y_{[1]}(\vec{x})$$

# Nearest Neighbor

$g(\vec{x})$  looks like a Voronoi diagram



- Properties of Nearest Neighbor (NN)
  - No training is needed
  - Good interpretability
  - In-sample error  $E_{in} = 0$
  - VC dimension is  $\infty$
- This seems to imply bad learning models from what we talk about so far? Why we care?
- What we really care about is  $E_{out}$ 
  - VC analysis:  $E_{out} \leq E_{in} + \text{Generalization error}$ 
    - We can infer  $E_{out}$  through  $E_{in}$  and model complexity
  - NN has nice guarantees outside of VC analysis

# Nearest Neighbor is 2-Optimal

- Given mild conditions, for nearest neighbor, when  $N \rightarrow \infty$ , with high probability,

$$E_{out} \leq 2E_{out}^*$$

- That is, we can not infer  $E_{out}$  from  $E_{in}$ , but we know it cannot be much worse than the **best anyone can do**.

## Informal intuitions:

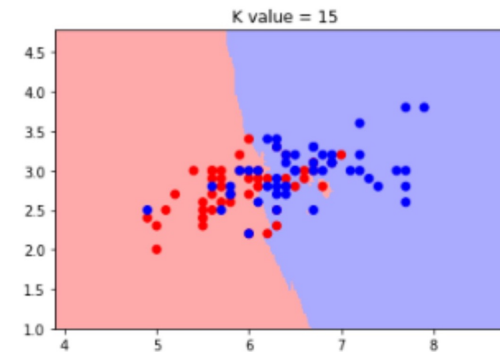
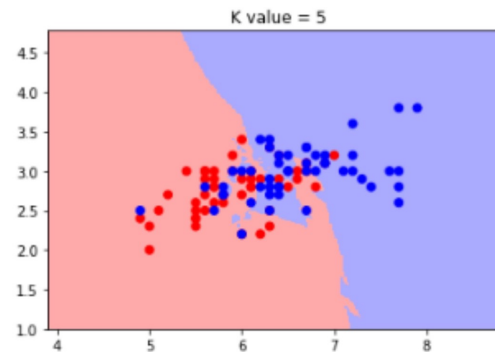
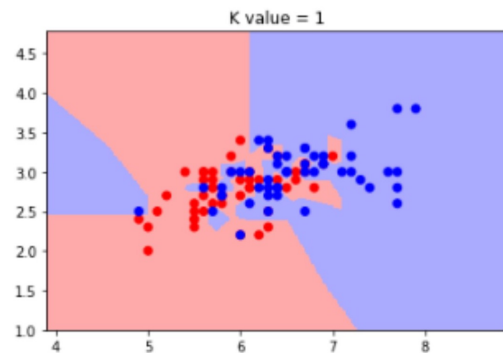
Assumption for nearest neighbor:

- Target function is continuous
- That is, nearby points have similar label distributions

As  $N$  goes large, there are “close enough” points for prediction


# $k$ -Nearest Neighbor (K-NN)

- Nearest neighbor
  - Influenced heavily by noisy data
- $k$ -nearest neighbor (K-NN)
  - $g(\vec{x}) = \text{sign}\left(\sum_{i=1}^k y_{[i]}(\vec{x})\right)$
  - ( $k$  is often odd for binary classification)




# Impacts of $k$ in $k$ -Nearest Neighbor


- $k = 1$ : the nearest neighbor hypothesis
  - many, complicated decision boundaries; may overfit
- $k = N$ ,  $g$  predicts the most common label in the training dataset
  - no decision boundaries; may underfit
- $k$  controls the complexity of the hypothesis set
  - $k$  affects how well the learned hypothesis will generalize



1-NN rule



21-NN rule



127-NN rule



# How to Choose $k$

- Making the choice of  $k$  a function of  $N$ , denoted by  $k(N)$
- Theorem:
  - For  $N \rightarrow \infty$ , if  $k(N) \rightarrow \infty$  and  $\frac{k(N)}{N} \rightarrow 0$
  - Then  $E_{in}(g) \rightarrow E_{out}(g)$  and  $E_{out}(g) \rightarrow E_{out}(g^*)$
- Example:  $k(N) = \sqrt{N}$  satisfies the condition

Informal intuitions:

$\frac{k(N)}{N} \rightarrow 0$ : all the nearest neighbors are next to  $\vec{x}$

$k(N) \rightarrow \infty$ : select infinitely many neighbors

=> We can almost reconstruct the target function

# How to Choose $k$

- Making the choice of  $k$  a function of  $N$ , denoted by  $k(N)$
- Theorem:
  - For  $N \rightarrow \infty$ , if  $k(N) \rightarrow \infty$  and  $\frac{k(N)}{N} \rightarrow 0$
  - Then  $E_{in}(g) \rightarrow E_{out}(g)$  and  $E_{out}(g) \rightarrow E_{out}(g^*)$
- Example:  $k(N) = \sqrt{N}$  satisfies the condition
- Practical rules of thumb:
  - Small  $k$  (e.g.,  $k = 3$ ) is often a good enough choice
  - Using validation to choose  $k$

# Summary of $k$ -NN So Far

- Pros
  - Simple algorithm
  - Good interpretations
  - Nice theoretical guarantee
  - Easy to adapt to regression (average of nearest neighbors) and multi-class classification (majority voting)
- Cons
  - Curse of dimensionality
  - Computational issue
    - each prediction requires  $O(N)$  computation

# Today's Lecture

The notes are not intended to be comprehensive. They should be accompanied by lectures and/or textbook.  
Let me know if you spot errors.

# Curse of Dimensionality

- Generally, higher dimensionality implies harder learning (think VC)
- Things are worse with similarity-based methods
  - Rely on assumptions that nearby points have similar labels
  - As the dimension grows, most of the points will not be nearby to each other...

# Illustration of Curse of Dimensionality

- Think about Euclidean distance:  $d(\vec{x}, \vec{x}') = \|\vec{x} - \vec{x}'\|$
- Illustration
  - Consider the space  $[0,1]^d$  (a hypercube with length of each side = 1)
  - What's the side length  $\ell$  of a hypercube that takes up 1% of the space?
    - $d = 1$ :  $\ell = 0.01$
    - $d = 2$ :  $\ell = 0.1$
    - ...
    - $d = 100$ :  $\ell^d = 0.01 \Rightarrow \ell \approx 0.95$
    - $d = 1000$ :  $\ell^d = 0.01 \Rightarrow \ell \approx 0.9954$

# Illustration of Curse of Dimensionality

- Consider the distance to the origin when  $d = 100$ 
  - Consider the case that the value of each dimension is uniformly drawn
  - Only 1% of the points will be in the hypercube  $[0,0.95]^{100}$
  - Most of the points will be far away from the origin
  - Most of the points will be far away from each other
- No simple solutions....
  - Getting more data
  - Reduce dimensions (see LFD 9.2)

# Computational Issues for k-NN




# Computational Issues

- $k$ -Nearest Neighbor is computationally demanding
  - Need to store all data points: space complexity  $O(Nd)$
  - For each prediction for  $\vec{x}$ 
    - Calculate the distance to every point in  $D$
    - Find the  $k$  closest points
    - Time complexity  $O(Nd + N \log k)$
- Two general approaches:
  - Reduce the number of data points
  - Store the data in some data structure to speed up searching
  - See LFD 6.2.3 for more discussion

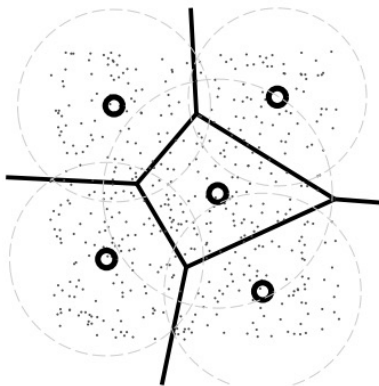
# Computational Issues: Potential Solutions

- Reduce the number of data points



- Intuition: remove points that will not impact the decision boundary.
- Generally a hard problem. But there are some heuristic approaches.


- Store the data in some data structure to speed up searching



- Intuition: Cluster data points
- For a new data point, first find a nearest cluster. Then find the nearest points within that cluster

# Reduce the Amount of Data


- Consider the following example



- Goal: Drop data that doesn't impact the decision boundary
  - Find  $S \subseteq D$ , such that  $g_S(\vec{x}) = g_D(\vec{x}) \forall \vec{x}$
  - Generally a hard goal to achieve


# A Probably Easier Goal

- Instead of making the **decision boundary** consistent



Find  $S \in D$ , such that  $g_S(\vec{x}) = g_D(\vec{x}) \forall \vec{x}$


- Aim for making the prediction on the **training data** consistent



Find  $S \in D$ , such that  $g_S(\vec{x}) = g_D(\vec{x}) \forall \vec{x} \in D$

# Condensed Nearest Neighbor (CNN)


- Goal: Find  $S \in D$ , such that  $g_S(\vec{x}) = g_D(\vec{x}) \forall \vec{x} \in D$



- CNN: An iterative algorithm
  - Random initialize a subset of data points  $S$
  - While (there exists point  $\vec{x}_* \in D$  such that  $g_S(\vec{x}_*) \neq g_D(\vec{x}_*)$ )
    - Let  $y_*$  be  $g_D(\vec{x}_*)$
    - Find the nearest point  $\vec{x}' \in D \setminus S$  with the label  $y_*$
    - Insert  $\vec{x}'$  to  $S$
- No theoretical guarantees
- Reasonable empirical performance

# Speed Up the Search for Nearest Neighbor

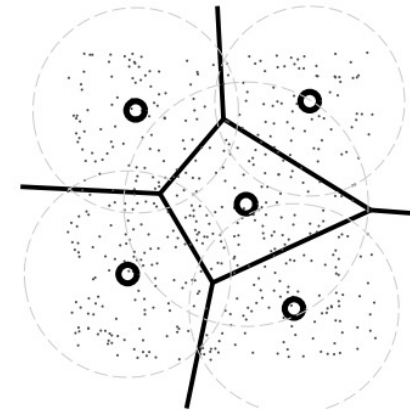
- For each prediction for  $\vec{x}$ , we need to find the nearest point in  $D$
- Can we speed up this process?



- $S_1$  and  $S_2$  are two clusters
  - Centers  $\vec{\mu}_1, \vec{\mu}_2$  and radii  $r_1, r_2$
- Let  $S_1$  be the nearest cluster for  $\vec{x}$
- Let  $\vec{x}'_{[1]}$  be the nearest point in  $S_1$
- Distance from  $\vec{x}$  to any point in  $S_2$  is at least  $\|\vec{x} - \vec{\mu}_2\| - r_2$
- If  $\|\vec{x} - \vec{x}'_{[1]}\| \leq \|\vec{x} - \vec{\mu}_2\| - r_2$ 
  - we don't need to search points in  $S_2$

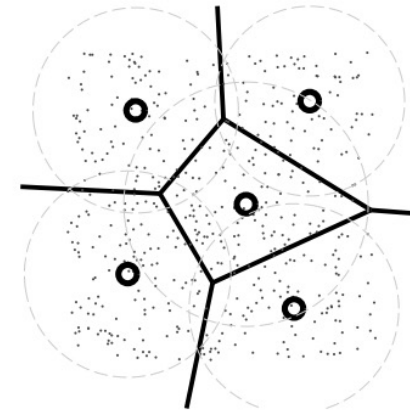
# Lloyd's Algorithm: Construct Clusters

- Goal: Cluster data into  $K$  clusters
  - Again, hard in general, but we can use greedy-based approaches



# Lloyd's Algorithm: Construct Clusters

- Goal: Cluster data into  $K$  clusters
  - Again, hard in general, but we can use greedy-based approaches
- Lloyd's Algorithm
  1. Randomly pick  $K$  points as centers
  2. Create the Voronoi regions as clusters
  3. Update the centers (calculating the mean)
  4. Update the region
  5. Repeat 3 and 4
- This is the first (and probably the only) unsupervised learning algorithm we talk about in this course.






# Final Remark:


## Nearest Neighbor is Nonparametric

Nearest Neighbor



no parameters  
expressive/flexible  
 $g(\mathbf{x})$  needs data  
generic, can model anything

Linear models



$(d + 1)$  parameters  
rigid, always linear  
 $g(\mathbf{x})$  needs only weights  
specialized

# Radial Basis Functions (RBF)

# Instance-Based Learning

- Make predictions based on data instances
- $k$ -nearest neighbor (K-NN)
  - $g(\vec{x}) = \text{sign}(\sum_{i=1}^k y_{[i]}(\vec{x}))$
  - Predict according to the nearest neighbors
- Radial Basis Functions (RBF)
  - Focus of today
- Kernel SVM (Topic of Next Week)
  - $g(\vec{x}) = \text{sign}(\sum_{\alpha_n^* > 0} \alpha_n^* y_n K(\vec{x}_n, \vec{x}) + b^*))$
  - Predict according to support vectors

# Radial Basis Functions

- Think about  $k$ -nearest neighbor (K-NN) again
  - $g(\vec{x}) = \text{sign}(\sum_{i=1}^k y_{[i]}(\vec{x}))$ 
    - Make predictions based on  $k$  nearest data points
    - Each of the  $k$  data points has the same weight
- Natural questions:
  - Can we use more (or even all) data?
  - **Weight** them based on **how close** data points are to  $\vec{x}$

# Radial Basis Functions

- Given dataset  $D = \{\vec{x}_1, \dots, \vec{x}_N\}$
- Task: Make a prediction on  $\vec{x}$

- Radial Basis Function:

- $g(\vec{x}) = \frac{1}{Z(\vec{x})} \sum_{n=1}^N \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right) y_n$

- $\phi(s)$ : a monotonically decreasing function

- This is for regression. We can take a sign and make it a binary classification.
- $Z(\vec{x}) = \sum_{m=1}^N \phi\left(\frac{\|\vec{x} - \vec{x}_m\|}{r}\right)$  is for normalization

- It's called **radial** basis function since it takes the **distance** to the points as the basis function

# Radial Basis Functions

- Radial Basis Function:

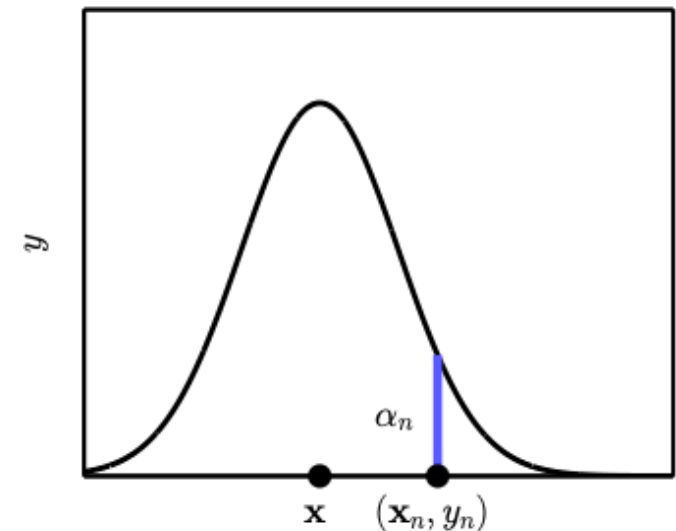
- $g(\vec{x}) = \sum_{n=1}^N \frac{1}{Z(\vec{x})} \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right) y_n$

- Example of  $\phi$

- Gaussian RBF:  $\phi(s) = e^{-s}$

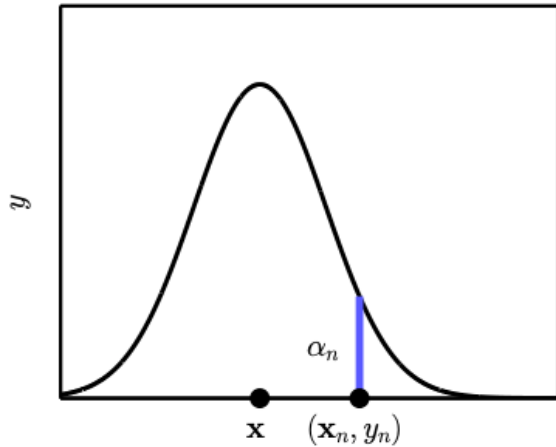
- Intuitions

- The impact of  $\vec{x}_n$  to  $\vec{x}$  is higher if it's closer to  $\vec{x}$
  - The role of  $r$  is similar to  $k$  in  $k$ -NN
    - $r \rightarrow 0$  : 1-NN
    - $r \rightarrow \infty$  :  $N$ -NN (i.e.,  $k = N$ )



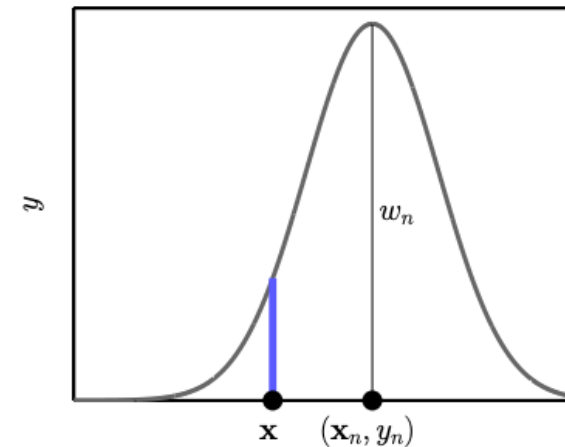
# Changing the Viewpoints

- $\phi$  centered around  $\vec{x}$



$$g(\vec{x}) = \sum_{n=1}^N \frac{y_n}{Z(\vec{x})} \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$$

- $\phi$  centered around  $\vec{x}_n$



$$g(\vec{x}) = \sum_{n=1}^N w_n(\vec{x}) \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$$

# From Nonparametric to Parametric RBF

- Nonparametric RBF

- $g(\vec{x}) = \sum_{n=1}^N \frac{y_n}{Z(\vec{x})} \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$
- $g(\vec{x}) = \sum_{n=1}^N w_n(\vec{x}) \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$
- The hypothesis is defined by dataset

- Parametric RBF hypothesis set

- $h(\vec{x}) = \sum_{n=1}^N w_n \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$
- Learn  $w_n$  from data



# Parametric RBF => Linear Models

- Parametric RBF is linear model with nonlinear transformation

- $h(\vec{x}) = \sum_{n=1}^N w_n \phi\left(\frac{\|\vec{x} - \vec{x}_n\|}{r}\right)$

- The projection  $\Phi(\vec{x}) = \begin{bmatrix} \phi\left(\frac{\|\vec{x} - \vec{x}_1\|}{r}\right) \\ \phi\left(\frac{\|\vec{x} - \vec{x}_2\|}{r}\right) \\ \vdots \\ \phi\left(\frac{\|\vec{x} - \vec{x}_N\|}{r}\right) \end{bmatrix}$

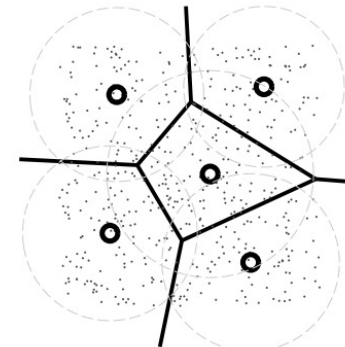
- We can apply what we learned in linear models to learn  $w_n$
- However, this seems to be **overfitting** ( $N$  parameters for  $N$  points)

# From $N$ points to $K$ points

- Use only  $K$  points  $(\vec{\mu}_1, \dots, \vec{\mu}_K)$

- $h(\vec{x}) = \sum_{k=1}^K w_k \phi\left(\frac{\|\vec{x} - \vec{\mu}_k\|}{r}\right)$

- Which  $K$  points?
  - We can find  $K$  representative points
  - Use clustering algorithms, e.g., Lloyd algorithm as introduced earlier
    1. Randomly pick  $K$  points as centers
    2. Create the Voronoi regions as clusters
    3. Update the centers (calculating the mean)
    4. Update the region
    5. Repeat 3 and 4




# More Discussion on RBF

- $h(\vec{x}) = \sum_{k=1}^K w_k \phi\left(\frac{\|\vec{x} - \vec{\mu}_k\|}{r}\right)$
- Connection to linear models
  - Parametric RBF is essentially linear model with nonlinear transformation
- Connection to nearest neighbor
  - Radial Basis Function is defined by “similarity”
  - A prediction for a point is based on the “similarity” of the points to be predicted and other points

# Support Vector Machines (SVM)


# What Do We Know about Linear Classification?

- What we discussed so far:
  - PLA: Find a linear separator that separates the data within finite steps, if data is linear separable.
  - Pocket algorithm: empirically keep the best separator during PLA.
  - Surrogate loss: Using logistic regression for linear classification.
- Challenges
  - Binary classification error is hard to optimize
  - We cannot use “gradient descent” type of algorithm minimize  $E_{in}$ .
- Support vector machines (SVM) tries to look at things a bit differently.



# Linear Classification


- Which separator would you choose?



Probably the right one.  
Why?

# Linear Classification


- Which separator would you choose?



More robust to noise (e.g., measurement error of  $\vec{x}$ )

# Linear Classification

- Which separator would you choose?



Margin: shortest distance from the separator to the points in  $D$   
(Informal argument)

Higher margin  $\Rightarrow$  more “constrained” hypothesis  $\Rightarrow$  lower VC dimension



# Support Vector Machine

- Goal:
  - Find the **max-margin** linear separator that separates the data
  - Recall the goal of PLA: Find the linear separator that separates the data

- Notations:

Notations we used so far:

- $\vec{x} = (x_0, x_1, \dots, x_d)$
- $\vec{w} = (w_0, w_1, \dots, w_d)$
- Linear separator  
$$h(\vec{x}) = \text{sign}(\vec{w}^T \vec{x})$$

Notations we will use in SVM

- $\vec{x} = (x_1, \dots, x_d)$
- $\vec{w} = (w_1, \dots, w_d)$
- Linear separator  
$$h(\vec{x}) = \text{sign}(\vec{w}^T \vec{x} + b)$$

Separating the bias/intercept  $b$  is important for us to characterize the margin.

We will use  $(\vec{w}, b)$  to characterize the hypothesis