

# Logistics

- Submit peer reviews **by 6pm.**
- Milestone 2: Due Nov 4 (tomorrow)
  - Describe your progress in no more than 2 pages

# Lecture 18

## Fairness in AI

Instructor: Chien-Ju (CJ) Ho

## ROBO RECRUITING

# Can an Algorithm Hire Better Than a Human?



**Claire Cain Miller** @clairecm JUNE 25, 2015



Hiring and recruiting might seem like some of the least likely jobs to be automated. The whole process seems to need human skills that computers lack, like making conversation and reading social cues.

But people have biases and predilections. They make hiring decisions, often unconsciously, based on similarities that have nothing to do with the job requirements — like whether an applicant has a friend in common, went to the same school or likes the same sports.

That is one reason researchers say traditional job searches are broken. The question is how to make them better.

A new wave of start-ups — including [Gild](#), [Entelo](#), [Textio](#), [Doxa](#) and [GapJumpers](#) — is trying various ways to automate hiring. They say that software can do the job more effectively and efficiently than people can. Many people are beginning to buy into the idea. Established headhunting firms like Korn Ferry are [incorporating algorithms](#) into their work, too.

If they succeed, they say, hiring could become faster and less expensive, and their data could lead recruiters to more highly skilled people who are better matches for their companies. Another potential result: a more diverse workplace. The software relies on data to surface candidates from a wide

## RECENT COMMENTS

**Mayurakshi Ghosh** January 7, 2016

Hi Claire, excellent article and really insightful facts on algorithm recruitment. I completely agree how you mentioned the role played by...

**Deborah Bishop** July 2, 2015

This is a very interesting article. Perhaps distinguishing between different aspects in the process of bringing talent into your...

**Yeti** July 2, 2015

Including talents that would be rejected by the subjective biased boss or colleagues does not fuarantee their integration. They will be...

[SEE ALL COMMENTS](#) [WRITE A COMMENT](#)

## HIDDEN BIAS

# When Algorithms Discriminate



Claire Cain Miller @clairecm JULY 9, 2015



The online world is shaped by forces beyond our control, determining the stories we read on Facebook, the people we meet on OkCupid and the search results we see on Google. Big data is used to make decisions about health care, employment, housing, education and policing.

But can computer programs be discriminatory?

There is a widespread belief that software and algorithms that rely on data [are objective](#). But software is not free of human influence. Algorithms are written and maintained by people, and machine learning algorithms adjust what they do based on people's behavior. As a result, say researchers in computer science, ethics and law, algorithms can [reinforce human prejudices](#).

Google's online advertising system, for instance, showed an ad for high-income jobs to men much more often than it showed the ad to women, [a new study](#) by Carnegie Mellon University researchers found.

[Research from Harvard University](#) found that ads for arrest records were significantly more likely to show up on searches for distinctively black names or a historically black fraternity. The [Federal Trade Commission said](#) advertisers are able to target people who live in low-income neighborhoods

## RECENT COMMENTS

tom July 10, 2015

Discrimination against women persists in other ways. Take the obituary column of the NYT - on a good week, you will find obits for perhaps...

SierramanCA July 10, 2015


"There is a widespread belief that software and algorithms that rely on data are objective." says Ms. Miller. Well, Ms. Miller, two things:1...

Dalglish July 10, 2015

Algorithms are written by people. People are biased, not objective. Daniel Kahneman et al. have proven this.

[SEE ALL COMMENTS](#)

# BRIEF HISTORY OF FAIRNESS IN ML



Isn't the point of ML to discriminate?

Want to avoid “unjustified” discrimination

# Example: Loan Applications

- By law, the banks can't discriminate people according to their race.
- First natural approach (fairness through blindness)
  - remove the race attribute from the data
- Guess what happened?
  - Redlining





# What should we do?

- From computer scientists / engineers' point of view....
- Give me an operational definition of fairness, I'll implement a system that satisfy it!
- How should we define fairness?
- Is it even possible to define a universal fairness notion?



# This Lecture: Group Fairness

- Consider social groups defined by **sensitive attributes**, such as gender, race, sexual orientation, etc.
- If we believe a decision shouldn't be depending on these sensitive attributes, then we want the **outcome** for different groups to be similar.

Won't Recidivate	<b>TN1</b>	<b>FP1</b>
Will Recidivate	<b>FN1</b>	<b>TP1</b>
	Labeled Low-Risk	Labeled High-Risk

Won't Recidivate	<b>TN2</b>	<b>FP2</b>
Will Recidivate	<b>FN2</b>	<b>TP2</b>
	Labeled Low-Risk	Labeled High-Risk

- Defendant: the probability that I'm incorrectly classified high-risk is independent of my race.
  - Equal False Positive Rate:  $\frac{FP1}{FP1 + TN1} = \frac{FP2}{FP2 + TN2}$
- Defendant: the probability that I'm incorrectly classified as low-risk is independent of my race.
  - Equal False Negative Rate:  $\frac{FN1}{FN1 + TP1} = \frac{FN2}{FN2 + TP2}$
- Decision-maker: the ratio of people who recidivated among the ones labeled high-risk is independent of race.
  - Equal predictive value:  $\frac{TP1}{TP1 + FP1} = \frac{TP2}{TP2 + FP2}$

## Impossibility Result [Kleignberg et al. 2016]

The above three conditions cannot be satisfied simultaneously, unless the predictor is perfect or the two groups are the same.

# There are more definitions of fairness

- If there are **intrinsic** differences between the groups, is it fair to ignore them?
- Another setup
  - **A**: Sensitive attributes (e.g., race)
  - **Y**: True labels (e.g., commit a crime in the future)
  - **C**: Predictions (e.g., predictions of recidivism)
- Criteria:
  - **C** independent of **A**
  - **C** independent of **A** conditional on **Y**
  - **Y** independent of **A** conditional on **C**

The same impossibility results exist!

# There are more definitions of fairness

- FAT\*18 Tutorial: 21 Definitions of Fairness and Their Politics. Arvind Narayanan.

## Ways to evaluate binary classifiers

		True condition			
		Condition positive	Condition negative		
Predicted condition	Total population			Prevalence = $\frac{\sum \text{Condition positive}}{\sum \text{Total population}}$	Accuracy (ACC) = $\frac{\sum \text{True positive} + \sum \text{True negative}}{\sum \text{Total population}}$
	Predicted condition positive	<b>True positive,</b> Power	<b>False positive,</b> Type I error	Positive predictive value (PPV), Precision = $\frac{\sum \text{True positive}}{\sum \text{Predicted condition positive}}$	False discovery rate (FDR) = $\frac{\sum \text{False positive}}{\sum \text{Predicted condition positive}}$
	Predicted condition negative	<b>False negative,</b> Type II error	<b>True negative</b>	False omission rate (FOR) = $\frac{\sum \text{False negative}}{\sum \text{Predicted condition negative}}$	Negative predictive value (NPV) = $\frac{\sum \text{True negative}}{\sum \text{Predicted condition negative}}$
		True positive rate (TPR), Recall, Sensitivity, probability of detection = $\frac{\sum \text{True positive}}{\sum \text{Condition positive}}$	False positive rate (FPR), Fall-out, probability of false alarm = $\frac{\sum \text{False positive}}{\sum \text{Condition negative}}$	Positive likelihood ratio (LR+) = $\frac{\text{TPR}}{\text{FPR}}$	Diagnostic odds ratio (DOR) = $\frac{\text{LR+}}{\text{LR-}}$  F <sub>1</sub> score = $\frac{2}{\frac{1}{\text{Recall}} + \frac{1}{\text{Precision}}}$
		False negative rate (FNR), Miss rate = $\frac{\sum \text{False negative}}{\sum \text{Condition positive}}$	True negative rate (TNR), Specificity (SPC) = $\frac{\sum \text{True negative}}{\sum \text{Condition negative}}$	Negative likelihood ratio (LR-) = $\frac{\text{FNR}}{\text{TNR}}$	

364 impossibility theorems?




# Individual Fairness

- Cucumbers and Grapes: <https://youtu.be/-KSryJXDpZo>



# Individual Fairness

- Similar people should be treated similarly
- Challenges
  - What do we mean by similar people
    - Need to define some kind of “distance” measure
  - What do we mean by being treated similarly
    - Decisions based on **threshold** won't work
    - Need to impose some “smooth” notion
    - Randomization is often required



# Counterfactual Fairness

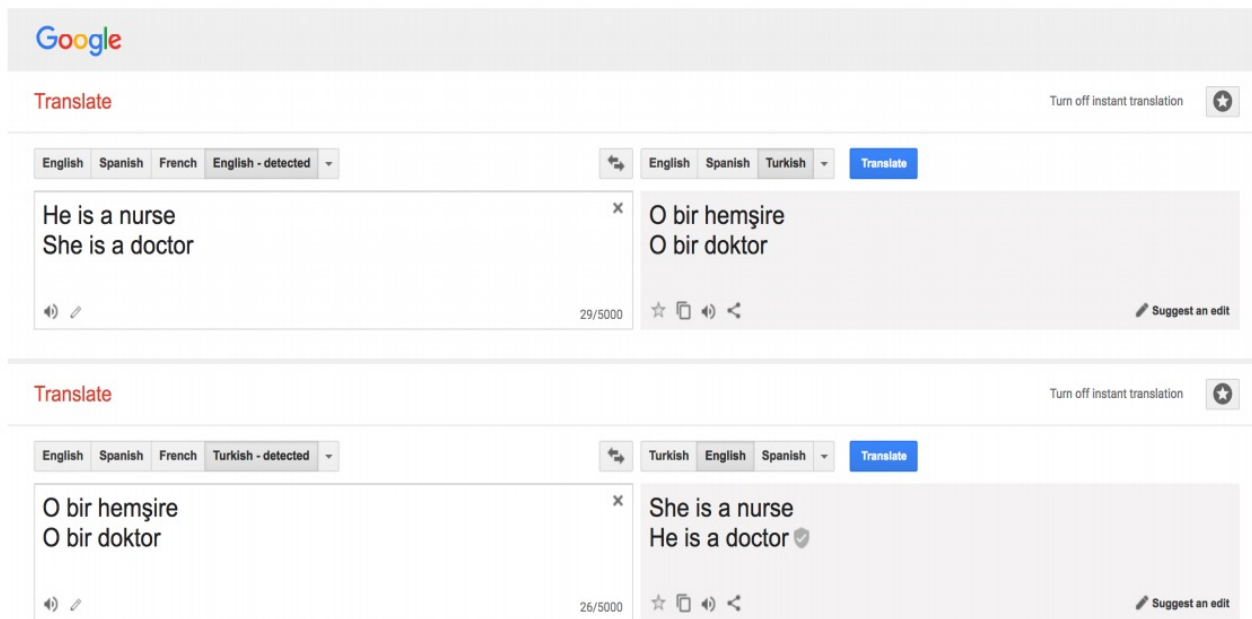
- A decision is fair towards an individual if it gives the same predictions in
  - (a) the observed world and
  - (b) a world where the individual had always belonged to a different demographic group



# Procedural Fairness (Procedural Justice)



# Representation



[Kay et al., 2015]

[Caliksan et al., 2017]

# Take-Aways

- As AI/ML becomes more ubiquitous in our daily decision making, ethical issues are getting more important as well.
- **Being aware** of the issues is the important first step!
- "Solving" the issues (if at all possible) requires communications among people in different disciplines.
- More references:
  - NIPS17 Tutorial: <https://vimeo.com/248490141>
  - FAT\*18 Tutorial: <https://www.youtube.com/watch?v=jlXluYdnnyk>
  - FAT\*19 Putting fairness in practice: <https://algorithmicbiasinpractice.wordpress.com>